
14.273 Industrial Organization: Pset4

Dave Holtz, Jeremy Yang

May 18, 2017

1. Model setup.

Following the notations in Rust (1987), HZ's flow utility is:

$$u(x_t, i_t, \theta_1) + \epsilon_t(i_t) = \begin{cases} -RC - c(0, \theta_1) + \epsilon_t(1) & i_t = 1 \\ -c(x_t, \theta_1) + \epsilon_t(0) & i_t = 0 \end{cases}$$

where RC is the replacement cost, x_t is the observed state variable for mileage, $c(\cdot)$ is cost function and i_t is the decision to replace engine and $\epsilon_t(\cdot)$ is action specific and type I extreme value distributed structural error (or unobserved state variable).

The state transition probability is given by:

$$\theta_{3j} = \mathbb{P}(x_{t+1} = x_t + j | x_t, i_t = 0)$$

$j \in \{0, 1, 2\}$ and if $i_t = 1$ then $x_{t+1} = 0$ with probability 1.

HZ chooses i_t in every period t to maximize an infinite sum of discounted flow utilities. The maximal value is defined as the value function (suppress the dependency on θ_1, θ_3):

$$V(x_1, \epsilon_1) := \max_{i_t, t \in \{1, 2, \dots\}} \mathbb{E} \left[\sum_{t=1}^{\infty} \beta^{t-1} (u(x_t, i_t, \theta_1) + \epsilon_t(i_t)) \right]$$

Rewrite the value function as in the Bellman optimality form:

$$V(x_t, \epsilon_t) = \max_{i_t} (u(x_t, i_t, \theta_1) + \epsilon_t(i_t)) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t]$$

where the expectation is with respect to (conditional) state transition probability of both x and ϵ , see Rust (1987) equation (4.5). The Bellman equation breaks the dynamic optimization problem into an infinite series of static choices.

2. (1) The choice specific value function can be derived by plugging a specific action into the value function:

$$\tilde{V}(x_t, \epsilon_t, i_t) = \begin{cases} -RC - c(0, \theta_1) + \epsilon_t(1) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 1] \\ -c(x_t, \theta_1) + \epsilon_t(0) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 0] \end{cases}$$

$$V(x_t, \epsilon_t) = \max\{\tilde{V}(x_t, \epsilon_t, 1), \tilde{V}(x_t, \epsilon_t, 0)\}$$

HZ's decision is about trading off the total (future) cost of maintaining an old engine and the lump sum cost of replacing to a new one. The time to replace is the stopping time in this problem, so it can be thought as an optimal stopping time problem where the optimal policy is characterized by a cutoff in x , HZ would choose to replace the engine if x is above that threshold (the threshold depends on realized value of ϵ).

- (2) It's clear from 2 (1) that the optimal stopping rule is:

$$\begin{aligned} & -RC - c(0, \theta_1) + \epsilon_t(1) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 1] > \\ & -c(x_t, \theta_1) + \epsilon_t(0) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 0] \end{aligned}$$

or,

$$\tilde{V}(x_t, \epsilon_t, 1) > \tilde{V}(x_t, \epsilon_t, 0)$$

therefore, because the errors are type I extreme value distributed:

$$\mathbb{P}(i_t = 1 | x_t) = \frac{\exp(u(x_t, 1, \theta_1) + \beta \mathbb{E}[V_{t+1} | x_t, i_t = 1])}{\sum_{k=\{0,1\}} \exp(u(x_t, k, \theta_1) + \beta \mathbb{E}[V_{t+1} | x_t, i_t = k])} \quad (2.1)$$

where $u(x_t, i_t, \theta_1)$ is defined in 1 and for convenience:

$$V_{t+1} := V(x_{t+1}, \epsilon_{t+1})$$

- (3) For discrete x , under the assumption that the errors are type I extreme value distributed, we have (Rust (1987) equation (4.14)):

$$EV(x, i) = \sum_y \log\left\{ \sum_j \exp[u(y, j) + \beta EV(y, j)] \right\} \cdot p(y | x, i) \quad (2.2)$$

where

$$EV(x, i) := \mathbb{E}[V_{t+1} | x_t, i_t]$$

and x, i are the state and choice of current period and y, j are the state and choice of the next period. Also note that here the transition probability does not depend on x_t but only on j (or Δx). To compute expected value function, we first need to estimate transition probability from the data, this can be done simply by counting:

$$\hat{\theta}_{30} = \frac{\sum_b \sum_t 1_{\{x_{bt+1} - x_{bt} = 0, i_{bt} = 0\}}}{\sum_b \sum_t 1_{\{i_{bt} = 0\}}}$$

$$\hat{\theta}_{31} = \frac{\sum_b \sum_t 1_{\{x_{bt+1}-x_{bt}=1, i_{bt}=0\}}}{\sum_b \sum_t 1_{\{i_{bt}=0\}}}$$

$$\hat{\theta}_{32} = \frac{\sum_b \sum_t 1_{\{x_{bt+1}-x_{bt}=2, i_{bt}=0\}}}{\sum_b \sum_t 1_{\{i_{bt}=0\}}}$$

we compute the expected value function in the inner loop of the nested fixed point algorithm (holding the value of θ fixed), we first guess the initial values of $EV(x, i)$ for all possible values of x, i and use the equation (2.2) to iterate expected value function until it converges. The criterion is:

$$\max_{x,i} |EV^{T+1}(x, i) - EV^T(x, i)| < \eta$$

The plot for $x = 1 - 30$ at the true value of parameters are shown in Figure 1.

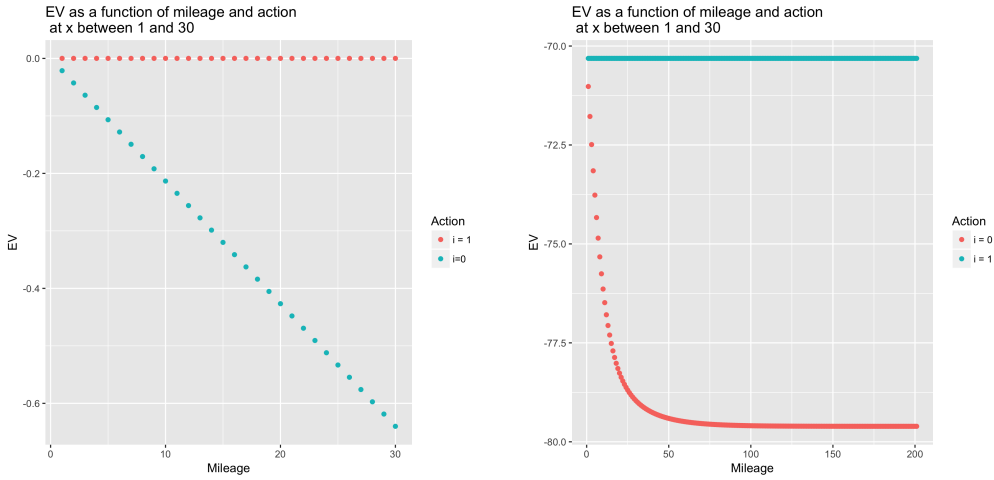


Figure 1: Expected Value Function for $i = 0$ and $i = 1$. Left panel shows results using iterative method, right panel shows provided Rust results.

Interestingly, our EV results are linear in mileage, which is probably not expected. Despite a good amount of debugging, we have been unable to identify a problem. However, it's also unclear how our calculated results for *just* 30 states should compare to the provided EV results, which provide information on 200 states. The first 30 states of the provided Rust EV estimates are decreasing in approximately linear fashion, suggesting our estimates might not be *so* bad. However, the order of magnitude of our EV values (e.g., 10^{-1}) is much smaller than the order of magnitude of EV values in the provided dataset (e.g., ~ 70), suggesting something is probably wrong. However, we don't have any more time to debug this, so we simply moved on.

- (4) The provided dataset contains mileage and engine replacement information for 100 buses over 1,000 periods. The table below shows the mean mileage, maximum mileage, minimum mileage, standard deviation of the mileage, the average mileage at engine replacement across all buses and periods, and the

average number of engine replacements for a particular bus over the 1,000 periods.

| avg miles | max miles | min miles | s.d. miles | avg replace miles | avg replacements |
|-----------|-----------|-----------|------------|-------------------|------------------|
| 8.245 | 33.000 | 0.000 | 5.709 | 15.953 | 52.980 |

We might also be interested in understanding how each of these summary statistics vary across buses. For instance, maybe some buses have their engines replaced much more often. In order to study this, Figure 2 shows the distributions of average mileage, maximum mileage, s.d. mileage, avg miles at replacement, and number of replacements across the 100 buses in the sample. In general, these distributions are quite concentrated, suggesting that there are not systematic differences across buses.

The final, bottom right plot in 2 also shows the empirically observed conditional choice probability as a function of state (mileage) that Harold Zurcher actually acts on. At a high level, Zurcher's has to make the investment decision of when to replace a given bus's engine. The mean replacement mileage plot suggests that on average he replaces a bus's engine after about 80,000 miles. The conditional choice probability plot suggests that the likelihood he increases the engine is practically zero until the bus hits 50,000 miles, after which the probability that the bus has its engine replaced climbs quickly. By the time a bus has 150,000 miles on it, it has a 50% probability of having its engine changed in a given time period.

3. (1) In the outer loop we search over a grid of values for (θ_1, β, RC) , and compute the log likelihood function:

$$\log L = \sum_b \left\{ \sum_t \log \mathbb{P}(i_{bt} | x_{bt}) + \sum_t \log \mathbb{P}(x_{bt} | x_{bt-1}, i_{t-1}) \right\}$$

where b indexes for bus and t indexes for time period. We compute a log likelihood for each combination of values for (θ_1, β, RC) and choose the set of parameters that maximizes the log-likelihood of the data. The maximum likelihood parameters obtained with the Rust method are:

$$\begin{aligned} \theta_1 &= 0.1 \\ \beta &= 0.99 \\ RC &= 6 \end{aligned}$$

- (2) In Hotz-Miller's approach, we will estimate the choice specific value function (as opposed to the expected value function as in Rust). We start by noting that conditional choice probability is observed directly from the data:

$$\hat{\mathbb{P}}(i = 1 | x) = \frac{\sum_b \sum_t 1_{\{i_{bt}=1, x_{bt}=x\}}}{\sum_b \sum_t 1_{\{x_{bt}=x\}}}$$

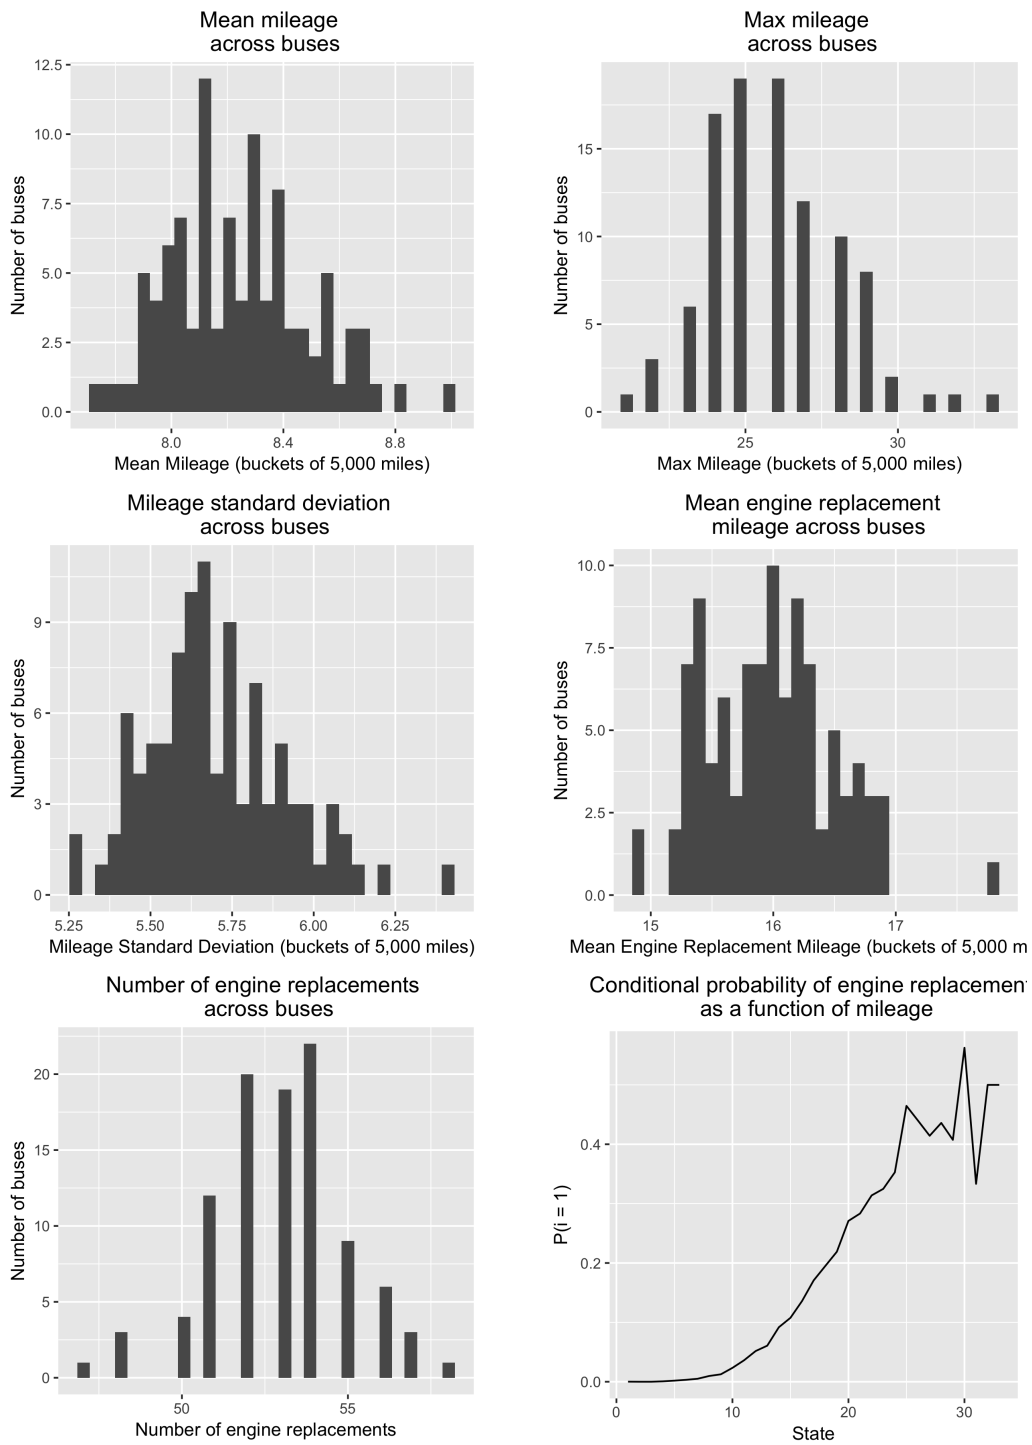


Figure 2: Distribution of various summary statistics across buses, as well as the empirical conditional choice probability for Zurcher.

The choice-specific value function (minus the structural error, and suppressing the dependency on θ_1, θ_3) can be presented recursively in the following form:

$$\tilde{V}(x_t, i_t) = u(x_t, i_t) + \beta \mathbb{E}_{x_{t+1}} [\mathbb{E}_{i_{t+1}} [\mathbb{E}_{\epsilon_{t+1}} [u(x_{t+1}, i_{t+1}) + \epsilon_{t+1} + \beta(\dots) | i_{t+1}, x_{t+1}] | x_{t+1}] | x_t, i_t]$$

where (\dots) represents higher (two and above) period forward expectations. In principle it's an infinite loop but in practice we need to stop at some T , for example, when $T = 2$, (\dots) simplifies to:

$$(\dots) = \mathbb{E}_{x_{t+2}} [\mathbb{E}_{i_{t+2}} [\mathbb{E}_{\epsilon_{t+2}} [u(x_{t+2}, i_{t+2}) + \epsilon_{t+2} | i_{t+2}, x_{t+2}] | x_{t+2}] | x_{t+1}, i_{t+1}]$$

For simplicity, in the code we use one-period forward simulation where:

$$\tilde{V}(x_t, i_t) = u(x_t, i_t) + \beta \mathbb{E}_{x_{t+1}} [\mathbb{E}_{i_{t+1}} [\mathbb{E}_{\epsilon_{t+1}} [u(x_{t+1}, i_{t+1}) + \epsilon_{t+1} | i_{t+1}, x_{t+1}] | x_{t+1}] | x_t, i_t]$$

it is estimated as:

$$\hat{V}(x_t, i_t) = \frac{1}{S} \sum_s [u(x_t, i_t) + \beta [u(x_{t+1}^s, i_{t+1}^s) + \gamma - \log(\hat{\mathbb{P}}(i_{t+1}^s | x_{t+1}^s))]]$$

where x_{t+1}^s is drawn from the transition probability $\hat{\theta}_{30}, \hat{\theta}_{31}, \hat{\theta}_{32}$, and i_{t+1}^s is drawn from $\hat{\mathbb{P}}(i | x)$, γ is the Euler's constant. We only go one period forward because we only observe data for states up to $x_t = 33$. It is possible for larger T that we would encounter a state that is not in our dataset. When this occurs, it's unclear what value should be used as the conditional choice probability. While we avoid this issue by setting $T = 2$, this does reduce the precision of our estimates.

- (3) In order to determine which engine HZ prefers, we simply need to look at HZ's value function for both engines at $t = 0$ (which corresponds to $x_t = 0$ for all buses). There are a number of different mileage evolution paths that any given bus could take. However, the ex ante value function at time = 0 provides a weighted average of all of these scenarios. So at the outset, he will prefer whichever engine provides the most value in expectation.

[paragraph with the results / which engine is preferred].

- (4) We want to compute HZ's demand function for the two buses, which we will denote as engine 1 ($\theta_1 = 0.09, RC = 6$) and engine 2 ($\theta_1 = 0.02, RC = 20$) as a function of RC. In order to do so, we obtain conditional choice probability estimates, $\hat{\mathbb{P}}(i = 1 | x)$ by using the Rust method to iterate EV values. We use the Rust methodology because the Hotz and Miller methodology depends on the observed conditional choice probabilities, which we know do *not* correspond to the counterfactual engine 2.

With those conditional choice probability estimates for the two engines in hand, we run 1,000 simulations of a bus's state transitions (and HZ's corresponding engine replacement decisions) over the first 15 periods. This allows

us to get an expected, per-bus demand for engines over the first 15 periods. In order to get the expected demand that HZ has for engines across all buses, we simply multiply this figure by 100. So the expected demand (as a function of period t) is:

$$D(t) = 100 \times \sum_{x_t} \hat{\mathbb{P}}(i = 1 | x = x_t) \hat{\mathbb{P}}(x = x_t | t) \quad (1)$$

HZ's demand for engines as a function of the period, t for a few values of RC can be found in Figure 3. The average per-period demand for the two engines (averaged across 15 periods) for different values of RC can be found in Figure 4. it's worth noting that the demand curves for the two engines appear almost identical - although you cannot distinguish them in the figure, there are small differences (on the order of a tenths of an engine. This could be a true difference, or it could be simulation error. Although we're not sure why these demand curves are so similar, we have two hypotheses:

1. Whatever our EV estimation issue, it is rearing its ugly head again and making these demand curves very similar.
 2. The increase in RC from engine 1 to engine 2 is almost perfectly offset by the decrease in θ_1 , creating two extremely similar demand curves.
- (5) To determine the total value of the engines, assuming marginal cost RC , we can simply compute the total area to the right of a given RC in a demand curve that looks like Figure 4. This area will give the total surplus that HZ gets from the engine in a given period. In order to get a total value, we simply multiply this by the number of periods we want to consider. Mathematically, the total value is

$$V_{engine}(RC) = n \times \int_{RC}^{\infty} D_{engine}(p) dp, \quad (2)$$

where n is the number of periods and $D(p)$ is the amount of demand that HZ would have a given RC.

For engine 1, the total value is \llbracket . For engine 2, the total value is \llbracket . R&D costs are essentially a fixed cost, so the maximum HZ should be willing to pay for the engine with $\theta_1 = 0.02$ is the difference in the total values of the two engines for a given value of RC. Given $RC = 20$ for the second engine and $RC = 6$ (our estimate) for the first engine, the maximum HZ would be willing to pay is \llbracket .

APPENDIX: CODE

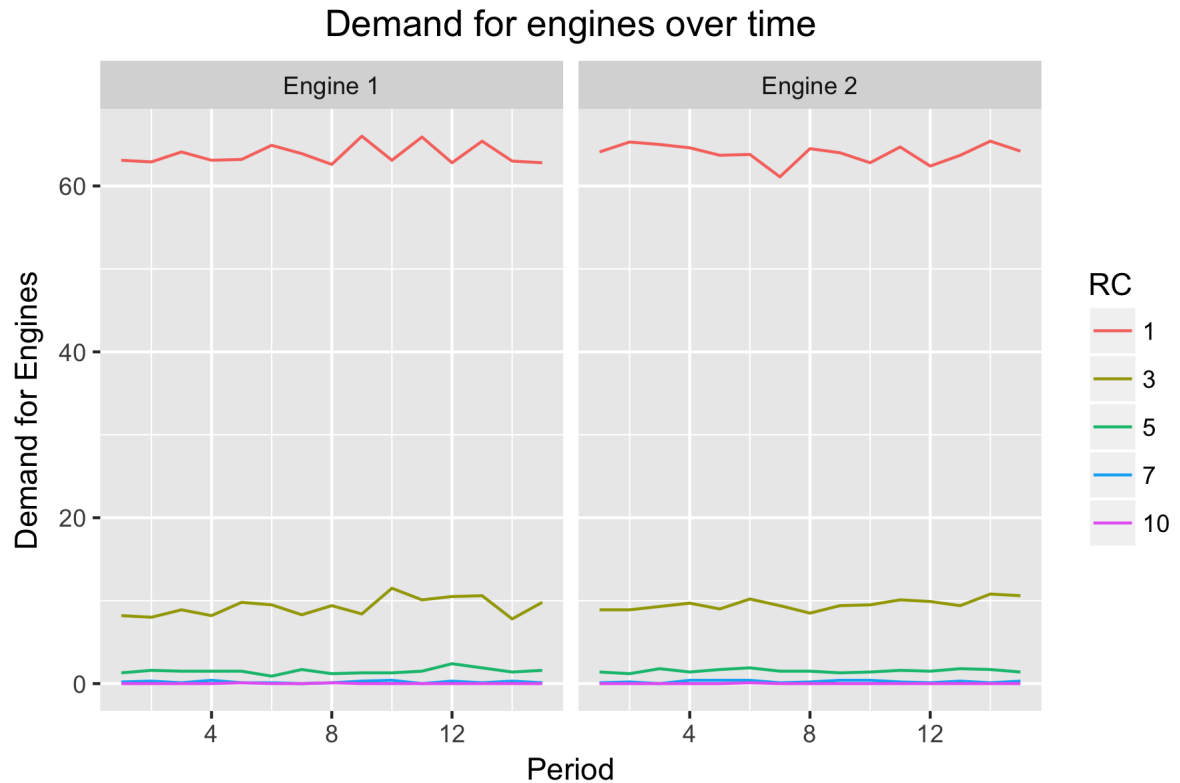


Figure 3: The demand for engines across a fleet of 100 buses as a function of period (over the first 15 periods) for different values of RC. Unsurprisingly, when RC is lower, HZ is much more willing to change bus engines.

```

1 ##### This code uses the methodologies of both Rust (1987) and Hotz and
2 ##### Miller (1993) to estimate the parameters of a single
3 ##### agent dynamic problem where an agent (Harold Zurcher) must choose when
4 ##### to have the engines replaced in a fleet of buses.
5
6 ## Import libraries
7 library(R.matlab)
8 library(ggplot2)
9 library(dplyr)
10
11 ## Read in data
12 setwd('~ /Dropbox_ /MIT /Spring_2017 /14.273 /HW4 /273-pset4 /')
13 data <- readMat('.. /rust.mat')
14 gamma = .577
15
16 ## Extract bus replacement events
17 i <- data$it
18 ## Extract bus mileage counts (in increments of 5,000 miles)
19 x <- data$xt

```





Figure 4: Aggregate per-period demand for new engines across a fleet of 100 buses as a function of RC. For engine 1, $\theta_1 = 0.09$. For engine 2, $\theta_1 = 0.02$.

```

19
20 ## Buses transition from different mileage states , and can jump forward
    zero, one, or two 5,000 mile buckets. This block
21 ## of code estimates the transition probabilities empirically from the data
    .
22
23 ### Initialize empty vectors to hold a count of how many times each jump
    happens
24 zero <- vector()
25 one <- vector()
26 two <- vector()
27
28 ### Loop through the 100 buses in the dataset
29 for (k in 1:100) {
30
31     ### Given a bus k, grab the mileage counts
32     xk <- x[,k]
33     ### Also grab the engine replacement events
34     ik <- i[,k]
35     ### Get a modified array which gives the change in mileage buckets
        from period j to period j+1
36     jk <- xk[-1]-xk[-1000]

```

```

37
38     ### We only care about periods where i=0 for transition
        probabilities , since i=1 will always send
39     ### x back to 0. This selects out only time periods for this bus
        where i = 0
40     j <- jk[ik==0]
41
42     ### This counts up how many times the mileage bucket counter , x,
        moves up by 0, 1, or 2 when i=0
43     zero[k] <- length(j[j==0])
44     one[k] <- length(j[j==1])
45     two[k] <- length(j[j==2])
46 }
47
48 ## Estimate the x_t-independent transition probabilities by dividing the
        number of times for each transition by the
49 ## total number of transitions
50 theta_30 = sum(zero)/(sum(zero)+sum(one)+sum(two))
51 theta_31 = sum(one)/(sum(zero)+sum(one)+sum(two))
52 theta_32 = sum(two)/(sum(zero)+sum(one)+sum(two))
53
54 #####
55 # Question 2.3 #
56 #####
57
58 ##### We'll now take the true values of the parameter values as given , and
        use the method described in Rust (1987) to iteratively
59 ##### estimate the value function (or in this case , the EV function).
60
61 ## Initialize parameters to their true values
62 theta_1 = .05
63 theta_30 = .3
64 theta_31 = .5
65 theta_32 = .2
66 beta =.99
67 RC = 10
68
69 ### Define the linear cost function. If an engine is not replaced , the bus
        incurs cost theta_1*x, so cost
70 ### increases linearly as a bus gets older.
71 cost <- function(x){
72     return (theta_1*x)}
73
74 ### Define the utility function at mileage x from action i. If the agent
        chooses to replace the engine in a bus ,
75 ### it costs RC. If they choose _not_ to replace the engine , they incur the
        cost of running the bus at mileage x.
76 u <- function(x,i){
77     -RC*i - cost(x*(1-i))
78 }
79
80
81 ##### The value function can be estimated through an iteration procedure. We
        start with some initial guess for EV,

```

```

82 ##### calculate EV with an expression that includes our initial guess of EV,
    and continue iterating until the difference
83 ##### between subsequent EV estimates becomes small.
84
85 ##### value.Iterate is a function to iteratively update the value function
    according to the methodology in Rust. The function
86 ##### takes as an argument a current estimate of EV, and returns an updated
    estimate of EV. EV is an x by d matrix – we want the
87 ##### EV values for each decision d at every possible current mileage value
    x.
88 value.Iterate <- function(EV){
89
90     ##### First iterate through each of the 30 x states
91     for (x in 1:31){
92         ## Update the EV value corresponding to not replacing the engine. There
            are three contributions here – one from the
93         ## j = 0 case, one from the j = 1 case, and one from the j = 2 case.
            Note the indexing here. When x =1, the state is
94         ## equal to 0 (this is the x that needs to be passed into u()), but we
            want to grab the EV corresponding to the 1st entry.
95         EV2[x,1] <- log(exp(u(x-1,0)+beta*EV[x,1])+exp(u(x-1,1)+beta*EV[x,2]))*
            theta_30 + gamma
96         + log(exp(u(x,0)+beta*EV[x+1,1])+exp(u(x,1)+beta*EV[x+1,2]))*theta_31
            + gamma
97         + log(exp(u(x+1,0)+beta*EV[x+2,1])+exp(u(x+1,1)+beta*EV[x+2,2]))*theta_
            32 + gamma
98
99         ## Update the EV value corresponding to replacing the engine. When the
            engine is replaced, x at the next period will
100        ## deterministically reset to x = 0.
101        EV2[x,2] <- log(exp(u(0,0)+beta*EV[1,1])+exp(u(0,1)+beta*EV[1,2])) +
            gamma
102    }
103
104    ## Return the updated EV values.
105    return(EV2)
106 }
107
108 ##### Set a critical value for to measure the deviation between iterative
    updates of EV. The distance between the two EV matrices
109 ##### is the infinity norm of the difference
110 cri <- 10^(-8)
111
112 ##### Set an initial value for the EV matrix (all 0s, EV), and another EV
    object to hold the updated estimates, EV2.
113 EV <- matrix(100,33,2)
114 EV2 <- matrix(0,33,2)
115
116 ## While the infinity norm is less than the threshold, iterate
117 while(max(abs(EV-EV2))>cri){
118
119     ##### Set the current EV to the previous updated EV
120     EV <- EV2
121     ##### Compute a new updated EV by iterating on the current EV

```

```

122 EV2 <- value.Iterate(EV)
123 }
124
125 ### Do one last update to set EV equal to the last EV2
126 EV <- EV2
127
128 # get EV(x,i) for x=0,1,2,...,30
129 ### EV contains extra states, which we needed to compute the above
    computation. Throw them away.
130 EV <- EV[1:31,]
131
132 ### Plot the EV of both replacing the engine (i = 1) and not replacing the
    engine (i = 0) at every x
133 ### between 1 and 30
134 df <- data.frame('x'=c(1:30, 1:30), 'EV'=c(EV[2:31,1], EV[2:31,2]), 'Action'
    = c(rep('i_0', 30), rep('i_1', 30)))
135
136 ### Generate a plot that compares the EV of replacing the engine and not
    replacing the engine
137 ev_plot <- ggplot(df, aes(x=x, y=EV, color=Action)) + geom_point() + xlab('
    Mileage') + ylab('EV') +
138   ggtitle('EV_as_a_function_of_mileage_and_action\nat_x_between_1_and_30'
    ) +
139   theme(plot.title = element_text(hjust = 0.5))
140 ggsave(ev_plot, file='ev_plot.png', height=6, width=6, units='in')
141
142 ### This is a plot to see the EV data in the attached rust matlab file. The
    state space is different than ours (200 states),
143 ### so its hard to compare. Our's is linear (seems wrong), whereas the
    provided data is not. However, the first 30 states
144 ### _do_ look approximately linear, so maybe we're not so far off.
145
146 df_rust <- data.frame('x'=c(seq(1,201), seq(1, 201)), 'EV'=c(data$EV[,1],
    data$EV[,2]), 'Action' = c(rep('i_0', 201),
147
148
149
150 ev_plot_rust <- ggplot(df_rust, aes(x=x, y=EV, color=Action)) + geom_point
    () + xlab('Mileage') + ylab('EV') +
    ggtitle('Rust_dataset_EV_as_a_function_of_mileage_and_action\nat_x_
    between_1_and_201') +
    theme(plot.title = element_text(hjust = 0.5))

```

```

151 | ggsave(ev_plot_rust, file='ev_plot_rust.png', height=6, width=6, units='in'
152 | )
153 | #####
154 | # Question 2.4 #
155 | #####
156 |
157 | ### Calculate the mean mileage, mean time to engine replacement, max
158 |     mileage, min mileage, and sd mileage over the whole sample
159 | mean_x <- mean(x)
160 | mean_engine_replacement_age <- mean(x[i == 1])
161 | max_x <- max(x)
162 | min_x <- min(x)
163 | sd_x <- sd(x)
164 | avg_replacements <- mean(apply(i, 2, function(x) {sum(x)}))
165 |
166 | aggregate_stats <- c(mean_x, max_x, min_x, sd_x, mean_engine_)
167 |     round(., 3) %>%
168 |     kable(., format='latex')
169 |
170 | ### Calculate the per bus mean mileage, mean time to engine replacement,
171 |     max mileage, min mileage, and sd mileage
172 | mean_x_per_bus <- apply(x, 2, function(x) {mean(x)})
173 | max_x_per_bus <- apply(x, 2, function(x) {max(x)})
174 | min_x_per_bus <- apply(x, 2, function(x) {min(x)})
175 | sd_x_per_bus <- apply(x, 2, function(x) {sd(x)})
176 | mean_engine_replacement_per_bus <- apply(x*i, 2, function(x) {sum(x)/sum(x
177 |     != 0)})
178 | replacements <- apply(i, 2, function(x) {sum(x)})
179 |
180 | ### Collate per bus information into a dataframe
181 | per_bus_statistics <- data.frame(bus = seq(1, 100, 1),
182 |     mean_x_per_bus = mean_x_per_bus,
183 |     max_x_per_bus = max_x_per_bus,
184 |     min_x_per_bus = min_x_per_bus,
185 |     sd_x_per_bus = sd_x_per_bus,
186 |     mean_engine_replacement_per_bus = mean_
187 |         engine_replacement_per_bus,
188 |     replacements = replacements)
189 |
190 | ### Create some plots
191 | mean_mileage_plot <- ggplot(per_bus_statistics, aes(x=mean_x_per_bus)) +
192 |     geom_histogram() +
193 |     xlab('Mean_Mileage_(buckets_of_5,000_miles)') + ylab('Number_of_buses') +
194 |     ggtitle('Mean_mileage_\n_across_buses') +
195 |     theme(plot.title = element_text(hjust = 0.5))
196 | ggsave(mean_mileage_plot, file='mean_mileage_plot.png', height=4, width=4,
197 |     units='in')
198 |
199 | max_mileage_plot <- ggplot(per_bus_statistics, aes(x=max_x_per_bus)) + geom
200 |     _histogram() +

```

```

195   xlab('Max_Mileage_(buckets_of_5,000_miles)') + ylab('Number_of_buses') +
196   ggtitle('Max_mileage_\n_across_buses') +
197   theme(plot.title = element_text(hjust = 0.5))
198   ggsave(max_mileage_plot, file='max_mileage_plot.png', height=4, width=4,
199         units='in')
200
201   sd_mileage_plot <- ggplot(per_bus_statistics, aes(x=sd_x_per_bus)) + geom_
202     histogram() +
203     xlab('Mileage_Standard_Deviation_(buckets_of_5,000_miles)') + ylab('
204       Number_of_buses') +
205     ggtitle('Mileage_standard_deviation_\n_across_buses') +
206     theme(plot.title = element_text(hjust = 0.5))
207     ggsave(sd_mileage_plot, file='sd_mileage_plot.png', height=4, width=4,
208           units='in')
209
210   time_to_engine_replacement_plot <- ggplot(per_bus_statistics, aes(x=mean_
211     engine_replacement_per_bus)) + geom_histogram() +
212     xlab('Mean_Engine_Replacement_Mileage_(buckets_of_5,000_miles)') + ylab('
213       Number_of_buses') +
214     ggtitle('Mean_engine_replacement_\n_mileage_across_buses') +
215     theme(plot.title = element_text(hjust = 0.5))
216     ggsave(time_to_engine_replacement_plot, file='time_to_engine_replacement_
217       plot.png', height=4, width=4, units='in')
218
219   replacements_plot <- ggplot(per_bus_statistics, aes(x=replacements)) + geom
220     _histogram() +
221     xlab('Number_of_engine_replacements') + ylab('Number_of_buses') +
222     ggtitle('Number_of_engine_replacements_\n_across_buses') +
223     theme(plot.title = element_text(hjust = 0.5))
224     ggsave(replacements_plot, file='replacements_plot.png', height=4, width=4,
225           units='in')
226
227   #####
228   # Question 3.1 #
229   #####
230
231   ### This code will estimate the parameters beta, theta_1, and RC using the
232     nested fixed-point algorithm
233     described in Rust.
234
235   ### A function to compute the probability of Zurcher's choices using the
236     EVs we calculate using the EV calculation
237     framework above. The probability of choosing different actions
238     basically acts like a multichoice logit function.
239   choice.prob.Estimate <- function(){
240
241     ### Initialize an empty matrix to hold choice probability estimates.
242     p_i <- matrix(0,31,2)
243
244     ### Iterate through all of the possible mileage states, x.
245     for (x in 1:31){
246       ## For each mileage state x, calculate the probability that Zurcher
247         will choose i = 0.

```

```

235     p_i[x,1] <- exp(u(x-1,0)+beta*EV[x,1])/(exp(u(x-1,0)+beta*EV[x,1])+exp(
236       u(x-1,1)+beta*EV[x,2]))
237     ## Calculate the probability of choosing i = 1 at each state, which is
238     just 1 - P(i = 0).
239     p_i[x,2] <- 1- p_i[x,1]
240   }
241   #### Return the updated p_i object
242   return(p_i)
243 }
244 #### A function to calculate the total log likelihood of the observed data
245 given a set of parameters. This method assumes that
246 the probabilities across periods and buses are independent, so we can
247 just add up all of the log probabilities.
248 log_likelihood.Compute <- function() {
249   #### Initialize 0-valued variables to hold the log choice probability, the
250   log transition probability,
251   and the sum of the two.
252   log_choice_prob <- 0
253   log_transition_prob <- 0
254   total <- 0
255   #### Iterate over buses
256   for (bus in 1:100) {
257     #### Iterate over time periods
258     for (t in 1:999) {
259       #### We special case mileage states greater than 30, since they are a
260       bit strange in our data. Otherwise, we calculate
261       the choice probability using the current value of p_i according
262       to the EV values we calculated to get the
263       choice probability. Take the log and add it to the current
264       running value.
265       if (x[t,bus] <= 30){
266         log_choice_prob <- log(p_i[x[t,bus]+1,i[t,bus]+1]) + log_choice_
267         prob
268       #### Do the same thing for our special cased, x > 30 case.
269       } else {
270         log_choice_prob <- log(p_i[31,i[t,bus]+1]) + log_choice_prob
271       }
272       #### Calculate over the transitions for each bus the sum of the log
273       transition probabilities. We have our estimates of
274       theta_3 given the empirical transition probabilities. So we can
275       just grab that for each observed transition and add it
276       to the total log transition probability.
277       #### First we do the j = 0 case.
278       if (x[t+1,bus]-x[t,bus]==0) {
279         log_transition_prob <- log(theta_30)+log_transition_prob
280       #### Then the j = 1 case.
281       } else if (x[t+1,bus]-x[t,bus]==1) {
282         log_transition_prob <- log(theta_31)+log_transition_prob

```

```

278     ### And finally the j = 2 case.
279   } else if (x[t+1,bus]-x[t,bus]==2) {
280     log_transition_prob <- log(theta_32)+log_transition_prob
281   }
282 }
283
284   ### Now, get the total log likelihood by adding up all of the
285   transition components and the choice components.
286   total <- (log_choice_prob+log_transition_prob) + total
287 }
288 return (total)
289 }
290
291 ### Now we're actually going to use the nested fixed point algorithm to get
292 the maximum likelihood estimates of the parameters
293 that we care about. This process has three steps.
294
295 ### Step 1: We would calculate theta_30, theta_31, and theta_31 directly
296 from the data. This step is not in the loop, and we've
297 actually already done this and it doesn't change, so we don't need to
298 do it again.
299
300 ### Step 2: Next, we are going to set up a grid over values of theta_1,
301 beta, and RC that we will calculate the
302 log likelihood to determine the maximum likelihood parameter values. We
303 'll also initialize a dataframe
304 to hold the parameter values and the log likelihoods.
305
306 theta_1_range <- seq(.01,.10,.01)
307 beta_range <- seq(.90,.99,.01)
308 RC_range <- seq(6,15,1)
309 likelihood <- data.frame( 'theta_1'=rep(0), 'beta'=rep(0), 'RC'=rep(0), 'log.
310   likelihood'=rep(0))
311
312 ### Step 3: Now we actually do the nested fixed point computation.
313
314 ### Loop through theta_1
315 for (theta_1 in theta_1_range) {
316   ### Loop through beta
317   for (beta in beta_range) {
318     ### Loop through RC
319     for (RC in RC_range) {
320       print(paste(c(theta_1, beta, RC), collapse='_'))
321
322       ### Initialize the EV functions to the initial values we used above.
323       EV <- matrix(100,33,2)
324       EV2 <- matrix(0,33,2)
325
326       ### Iteratively compute the EV values.
327       while (max(abs(EV-EV2))>cri) {
328         EV <- EV2
329         EV2 <- value.Iterate(EV)
330       }
331     }
332   }
333 }

```



```

325     EV <- EV2
326     EV <- EV[1:31,]
327
328     ### Given these values of EV, calculated the choice probabilities
329     p_i <- choice.prob.Estimate()
330
331     ### Given the EV values, the choice probabilities and the parameters,
332     calculate
333     ### the log-likelihood of the data.
334     likelihood <- rbind(likelihood, c(theta_1, beta, RC, log.likelihood.
335     Compute()))
336   }
337 }
338
339 ### Retrieve the row in the likelihood dataframe corresponding to the
340 maximum likelihood estimate
341 likelihood <- likelihood[-1,]
342 parameter_estimates <- likelihood[which.max(likelihood[,4]),]
343
344 ### Use these parameters and get the relevant estimate of EV and p_i
345 theta_1 = parameter_estimates$theta_1
346 beta = parameter_estimates$beta
347 RC = parameter_estimates$RC
348
349 EV <- matrix(100,33,2)
350 EV2 <- matrix(0,33,2)
351 ### Iteratively compute the EV values.
352 while(max(abs(EV-EV2))>cri){
353   EV <- EV2
354   EV2 <- value.Iterate(EV)
355 }
356 EV <- EV2
357 EV <- EV[1:31,]
358
359 ### Given these values of EV, calculated the choice probabilities
360 p_i <- choice.prob.Estimate()
361
362 ### Given the EV values, the choice probabilities and the parameters,
363 calculate
364 ### the log-likelihood of the data.
365 likelihood <- rbind(likelihood, c(theta_1, beta, RC, log.likelihood.Compute()))
366
367 save(EV, p_i, likelihood, parameter_estimates, file='rust_estimate.Rdata')
368
369 #####
370 # Question 3.2 #
371 #####
372
373 ### Now we will get estimates of the parameters using the Hotz and Miller
374 conditional choice probability approach. This will
375 ### allow us to compare these parameter estimates to those obtained using
376 the Rust approach.
377
378 ### First, we need to calculate the probability of the agent choosing
379 either i = 0 or i = 1 based on the state that they find

```

```

372 ##### a given bus in , x, at some time period t. This will be the baseline
      that we use to try and find the best parameter values
373 ##### (i.e., which parameter values minimize the infinity norm between these
      true probabilities and the estimated probabilities)
374
375 ##### The probability matrix
376 p_ix <- matrix(0,33,2)
377 ##### The vector of how often the agent chooses i=1 given state x
378 ones <- vector()
379 ##### The vector of how often the agent finds a bus in state x
380 total <- vector()
381
382 ##### Loop through the states
383 for (state in 0:32){
384
385     ##### For a given state , a will track how many times i = 1 and b will track
      how many times that state occurs.
386     ##### Initialize them to 0 for the given state.
387     a <- 0
388     b <- 0
389
390     ##### Loop over the buses
391     for (bus in 1:100){
392
393         ##### Increment how many times the agent chooses i = 1 in state x
394         a <- sum(i[which(x[,bus]==state),bus]) + a
395         ##### Increment how many times the state x occurs
396         b <- length(i[which(x[,bus]==state),bus]) + b
397     }
398
399     ##### Add the most recent estimates to the vector.
400     ones[state+1] <- a
401     total[state+1] <- b
402 }
403
404 ##### Based on the ones and total vectors , updated the choice probability
      matrix.
405 p_ix[,1] <- 1-ones/total
406 p_ix[,2] <- ones/total
407
408 ##### Plot conditional choice probabilities
409 p_ix_df <- as.data.frame(p_ix)
410 p_ix_df$state <- as.numeric(rownames(p_ix_df))
411 names(p_ix_df) <- c('P(i=0)', 'P(i=1)', 'State')
412 ccp_plot <- p_ix_df %>%
413   ggplot(. , aes(x=State , y='P(i = 1)')) + geom_line() +
414   ggtitle('Conditional_probability_of_engine_replacement_\n as_a_function_
      of_mileage') +
415   theme(plot.title = element_text(hjust = 0.5))
416 ggsave(ccp_plot , file='ccp_plot.png' , height=4, width=4, units='in')
417
418 ##### The function below uses the Hotz and Miller method to estimate V and p_
      ix_hat for every state and period
419 ##### given a set of model parameters (beta , theta_1, and RC).

```

```

420 approximate.V_pihat <- function() {
421   ### Initialize an empty valuation matrix
422   V <- matrix(0,33,2)
423   ### Initialize an empty conditional choice probability matrix
424   p_ix_hat <- matrix(0,33,2)
425
426   ### Iterate through the states
427   for (state in 0:30){
428     ### Initialize a and b, which will basically track a running total of V
429     ### for different choices over simulations, to 0.
430     a = 0
431     b = 0
432     ### Iterate through the simulations. Note that ideal we would probably
433     ### want to go more than one time step into the
434     ### future. However, because of the limitations in our dataset, we only
435     ### go one time step forward. This is mainly because
436     ### it's unclear how we would draw i (the choice) for states that do
437     ### not appear in our data (i.e., x = 34).
438     for (s in 1:S){
439       ## Conditional on choosing i = 0, simulate the next state that a
440       ## given bus will end up in by drawing from the
441       ## transition probabilities.
442       x_prime_0 = state + sample(c(0,1,2),1,replace = T, prob = c(theta_30,
443         theta_31,theta_32))
444       ## Conditional on choosing i = 0 and ending up in some state in the
445       ## next time period, randomly simulate a draw from
446       ## i based on the conditional choice probabilities
447       i_prime_0 = sample(c(0,1),1,replace=T,prob = c(p_ix[x_prime_0+1,1],p_
448         ix[x_prime_0+1,2]))
449       # Figure out the expected utility from this truncated sequence of
450       # choices.
451       a = (u(state,0) + beta*(u(x_prime_0,i_prime_0)+gamma-log(p_ix[x_prime
452         _0+1,i_prime_0+1]))) + a
453
454       ## Conditional on choosing i = 1, we don't need to simulate the next
455       ## state that a bus will end up in. It will always
456       ## be x = 0. So we jump right to simulating the draw from i for x =
457       ## 0.
458       i_prime_1 = sample(c(0,1),1,replace=T,prob = c(p_ix[1,1],p_ix[1,2]))
459       ## Figure out the expected utility from this truncated sequence of
460       ## choices.
461       b = (u(state,1) + beta*(u(0,i_prime_1)+gamma-log(p_ix[1,i_prime_1+1])
462         )) + b
463     }
464
465     ## Set the value of V to be the average over all S of our simulations
466     ## for both the i = 0 and i = 1 choices.
467     V[state+1,1] = a/S
468     V[state+1,2] = b/S
469     ## Use the multinomial logit-esque probability expression to figure out
470     ## the probability of choosing i = 0 or i = 1
471     ## given that the bus is in state x.
472     p_ix_hat[state+1,1] <- exp(V[state+1,1])/(exp(V[state+1,1])+exp(V[state
473       +1,2]))

```

```

457   p_ix_hat[state+1,2] <- 1- p_ix_hat[state+1,1]
458 }
459
460 # Put final output into a list and return it
461 results <- list('V' = V, 'p_ix_hat' = p_ix_hat)
462 return(results)
463 }
464
465 ### Specify a number of constants that will be used in the Hotz and Miller
466   algorithm:
467 ### S: The number of "simulations" to do per state / decision
468 ### gamma: This should be Euler's constant
469 ### theta_1_range: The range of theta_1 values to test
470 ### beta_range: The range of beta values to test
471 ### RC_range: The range of RC values to test
472 S = 1000
473 theta_1_range <- seq(.01,.10,.01)
474 beta_range <- seq(.90,.99,.01)
475 RC_range <- seq(6,15,1)
476
477 ### Initialize a dataframe to hold different parameter combinations and the
478   infinity-norm between the actual conditional
479   choice probabilities and the estimated ones
480 difference <- data.frame('theta_1'=rep(0), 'beta'=rep(0), 'RC'=rep(0),
481   'difference'=rep(0))
482
483 ### Loop through theta_1
484 for (theta_1 in theta_1_range) {
485   ### Loop through theta_2
486   for (beta in beta_range) {
487     ### Loop through RC
488     for (RC in RC_range) {
489       # Check progress
490       print(paste(c(theta_1, beta, RC), collapse='_'))
491
492       ### Get estimates of V and P_ix_hat using the Hotz and Miller method
493       v_and_p_ix_hat <- approximate.V_pixhat()
494       V = v_and_p_ix_hat$V
495       p_ix_hat <- v_and_p_ix_hat$p_ix_hat
496
497       ### Now that we have a full conditional choice probability matrix,
498         calculate the infinity norm (i.e., largest
499       ### absolute difference between the empirical conditional choice
500       probabilities and those estimated with the
501       ### given parameters)
502       difference <- rbind(difference, c(theta_1, beta, RC, max(abs(p_ix[1:31,] -
503         p_ix_hat[1:31,]))))
504     }
505   }
506 }
507
508 ### Find the set of parameters that minimizes this difference
509 difference <- difference[-1,]
510 parameter_estimates <- difference[which.min(difference[,4]),]

```

```

505
506 ### Use these parameters and get the relevant estimate of V and p_ix_hat
507 theta_1 = parameter_estimates$theta_1
508 beta = parameter_estimates$beta
509 RC = parameter_estimates$RC
510 best_guesses <- approximate.V_pihat()
511 V <- best_guesses$V
512 p_ix_hat <- best_guesses$p_ix_hat
513
514 save(V, p_ix_hat, difference, parameter_estimates, file='hotz_and_miller_
    estimate.Rdata')
515
516 #####
517 # Question 3.3 #
518 #####
519
520
521
522 #####
523 # Question 3.4 #
524 #####
525
526 ### This function simulates, for one agent, a sequence of state transitions
    and also engine replacement decisions
527 simulate_sequence <- function(n_periods) {
528     ### Initialize empty vectors to hold states and engine replacement
        transitions
529     x_values <- rep(0, n_periods)
530     i_values <- rep(0, n_periods)
531     ### Every bus starts at state 0
532     x_values[1] <- 0
533     ### Go through the progression
534     for (j in 1:length(x_values)) {
535         ### Make a decision based on current state
536         i_values[j] = sample(c(0,1),1,replace=T,prob = c(p_ix_hat[x_values[j] +
            1,1],p_ix_hat[x_values[j] + 1,2]))
537         ### If decision is to not replace, continue on and increment x randomly
538         if (i_values[j] == 0) {
539             x_values[j+1] = x_values[j] + sample(c(0,1,2),1,replace = T, prob = c
                (theta_30,theta_31,theta_32))
540             ### If decision is to replace, reset state to 0
541             } else {
542                 x_values[j+1] = 0
543             }
544         }
545     ### Generate a decision for the last period, even though we never see the
        fruits of that decision
546     i_values[length(i_values)] = sample(c(0,1),1,replace=T,prob = c(p_ix_hat[
        x_values[length(x_values)] + 1,1],
547
        p_ix_hat[
            x_
            values
            |
            length

```

```

(x_
values
)] +
1,2])
)

548   ### Return the states and replacement decisions in a list
549   results <- list('x_values' = x_values, 'i_values' = i_values)
550   return(results)
551 }
552
553   ### Given a set of parameters, this function generates period-by-period
554   demand estimates for new buses (e.g.,
555   ### how many buses will get their engine replaced in each period)
556   estimate_demand <- function(n_sims, n_buses, n_periods) {
557
558     ### Initialize a vector to hold simulated demand
559     simulated_demand_total <- rep(0, n_periods)
560
561     ### Run a bunch of simulations and simulate engine replacement decisions
562     for (j in 1:n_sims) {
563       simulated_demand_total = simulated_demand_total + simulate_sequence(n_
564         periods)$i_values
565     }
566
567     ### Divide by the number of sims to get averages, multiply by number of
568     buses (this works because
569     ###buses are independent). Then return what we get.
570     return((n_buses/n_sims)*simulated_demand_total)
571 }
572
573   ### Get demand as a function of RC for the first bus
574
575   ## Specify the range of RCs, as well as constants.
576   RC_range = seq(0, 15, .25)
577   n_periods = 15
578   n_sims = 1000
579   n_buses = 100
580   load('hotz_and_miller_estimate.Rdata')
581
582   ## Initialize an empty dataframe to hold results
583   estimated_demand_df <- data.frame(time_period = c(),
584                                     RC = c(),
585                                     demand = c(),
586                                     engine = c())
587
588   # Loop through the RCs, then estimate the probabilities using the Rust
589   method, then do simulation.
590   for (j in RC_range) {
591     RC = j
592
593     ### Set a critical value for to measure the deviation between iterative
594     updates of EV. The distance between the two EV matrices
595     ### is the infinity norm of the difference
596     cri <- 10^(-8)

```

```

592
593 ##### Set an initial value for the EV matrix (all 0s, EV), and another EV
      object to hold the updated estimates, EV2.
594 EV <- matrix(100,33,2)
595 EV2 <- matrix(0,33,2)
596
597 ## While the infinity norm is less than the threshold, iterate
598 while(max(abs(EV-EV2))>cri){
599
600     ##### Set the current EV to the previous updated EV
601     EV <- EV2
602     ##### Compute a new updated EV by iterating on the current EV
603     EV2 <- value.Iterate(EV)
604 }
605
606 ##### Do one last update to set EV equal to the last EV2
607 EV <- EV2
608
609 # get EV(x,i) for x=0,1,2,..,30
610 ##### EV contains extra states, which we needed to compute the above
      computation. Throw them away.
611 EV <- EV[1:31,]
612
613 ##### Get estimated probability based on the above EV
614 p_ix_hat <- choice.prob.Estimate()
615 ##### Estimate demand using that probability
616 estimated_demand <- estimate_demand(n_sims, n_buses, n_periods)
617
618 ##### Add this estimate to a temp dataframe
619 estimated_demand_df_temp <- data.frame(time_period = seq(1, n_periods, 1)
      ,
620                                         RC = rep(RC, n_periods),
621                                         demand = estimated_demand,
622                                         engine = rep('Engine_1', n_periods))
623 ##### Collate temp dataframe to full dataframe
624 estimated_demand_df <- rbind(estimated_demand_df, estimated_demand_df_
      temp)
625
626 }
627
628 ##### Reset theta_1 to the "new engine", redo the exercise above.
629 theta_1 = .02
630
631 ##### Loop through RCs
632 for (j in RC_range) {
633     RC = j
634
635     ##### Set a critical value for to measure the deviation between iterative
      updates of EV. The distance between the two EV matrices
636     ##### is the infinity norm of the difference
637     cri <- 10^(-8)
638
639     ##### Set an initial value for the EV matrix (all 0s, EV), and another EV
      object to hold the updated estimates, EV2.

```

```

640 EV <- matrix(100,33,2)
641 EV2 <- matrix(0,33,2)
642
643 ## While the infinity norm is less than the threshold, iterate
644 while(max(abs(EV-EV2))>cri){
645
646     ### Set the current EV to the previous updated EV
647     EV <- EV2
648     ### Compute a new updated EV by iterating on the current EV
649     EV2 <- value.Iterate(EV)
650 }
651
652 ### Do one last update to set EV equal to the last EV2
653 EV <- EV2
654
655 # get EV(x,i) for x=0,1,2,...,30
656 ### EV contains extra states, which we needed to compute the above
        computation. Throw them away.
657 EV <- EV[1:31,]
658
659 ### Get probability estimates based on EV
660 p_ix_hat <- choice.prob.Estimate()
661 ### Estimate demand
662 estimated_demand <- estimate_demand(n_sims, n_buses, n_periods)
663
664 ### Add to temp dataframe
665 estimated_demand_df_temp <- data.frame(time_period = seq(1, n_periods, 1)
666                                     ,
667                                     RC = rep(RC, n_periods),
668                                     demand = estimated_demand,
669                                     engine = rep('Engine_2', n_periods)
670                                     ))
671
672 ### Collate to full dataframe
673 estimated_demand_df <- rbind(estimated_demand_df, estimated_demand_df_
674                             temp)
675 }
676
677 ### For a reduced set of RCs, see the period-by-period demand
678 per_period_demand_plot <- estimated_demand_df %>%
679   filter(RC %in% c(1, 3, 5, 7, 10)) %>%
680   mutate(RC = as.factor(RC)) %>%
681   ggplot(., aes(x=time_period, y=demand, color=RC)) + geom_line() +
682   facet_wrap(~engine) + xlab('Period') + ylab('Demand_for_Engines') +
683   ggtitle('Demand_for_engines_over_time') +
684   theme(plot.title = element_text(hjust = 0.5))
685 ggsave(per_period_demand_plot, file='per_period_demand_plot.png', height=4,
686       width=6, units='in')
687
688 ### Aggregate over periods to get demand as a function of RC for different
        thetas.
689 aggregate_demand_plot <- estimated_demand_df %>%
690   group_by(RC, engine) %>%
691   summarise(total_demand = sum(demand)/n_periods) %>%

```



```

687 | ungroup() %>% head()
688 | ggplot(., aes(x=RC, y=total_demand, color=engine)) + geom_line() + xlab('
689 | RC') +
689 | ylab('Average_per-period_demand') + ggtitle('Average_per-period_engine_
690 | demand_for_100_buses') +
690 | ggsave(aggregate_demand_plot, file='aggregate_demand_plot.png', height=4,
691 |         width=6, units='in')
692 |
693 | #####
694 | # Question 3.5 #
695 | #####

```

Listing 1: ./rust.R