
14.273 Industrial Organization: Pset4

Dave Holtz, Jeremy Yang

May 18, 2017

1. Model setup.

Following the notations in Rust (1987), HZ's flow utility is:

$$u(x_t, i_t, \theta_1) + \epsilon_t(i_t) = \begin{cases} -RC - c(0, \theta_1) + \epsilon_t(1) & i_t = 1 \\ -c(x_t, \theta_1) + \epsilon_t(0) & i_t = 0 \end{cases}$$

where RC is the replacement cost, x_t is the observed state variable for mileage, $c(\cdot)$ is cost function and i_t is the decision to replace engine and $\epsilon_t(\cdot)$ is action specific and type I extreme value distributed structural error (or unobserved state variable).

The state transition probability is given by:

$$\theta_{3j} = \mathbb{P}(x_{t+1} = x_t + j | x_t, i_t = 0)$$

$j \in \{0, 1, 2\}$ and if $i_t = 1$ then $x_{t+1} = 0$ with probability 1.

HZ chooses i_t in every period t to maximize an infinite sum of discounted flow utilities. The maximal value is defined as the value function (suppress the dependency on θ_1, θ_3):

$$V(x_1, \epsilon_1) := \max_{i_t, t \in \{1, 2, \dots\}} \mathbb{E} \left[\sum_{t=1}^{\infty} \beta^{t-1} (u(x_t, i_t, \theta_1) + \epsilon_t(i_t)) \right]$$

Rewrite the value function as in the Bellman optimality form:

$$V(x_t, \epsilon_t) = \max_{i_t} (u(x_t, i_t, \theta_1) + \epsilon_t(i_t)) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t]$$

where the expectation is with respect to (conditional) state transition probability of both x and ϵ , see Rust (1987) equation (4.5). The Bellman equation breaks the dynamic optimization problem into an infinite series of static choices.

2. (1) The choice specific value function can be derived by plugging a specific action into the value function:

$$\tilde{V}(x_t, \epsilon_t, i_t) = \begin{cases} -RC - c(0, \theta_1) + \epsilon_t(1) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 1] \\ -c(x_t, \theta_1) + \epsilon_t(0) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 0] \end{cases}$$

$$V(x_t, \epsilon_t) = \max\{\tilde{V}(x_t, \epsilon_t, 1), \tilde{V}(x_t, \epsilon_t, 0)\}$$

HZ's decision is about trading off the total (future) cost of maintaining an old engine and the lump sum cost of replacing to a new one. The time to replace is the stopping time in this problem, so it can be thought as an optimal stopping time problem where the optimal policy is characterized by a cutoff in x , HZ would choose to replace the engine if x is above that threshold (the threshold depends on realized value of ϵ).

- (2) It's clear from 2 (1) that the optimal stopping rule is:

$$\begin{aligned} & -RC - c(0, \theta_1) + \epsilon_t(1) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 1] > \\ & -c(x_t, \theta_1) + \epsilon_t(0) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 0] \end{aligned}$$

or,

$$\tilde{V}(x_t, \epsilon_t, 1) > \tilde{V}(x_t, \epsilon_t, 0)$$

therefore, because the errors are type I extreme value distributed:

$$\mathbb{P}(i_t = 1 | x_t) = \frac{\exp(u(x_t, 1, \theta_1) + \beta \mathbb{E}[V_{t+1} | x_t, i_t = 1])}{\sum_{k=\{0,1\}} \exp(u(x_t, k, \theta_1) + \beta \mathbb{E}[V_{t+1} | x_t, i_t = k])} \quad (2.1)$$

where $u(x_t, i_t, \theta_1)$ is defined in 1 and for convenience:

$$V_{t+1} := V(x_{t+1}, \epsilon_{t+1})$$

- (3) For discrete x , under the assumption that the errors are type I extreme value distributed, we have (Rust (1987) equation (4.14)):

$$EV(x, i) = \sum_y \log\left\{ \sum_j \exp[u(y, j) + \beta EV(y, j)] \right\} \cdot p(y | x, i) \quad (2.2)$$

where

$$EV(x, i) := \mathbb{E}[V_{t+1} | x_t, i_t]$$

and x, i are the state and choice of current period and y, j are the state and choice of the next period. Also note that here the transition probability does not depend on x_t but only on j (or Δx). To compute expected value function, we first need to estimate transition probability from the data, this can be done simply by counting:

$$\hat{\theta}_{30} = \frac{\sum_b \sum_t 1_{\{x_{bt+1} - x_{bt} = 0, i_{bt} = 0\}}}{\sum_b \sum_t 1_{\{i_{bt} = 0\}}}$$

$$\hat{\theta}_{31} = \frac{\sum_b \sum_t 1_{\{x_{bt+1}-x_{bt}=1, i_{bt}=0\}}}{\sum_b \sum_t 1_{\{i_{bt}=0\}}}$$

$$\hat{\theta}_{32} = \frac{\sum_b \sum_t 1_{\{x_{bt+1}-x_{bt}=2, i_{bt}=0\}}}{\sum_b \sum_t 1_{\{i_{bt}=0\}}}$$

we compute the expected value function in the inner loop of the nested fixed point algorithm (holding the value of θ fixed), we first guess the initial values of $EV(x, i)$ for all possible values of x, i and use the equation (2.2) to iterate expected value function until it converges. The criterion is:

$$\max_{x,i} |EV^{T+1}(x, i) - EV^T(x, i)| < \eta$$

The plot for $x = 1 - 30$ at the true value of parameters are shown in Figure 1.

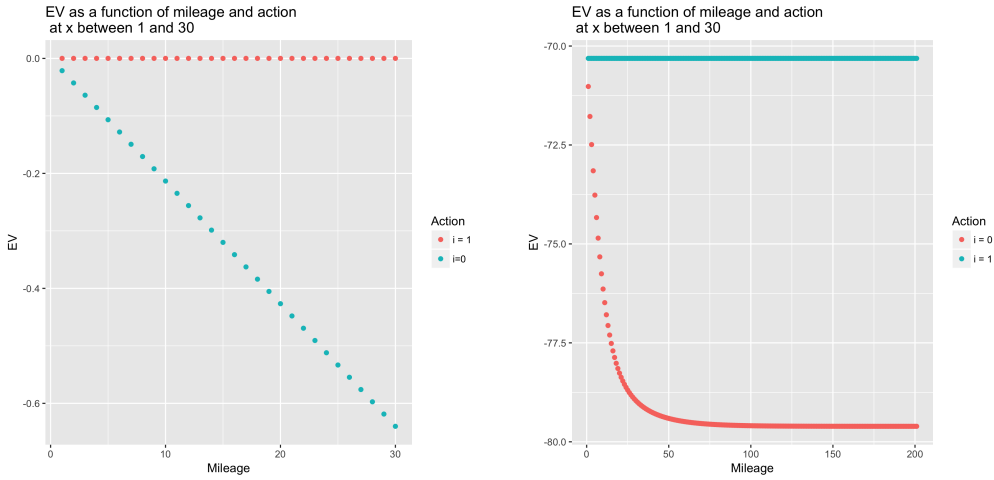


Figure 1: Expected Value Function for $i = 0$ and $i = 1$. Left panel shows results using iterative method, right panel shows provided Rust results.

Interestingly, our EV results are linear in mileage, which is probably not expected. Despite a good amount of debugging, we have been unable to identify a problem. However, it's also unclear how our calculated results for *just* 30 states should compare to the provided EV results, which provide information on 200 states. The first 30 states of the provided Rust EV estimates are decreasing in approximately linear fashion, suggesting our estimates might not be *so* bad. However, the order of magnitude of our EV values (e.g., 10^{-1}) is much smaller than the order of magnitude of EV values in the provided dataset (e.g., ~ 70), suggesting something is probably wrong. However, we don't have any more time to debug this, so we simply moved on.

- (4) The provided dataset contains mileage and engine replacement information for 100 buses over 1,000 periods. The table below shows the mean mileage, maximum mileage, minimum mileage, standard deviation of the mileage, the average mileage at engine replacement across all buses and periods, and the

average number of engine replacements for a particular bus over the 1,000 periods.

avg miles	max miles	min miles	s.d. miles	avg replace miles	avg replacements
8.245	33.000	0.000	5.709	15.953	52.980

We might also be interested in understanding how each of these summary statistics vary across buses. For instance, maybe some buses have their engines replaced much more often. In order to study this, Figure 2 shows the distributions of average mileage, maximum mileage, s.d. mileage, avg miles at replacement, and number of replacements across the 100 buses in the sample. In general, these distributions are quite concentrated, suggesting that there are not systematic differences across buses.

The final, bottom right plot in 2 also shows the empirically observed conditional choice probability as a function of state (mileage) that Harold Zurcher actually acts on. At a high level, Zurcher's has to make the investment decision of when to replace a given bus's engine. The mean replacement mileage plot suggests that on average he replaces a bus's engine after about 80,000 miles. The conditional choice probability plot suggests that the likelihood he increases the engine is practically zero until the bus hits 50,000 miles, after which the probability that the bus has its engine replaced climbs quickly. By the time a bus has 150,000 miles on it, it has a 50% probability of having its engine changed in a given time period.

3. (1) In the outer loop we search over a grid of values for (θ_1, β, RC) , and compute the log likelihood function:

$$\log L = \sum_b \left\{ \sum_t \log \mathbb{P}(i_{bt} | x_{bt}) + \sum_t \log \mathbb{P}(x_{bt} | x_{bt-1}, i_{t-1}) \right\}$$

where b indexes for bus and t indexes for time period. We compute a log likelihood for each combination of values for (θ_1, β, RC) and choose the set of parameters that maximizes the log-likelihood of the data. The maximum likelihood parameters obtained with the Rust method are:

$$\begin{aligned} \theta_1 &= 0.1 \\ \beta &= 0.99 \\ RC &= 6 \end{aligned}$$

- (2) In Hotz-Miller's approach, we will estimate the choice specific value function (as opposed to the expected value function as in Rust). We start by noting that conditional choice probability is observed directly from the data:

$$\hat{\mathbb{P}}(i = 1 | x) = \frac{\sum_b \sum_t 1_{\{i_{bt}=1, x_{bt}=x\}}}{\sum_b \sum_t 1_{\{x_{bt}=x\}}}$$

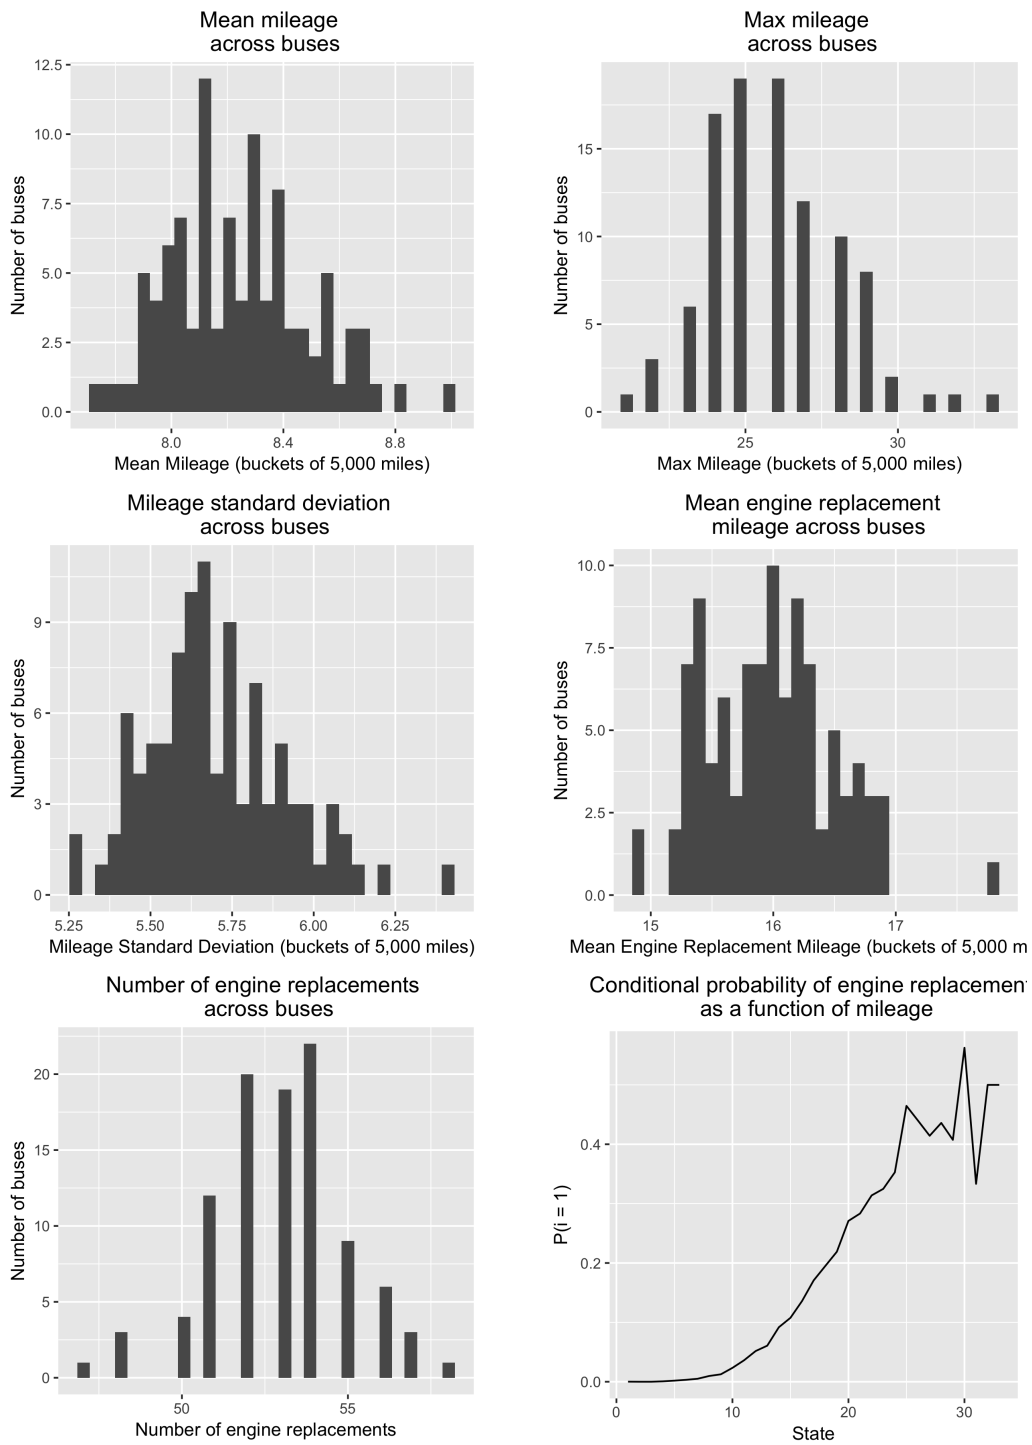


Figure 2: Distribution of various summary statistics across buses, as well as the empirical conditional choice probability for Zurcher.

The choice-specific value function (minus the structural error, and suppressing the dependency on θ_1, θ_3) can be presented recursively in the following form:

$$\tilde{V}(x_t, i_t) = u(x_t, i_t) + \beta \mathbb{E}_{x_{t+1}} [\mathbb{E}_{i_{t+1}} [\mathbb{E}_{\epsilon_{t+1}} [u(x_{t+1}, i_{t+1}) + \epsilon_{t+1} + \beta(\dots) | i_{t+1}, x_{t+1}] | x_{t+1}] | x_t, i_t]$$

where (\dots) represents higher (two and above) period forward expectations. In principle it's an infinite loop but in practice we need to stop at some T , for example, when $T = 2$, (\dots) simplifies to:

$$(\dots) = \mathbb{E}_{x_{t+2}} [\mathbb{E}_{i_{t+2}} [\mathbb{E}_{\epsilon_{t+2}} [u(x_{t+2}, i_{t+2}) + \epsilon_{t+2} | i_{t+2}, x_{t+2}] | x_{t+2}] | x_{t+1}, i_{t+1}]$$

For simplicity, in the code we use one-period forward simulation where:

$$\tilde{V}(x_t, i_t) = u(x_t, i_t) + \beta \mathbb{E}_{x_{t+1}} [\mathbb{E}_{i_{t+1}} [\mathbb{E}_{\epsilon_{t+1}} [u(x_{t+1}, i_{t+1}) + \epsilon_{t+1} | i_{t+1}, x_{t+1}] | x_{t+1}] | x_t, i_t]$$

it is estimated as:

$$\hat{\tilde{V}}(x_t, i_t) = \frac{1}{S} \sum_s [u(x_t, i_t) + \beta [u(x_{t+1}^s, i_{t+1}^s) + \gamma - \log(\hat{\mathbb{P}}(i_{t+1}^s | x_{t+1}^s))]]$$

where x_{t+1}^s is drawn from the transition probability $\hat{\theta}_{30}, \hat{\theta}_{31}, \hat{\theta}_{32}$, and i_{t+1}^s is drawn from $\hat{\mathbb{P}}(i|x)$, γ is the Euler's constant. We only go one period forward because we only observe data for states up to $x_t = 33$. It is possible for larger T that we would encounter a state that is not in our dataset. When this occurs, it's unclear what value should be used as the conditional choice probability. While we avoid this issue by setting $T = 2$, this does reduce the precision of our estimates.

- (3)
- (4) We want to compute HZ's demand function for the two buses, which we will denote as engine 1 ($\theta_1 = 0.09, RC = 6$) and engine 2 ($\theta_1 = 0.02, RC = 20$) as a function of RC. In order to do so, we obtain conditional choice probability estimates, $\hat{\mathbb{P}}(i = 1|x)$ by using the Rust method to iterate EV values. We use the Rust methodology because the Hotz and Miller methodology depends on the observed conditional choice probabilities, which we know do *not* correspond to the counterfactual engine 2.

With those conditional choice probability estimates for the two engines in hand, we run 1,000 simulations of a bus's state transitions (and HZ's corresponding engine replacement decisions) over the first 15 periods. This allows us to get an expected, per-bus demand for engines over the first 15 periods. In order to get the expected demand that HZ has for engines across all buses, we simply multiply this figure by 100. So the expected demand (as a function of period t) is:

$$D(t) = 100 \times \sum_{x_t} \hat{\mathbb{P}}(i = 1 | x = x_t) \hat{\mathbb{P}}(x = x_t | t) \quad (1)$$

HZ's demand for engines as a function of the period, t for a few values of RC can be found in Figure 3. The average per-period demand for the two engines (averaged across 15 periods) for different values of RC can be found in Figure 4. it's worth noting that the demand curves for the two engines appear almost identical - although you cannot distinguish them in the figure, there are small differences (on the order of a tenths of an engine. This could be a true difference, or it could be simulation error. Although we're not sure why these demand curves are so similar, we have two hypotheses:

1. Whatever our EV estimation issue, it is rearing its ugly head again and making these demand curves very similar.
2. The increase in RC from engine 1 to engine 2 is almost perfectly offset by the decrease in θ_1 , creating two extremely similar demand curves.

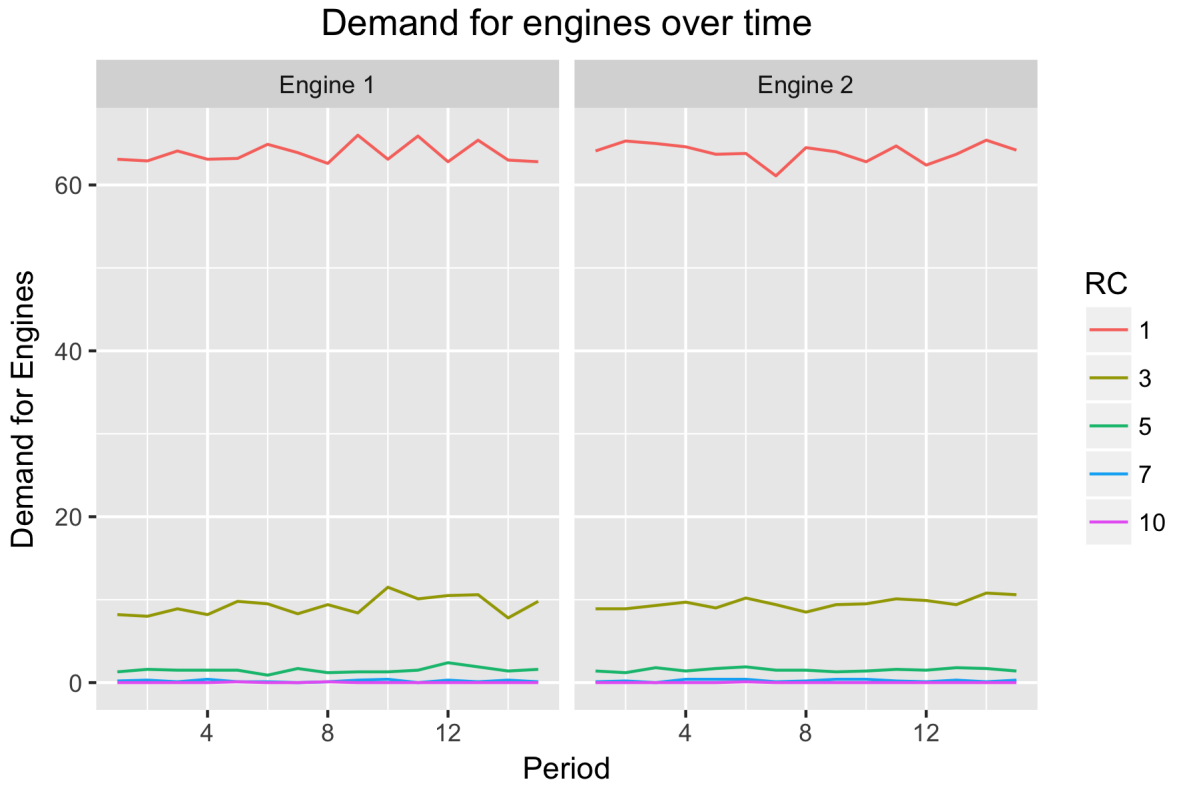


Figure 3: The demand for engines across a fleet of 100 buses as a function of period (over the first 15 periods) for different values of RC. Unsurprisingly, when RC is lower, HZ is much more willing to change bus engines.

(5)

APPENDIX: CODE




Figure 4: Aggregate per-period demand for new engines across a fleet of 100 buses as a function of RC. For engine 1, $\theta_1 = 0.09$. For engine 2, $\theta_1 = 0.02$.

```

1 ##### This code uses the methodologies of both Rust (1987) and Hotz and
   Miller (1993) to estimate the parameters of a single
2 ##### agent dynamic problem where an agent (Harold Zurcher) must choose when
   to have the engines replaced in a fleet of buses.
3
4 ## Import libraries
5 library(R.matlab)
6 library(ggplot2)
7 library(dplyr)
8
9 ## Read in data
10 setwd('~/.Dropbox/MIT/MIT/Spring_2017/14.273/HW4/273-pset4/')
11 data <- readMat('rust.mat')
12 gamma = .577
13
14 ## Extract bus replacement events
15 i <- data$it
16 ## Extract bus mileage counts (in increments of 5,000 miles)
17 x <- data$xt
18
19

```



```

20 ## Buses transition from different mileage states , and can jump forward
    zero , one , or two 5,000 mile buckets. This block
21 ## of code estimates the transition probabilities empirically from the data
    .
22
23 ##### Initialize empty vectors to hold a count of how many times each jump
    happens
24 zero <- vector()
25 one <- vector()
26 two <- vector()
27
28 ##### Loop through the 100 buses in the dataset
29 for (k in 1:100) {
30
31     ##### Given a bus k, grab the mileage counts
32     xk <- x[,k]
33     ##### Also grab the engine replacement events
34     ik <- i[,k]
35     ##### Get a modified array which gives the change in mileage buckets
        from period j to period j+1
36     jk <- xk[-1]-xk[-1000]
37
38     ##### We only care about periods where i=0 for transition
        probabilities , since i=1 will always send
39     ##### x back to 0. This selects out only time periods for this bus
        where i = 0
40     j <- jk[ik==0]
41
42     ##### This counts up how many times the mileage bucket counter , x,
        moves up by 0, 1, or 2 when i=0
43     zero[k] <- length(j[j==0])
44     one[k] <- length(j[j==1])
45     two[k] <- length(j[j==2])
46 }
47
48 ## Estimate the x_t-independent transition probabilities by dividing the
    number of times for each transition by the
49 ## total number of transitions
50 theta_30 = sum(zero)/(sum(zero)+sum(one)+sum(two))
51 theta_31 = sum(one)/(sum(zero)+sum(one)+sum(two))
52 theta_32 = sum(two)/(sum(zero)+sum(one)+sum(two))
53
54 #####
55 # Question 2.3 #
56 #####
57
58 ##### We'll now take the true values of the parameter values as given , and
    use the method described in Rust (1987) to iteratively
59 ##### estimate the value function (or in this case , the EV function).
60
61 ## Initialize parameters to their true values
62 theta_1 = .05
63 theta_30 = .3
64 theta_31 = .5

```

```

65 | theta_32 = .2
66 | beta =.99
67 | RC = 10
68 |
69 | ### Define the linear cost function. If an engine is not replaced, the bus
    | incurs cost theta_1*x, so cost
70 | ### increases linearly as a bus gets older.
71 | cost <- function(x){
72 |     return (theta_1*x)}
73 |
74 | ### Define the utility function at mileage x from action i. If the agent
    | chooses to replace the engine in a bus,
75 | ### it costs RC. If they choose _not_ to replace the engine, they incur the
    | cost of running the bus at mileage x.
76 | u <- function(x,i){
77 |     -RC*i - cost(x*(1-i))
78 | }
79 |
80 |
81 | ### The value function can be estimated through an iteration procedure. We
    | start with some initial guess for EV,
82 | ### calculate EV with an expression that includes our initial guess of EV,
    | and continue iterating until the difference
83 | ### between subsequent EV estimates becomes small.
84 |
85 | ##### value.Iterate is a function to iteratively update the value function
    | according to the methodology in Rust. The function
86 | ##### takes as an argument a current estimate of EV, and returns an updated
    | estimate of EV. EV is an x by d matrix – we want the
87 | ##### EV values for each decision d at every possible current mileage value
    | x.
88 | value.Iterate <- function(EV){
89 |
90 |     ### First iterate through each of the 30 x states
91 |     for (x in 1:31){
92 |         ## Update the EV value corresponding to not replacing the engine. There
            | are three contributions here – one from the
93 |         ## j = 0 case, one from the j = 1 case, and one from the j = 2 case.
            | Note the indexing here. When x =1, the state is
94 |         ## equal to 0 (this is the x that needs to be passed into u()), but we
            | want to grab the EV corresponding to the 1st entry.
95 |         EV2[x,1] <- log(exp(u(x-1,0)+beta*EV[x,1])+exp(u(x-1,1)+beta*EV[x,2]))*
            | theta_30 + gamma
96 |         + log(exp(u(x,0)+beta*EV[x+1,1])+exp(u(x,1)+beta*EV[x+1,2]))*theta_31
            | + gamma
97 |         + log(exp(u(x+1,0)+beta*EV[x+2,1])+exp(u(x+1,1)+beta*EV[x+2,2]))*theta_
            | 32 + gamma
98 |
99 |         ## Update the EV value corresponding to replacing the engine. When the
            | engine is replaced, x at the next period will
100 |         ## deterministically reset to x = 0.
101 |         EV2[x,2] <- log(exp(u(0,0)+beta*EV[1,1])+exp(u(0,1)+beta*EV[1,2])) +
            | gamma
102 |     }

```

```

103
104   ## Return the updated EV values.
105   return(EV2)
106 }
107
108 ##### Set a critical value for to measure the deviation between iterative
109         updates of EV. The distance between the two EV matrices
110 ##### is the infinity norm of the difference
111 cri <- 10^(-8)
112
113 ##### Set an initial value for the EV matrix (all 0s, EV), and another EV
114         object to hold the updated estimates, EV2.
115 EV <- matrix(100,33,2)
116 EV2 <- matrix(0,33,2)
117
118 ## While the infinity norm is less than the threshold, iterate
119 while(max(abs(EV-EV2))>cri){
120
121     ##### Set the current EV to the previous updated EV
122     EV <- EV2
123     ##### Compute a new updated EV by iterating on the current EV
124     EV2 <- value.Iterate(EV)
125 }
126
127 ##### Do one last update to set EV equal to the last EV2
128 EV <- EV2
129
130 # get EV(x,i) for x=0,1,2,...,30
131 ##### EV contains extra states, which we needed to compute the above
132         computation. Throw them away.
133 EV <- EV[1:31,]
134
135 ##### Plot the EV of both replacing the engine (i = 1) and not replacing the
136         engine (i = 0) at every x
137 ##### between 1 and 30
138 df <- data.frame('x'=c(1:30, 1:30), 'EV'=c(EV[2:31,1], EV[2:31,2]), 'Action'
139         = c(rep('i_0', 30), rep('i_1', 30)))
140
141 ##### Generate a plot that compares the EV of replacing the engine and not
142         replacing the engine
143 ev_plot <- ggplot(df, aes(x=x, y=EV, color=Action)) + geom_point() + xlab('
144         Mileage') + ylab('EV') +
145         ggtitle('EV as a function of mileage and action\n at x between 1 and 30'
146         ) +
147         theme(plot.title = element_text(hjust = 0.5))
148 ggsave(ev_plot, file='ev_plot.png', height=6, width=6, units='in')
149
150 ##### This is a plot to see the EV data in the attached rust matlab file. The
151         state space is different than ours (200 states),
152 ##### so its hard to compare. Our's is linear (seems wrong), whereas the
153         provided data is not. However, the first 30 states
154 ##### _do_ look approximately linear, so maybe we're not so far off.
155

```

```

146 df_rust <- data.frame('x'=c(seq(1,201), seq(1, 201)), 'EV'=c(data$EV[,1],
147   data$EV[,2]), 'Action' = c(rep('i_0', 201),

148 ev_plot_rust <- ggplot(df_rust, aes(x=x, y=EV, color=Action)) + geom_point
  () + xlab('Mileage') + ylab('EV') +
149   ggtitle('Rust_dataset_EV_as_a_function_of_mileage_and_action\nat_x_
    between_1_and_201') +
150   theme(plot.title = element_text(hjust = 0.5))
151 ggsave(ev_plot_rust, file='ev_plot_rust.png', height=6, width=6, units='in'
  )

152 #####
153 # Question 2.4 #
154 #####
155
156
157 ### Calculate the mean mileage, mean time to engine replacement, max
  mileage, min mileage, and sd mileage over the whole sample
158 mean_x <- mean(x)
159 mean_engine_replacement_age <- mean(x[i == 1])
160 max_x <- max(x)
161 min_x <- min(x)
162 sd_x <- sd(x)
163 avg_replacements <- mean(apply(i, 2, function(x) {sum(x)}))
164
165 aggregate_stats <- c(mean_x, max_x, min_x, sd_x, mean_engine_)
166 aggregate_stats %>%
167   round(., 3) %>%
168   kable(., format='latex')
169
170 ### Calculate the per bus mean mileage, mean time to engine replacement,
  max mileage, min mileage, and sd mileage
171 mean_x_per_bus <- apply(x, 2, function(x) {mean(x)})
172 max_x_per_bus <- apply(x, 2, function(x) {max(x)})
173 min_x_per_bus <- apply(x, 2, function(x) {min(x)})
174 sd_x_per_bus <- apply(x, 2, function(x) {sd(x)})
175 mean_engine_replacement_per_bus <- apply(x*i, 2, function(x) {sum(x)/sum(x
  != 0)})
176 replacements <- apply(i, 2, function(x) {sum(x)})
177
178

```

```

179 ### Collate per bus information into a dataframe
180 per_bus_statistics <- data.frame(bus = seq(1, 100, 1),
181                                   mean_x_per_bus = mean_x_per_bus,
182                                   max_x_per_bus = max_x_per_bus,
183                                   min_x_per_bus = min_x_per_bus,
184                                   sd_x_per_bus = sd_x_per_bus,
185                                   mean_engine_replacement_per_bus = mean_
186                                     engine_replacement_per_bus,
187                                   replacements = replacements)
188
189 ### Create some plots
190 mean_mileage_plot <- ggplot(per_bus_statistics, aes(x=mean_x_per_bus)) +
191   geom_histogram() +
192   xlab('Mean_Mileage_(buckets_of_5,000_miles)') + ylab('Number_of_buses') +
193   ggtitle('Mean_mileage_\n_across_buses') +
194   theme(plot.title = element_text(hjust = 0.5))
195 ggsave(mean_mileage_plot, file='mean_mileage_plot.png', height=4, width=4,
196        units='in')
197
198 max_mileage_plot <- ggplot(per_bus_statistics, aes(x=max_x_per_bus)) + geom_
199   _histogram() +
200   xlab('Max_Mileage_(buckets_of_5,000_miles)') + ylab('Number_of_buses') +
201   ggtitle('Max_mileage_\n_across_buses') +
202   theme(plot.title = element_text(hjust = 0.5))
203 ggsave(max_mileage_plot, file='max_mileage_plot.png', height=4, width=4,
204        units='in')
205
206 sd_mileage_plot <- ggplot(per_bus_statistics, aes(x=sd_x_per_bus)) + geom_
207   _histogram() +
208   xlab(' Mileage_Standard_Deviation_(buckets_of_5,000_miles)') + ylab('
209     Number_of_buses') +
210   ggtitle(' Mileage_standard_deviation_\n_across_buses') +
211   theme(plot.title = element_text(hjust = 0.5))
212 ggsave(sd_mileage_plot, file='sd_mileage_plot.png', height=4, width=4,
213        units='in')
214
215 time_to_engine_replacement_plot <- ggplot(per_bus_statistics, aes(x=mean_
216   engine_replacement_per_bus)) + geom_histogram() +
217   xlab('Mean_Engine_Replacement_Mileage_(buckets_of_5,000_miles)') + ylab('
218     Number_of_buses') +
219   ggtitle('Mean_engine_replacement_\n_mileage_across_buses') +
220   theme(plot.title = element_text(hjust = 0.5))
221 ggsave(time_to_engine_replacement_plot, file='time_to_engine_replacement_
222   plot.png', height=4, width=4, units='in')
223
224 replacements_plot <- ggplot(per_bus_statistics, aes(x=replacements)) + geom_
225   _histogram() +
226   xlab('Number_of_engine_replacements') + ylab('Number_of_buses') +
227   ggtitle('Number_of_engine_replacements_\n_across_buses') +
228   theme(plot.title = element_text(hjust = 0.5))
229 ggsave(replacements_plot, file='replacements_plot.png', height=4, width=4,
230        units='in')

```

```

218 #####
219 # Question 3.1 #
220 #####
221
222 ### This code will estimate the parameters beta, theta_1, and RC using the
223 nested fixed-point algorithm
224 described in Rust.
225
226 ### A function to compute the probability of Zurcher's choices using the
227 EVs we calculate using the EV calculation
228 framework above. The probability of choosing different actions
229 basically acts like a multichoice logit function.
230 choice.prob.Estimate <- function(){
231
232     ### Initialize an empty matrix to hold choice probability estimates.
233     p_i <- matrix(0,31,2)
234
235     ### Iterate through all of the possible mileage states, x.
236     for (x in 1:31){
237         ## For each mileage state x, calculate the probability that Zurcher
238         will choose i = 0.
239         p_i[x,1] <- exp(u(x-1,0)+beta*EV[x,1])/(exp(u(x-1,0)+beta*EV[x,1])+exp(
240             u(x-1,1)+beta*EV[x,2]))
241         ## Calculate the probability of choosing i = 1 at each state, which is
242         just 1 - P(i = 0).
243         p_i[x,2] <- 1- p_i[x,1]
244     }
245
246     ### Return the updated p_i object
247     return(p_i)
248 }
249
250 ### A function to calculate the total log likelihood of the observed data
251 given a set of parameters. This method assumes that
252 the probabilities across periods and buses are independent, so we can
253 just add up all of the log probabilities.
254 log.likelihood.Compute <- function() {
255
256     ### Initialize 0-valued variables to hold the log choice probability, the
257     log transition probability,
258     and the sum of the two.
259     log_choice_prob <- 0
260     log_transition_prob <- 0
261     total <- 0
262
263     ### Iterate over buses
264     for (bus in 1:100) {
265         ### Iterate over time periods
266         for (t in 1:999) {
267             ### We special case mileage states greater than 30, since they are a
268             bit strange in our data. Otherwise, we calculate
269             the choice probability using the current value of p_i according
270             to the EV values we calculated to get the

```

```

260     ### choice probability. Take the log and add it to the current
      running value.
261     if (x[t,bus] <= 30){
262         log_choice_prob <- log(p_i[x[t,bus]+1,i[t,bus]+1]) + log_choice_
      prob
263     ### Do the same thing for our special cased, x > 30 case.
264     } else {
265         log_choice_prob <- log(p_i[31,i[t,bus]+1]) + log_choice_prob
266     }
267
268     ### Calculate over the transitions for each bus the sum of the log
      transition probabilities. We have our estimates of
269     ### theta_3 given the empirical transition probabilities. So we can
      just grab that for each observed transition and add it
270     ### to the total log transition probability.
271
272     ### First we do the j = 0 case.
273     if (x[t+1,bus]-x[t,bus]==0) {
274         log_transition_prob <- log(theta_30)+log_transition_prob
275     ### Then the j = 1 case.
276     } else if (x[t+1,bus]-x[t,bus]==1) {
277         log_transition_prob <- log(theta_31)+log_transition_prob
278     ### And finally the j = 2 case.
279     } else if (x[t+1,bus]-x[t,bus]==2) {
280         log_transition_prob <- log(theta_32)+log_transition_prob
281     }
282 }
283
284     ### Now, get the total log likelihood by adding up all of the
      transition components and the choice components.
285     total <- (log_choice_prob+log_transition_prob) + total
286 }
287 return (total)
288 }
289
290 ### Now we're actually going to use the nested fixed point algorithm to get
      the maximum likelihood estimates of the parameters
291 ### that we care about. This process has three steps.
292
293 ### Step 1: We would calculate theta_30, theta_31, and theta_31 directly
      from the data. This step is not in the loop, and we've
294 ### actually already done this and it doesn't change, so we don't need to
      do it again.
295
296 ### Step 2: Next, we are going to set up a grid over values of theta_1,
      beta, and RC that we will calculate the
297 ### log likelihood to determine the maximum likelihood parameter values. We
      'll also initialize a dataframe
298 ### to hold the parameter values and the log likelihoods.
299
300 theta_1_range <- seq(.01,.10,.01)
301 beta_range <- seq(.90,.99,.01)
302 RC_range <- seq(6,15,1)

```

```

303 | likelihood <- data.frame( 'theta_1'=rep(0) , 'beta'=rep(0) , 'RC'=rep(0) , 'log.
    |   likelihood'=rep(0) )
304 |
305 | ### Step 3: Now we actually do the nested fixed point computation.
306 |
307 | #### Loop through theta_1
308 | for (theta_1 in theta_1_range) {
309 |   ### Loop through beta
310 |   for (beta in beta_range) {
311 |     ### Loop through RC
312 |     for (RC in RC_range) {
313 |       print(paste(c(theta_1, beta, RC), collapse='_'))
314 |
315 |       ### Initialize the EV functions to the initial values we used above.
316 |       EV <- matrix(100,33,2)
317 |       EV2 <- matrix(0,33,2)
318 |
319 |       ### Iteratively compute the EV values.
320 |       while(max(abs(EV-EV2))>cri){
321 |         EV <- EV2
322 |         EV2 <- value.Iterate(EV)
323 |       }
324 |
325 |       EV <- EV2
326 |       EV <- EV[1:31,]
327 |
328 |       ### Given these values of EV, calculated the choice probabilities
329 |       p_i <- choice.prob.Estimate()
330 |
331 |       ### Given the EV values, the choice probabilities and the parameters,
    |         calculate
332 |       ### the log-likelihood of the data.
333 |       likelihood <- rbind(likelihood ,c(theta_1,beta,RC,log.likelihood.
    |         Compute()))
334 |     }
335 |   }
336 | }
337 |
338 | ### Retrieve the row in the likelihood dataframe corresponding to the
    |   maximum likelihood estimate
339 | likelihood <- likelihood[-1,]
340 | parameter_estimates <- likelihood[which.max(likelihood[,4]) ,]
341 |
342 | ### Use these parameters and get the relevant estimate of EV and p_i
343 | theta_1 = parameter_estimates$theta_1
344 | beta = parameter_estimates$beta
345 | RC = parameter_estimates$RC
346 |
347 | EV <- matrix(100,33,2)
348 | EV2 <- matrix(0,33,2)
349 | ### Iteratively compute the EV values.
350 | while(max(abs(EV-EV2))>cri){
351 |   EV <- EV2
352 |   EV2 <- value.Iterate(EV)

```



```

353 }
354 EV <- EV2
355 EV <- EV[1:31,]
356 ### Given these values of EV, calculate the choice probabilities
357 p_i <- choice.prob.Estimate()
358 ### Given the EV values, the choice probabilities and the parameters,
    calculate
359 ### the log-likelihood of the data.
360 likelihood <- rbind(likelihood,c(theta_1,beta,RC,log.likelihood.Compute()))
361
362 save(EV, p_i, likelihood, parameter_estimates, file='rust_estimate.Rdata')
363
364 #####
365 # Question 3.2 #
366 #####
367
368 ### Now we will get estimates of the parameters using the Hotz and Miller
    conditional choice probability approach. This will
369 ### allow us to compare these parameter estimates to those obtained using
    the Rust approach.
370
371 ### First, we need to calculate the probability of the agent choosing
    either i = 0 or i = 1 based on the state that they find
372 ### a given bus in, x, at some time period t. This will be the baseline
    that we use to try and find the best parameter values
373 ### (i.e., which parameter values minimize the infinity norm between these
    true probabilities and the estimated probabilities)
374
375 ### The probability matrix
376 p_ix <- matrix(0,33,2)
377 ### The vector of how often the agent chooses i=1 given state x
378 ones <- vector()
379 ### The vector of how often the agent finds a bus in state x
380 total <- vector()
381
382 ### Loop through the states
383 for (state in 0:32){
384
385     ### For a given state, a will track how many times i = 1 and b will track
        how many times that state occurs.
386     ### Initialize them to 0 for the given state.
387     a <- 0
388     b <- 0
389
390     ### Loop over the buses
391     for (bus in 1:100){
392
393         ### Increment how many times the agent chooses i = 1 in state x
394         a <- sum(i[which(x[,bus]==state),bus]) + a
395         ### Increment how many times the state x occurs
396         b <- length(i[which(x[,bus]==state),bus]) + b
397     }
398
399     ### Add the most recent estimates to the vector.

```

```

400     ones[state+1] <- a
401     total[state+1] <- b
402 }
403
404 ### Based on the ones and total vectors, updated the choice probability
405     matrix.
406 p_ix[,1] <- 1-ones/total
407 p_ix[,2] <- ones/total
408
409 ### Plot conditional choice probabilities
410 p_ix_df <- as.data.frame(p_ix)
411 p_ix_df$state <- as.numeric(rownames(p_ix_df))
412 names(p_ix_df) <- c('P(i=0)', 'P(i=1)', 'State')
413 ccp_plot <- p_ix_df %>%
414     ggplot(., aes(x=State, y='P(i = 1)')) + geom_line() +
415     ggtitle('Conditional_probability_of_engine_replacement_\n_as_a_function_
416         of_mileage') +
417     theme(plot.title = element_text(hjust = 0.5))
418 ggsave(ccp_plot, file='ccp_plot.png', height=4, width=4, units='in')
419
420 ### The function below uses the Hotz and Miller method to estimate V and p_
421     ix_hat for every state and period
422 ### given a set of model parameters (beta, theta_1, and RC).
423 approximate.V_pihat <- function() {
424     ### Initialize an empty valuation matrix
425     V <- matrix(0,33,2)
426     ### Initialize an empty conditional choice probability matrix
427     p_ix_hat <- matrix(0,33,2)
428
429     ### Iterate through the states
430     for (state in 0:30){
431         ### Initialize a and b, which will basically track a running total of V
432         for different choices over simulations, to 0.
433         a = 0
434         b = 0
435         ### Iterate through the simulations. Note that ideal we would probably
436         want to go more than one time step into the
437         ### future. However, because of the limitations in our dataset, we only
438         go one time step forward. This is mainly because
439         ### it's unclear how we would draw i (the choice) for states that do
440         not appear in our data (i.e., x = 34).
441         for (s in 1:S){
442             ## Conditional on choosing i = 0, simulate the next state that a
443             given bus will end up in by drawing from the
444             ## transition probabilities.
445             x_prime_0 = state + sample(c(0,1,2),1,replace = T, prob = c(theta_30,
446                 theta_31,theta_32))
447             ## Conditional on choosing i = 0 and ending up in some state in the
448             next time period, randomly simulate a draw from
449             ## i based on the conditional choice probabilities
450             i_prime_0 = sample(c(0,1),1,replace=T,prob = c(p_ix[x_prime_0+1,1],p_
451                 ix[x_prime_0+1,2]))
452             # Figure out the expected utility from this truncated sequence of
453             choices.

```

```

442     a = (u(state,0) + beta*(u(x_prime_0,i_prime_0)+gamma-log(p_ix[x_prime_0+1,i_prime_0+1]))) + a
443
444     ## Conditional on choosing i = 1, we don't need to simulate the next
445     ## state that a bus will end up in. It will always
446     ## be x = 0. So we jump right to simulating the draw from i for x =
447     ## 0.
448     i_prime_1 = sample(c(0,1),1,replace=T,prob = c(p_ix[1,1],p_ix[1,2]))
449     ## Figure out the expected utility from this truncated sequence of
450     ## choices.
451     b = (u(state,1) + beta*(u(0,i_prime_1)+gamma-log(p_ix[1,i_prime_1+1]))
452         ) + b
453   }
454
455   ## Set the value of V to be the average over all S of our simulations
456   ## for both the i = 0 and i = 1 choices.
457   V[state+1,1] = a/S
458   V[state+1,2] = b/S
459   ## Use the multinomial logit-esque probability expression to figure out
460   ## the probability of choosing i = 0 or i = 1
461   ## given that the bus is in state x.
462   p_ix_hat[state+1,1] <- exp(V[state+1,1])/(exp(V[state+1,1])+exp(V[state+1,2]))
463   p_ix_hat[state+1,2] <- 1- p_ix_hat[state+1,1]
464 }
465
466 # Put final output into a list and return it
467 results <- list('V' = V, 'p_ix_hat' = p_ix_hat)
468 return(results)
469 }
470
471 #### Specify a number of constants that will be used in the Hotz and Miller
472 algorithm:
473 #### S: The number of "simulations" to do per state / decision
474 #### gamma: This should be Euler's constant
475 #### theta_1_range: The range of theta_1 values to test
476 #### beta_range: The range of beta values to test
477 #### RC_range: The range of RC values to test
478 S = 1000
479 theta_1_range <- seq(.01,.10,.01)
480 beta_range <- seq(.90,.99,.01)
481 RC_range <- seq(6,15,1)
482
483 #### Initialize a dataframe to hold different parameter combinations and the
484 infinity-norm between the actual conditional
485 choice probabilities and the estimated ones
486 difference <- data.frame('theta_1'=rep(0),'beta'=rep(0),'RC'=rep(0),'
487   difference'=rep(0))
488
489 #### Loop through theta_1
490 for (theta_1 in theta_1_range) {
491   #### Loop through theta_2
492   for (beta in beta_range) {
493     #### Loop through RC

```

```

485   for (RC in RC_range) {
486     # Check progress
487     print(paste(c(theta_1, beta, RC), collapse='_'))
488
489     ### Get estimates of V and P_ix_hat using the Hotz and Miller method
490     v_and_p_ix_hat <- approximate.V_pixhat()
491     V = v_and_p_ix_hat$V
492     p_ix_hat <- v_and_p_ix_hat$p_ix_hat
493
494     ### Now that we have a full conditional choice probability matrix,
495     calculate the infinity norm (i.e., largest
496     ### absolute difference between the empirical conditional choice
497     probabilities and those estimated with the
498     ### given parameters)
499     difference <- rbind(difference, c(theta_1, beta, RC, max(abs(p_ix[1:31,] -
500       p_ix_hat[1:31,])))
501   }
502 }
503
504 ### Find the set of parameters that minimizes this difference
505 difference <- difference[-1,]
506 parameter_estimates <- difference[which.min(difference[,4]),]
507
508 ### Use these parameters and get the relevant estimate of V and p_ix_hat
509 theta_1 = parameter_estimates$theta_1
510 beta = parameter_estimates$beta
511 RC = parameter_estimates$RC
512 best_guesses <- approximate.V_pixhat()
513 V <- best_guesses$V
514 p_ix_hat <- best_guesses$p_ix_hat
515
516 save(V, p_ix_hat, difference, parameter_estimates, file='hotz_and_miller_
517   estimate.Rdata')
518
519 #####
520 # Question 3.3 #
521 #####
522
523 #####
524 # Question 3.4 #
525 #####
526
527 ### This function simulates, for one agent, a sequence of state transitions
528 and also engine replacement decisions
529 simulate_sequence <- function(n_periods) {
530   ### Initialize empty vectors to hold states and engine replacement
531   transitions
532   x_values <- rep(0, n_periods)
533   i_values <- rep(0, n_periods)
534   ### Every bus starts at state 0
535   x_values[1] <- 0

```

```

533 ##### Go through the progression
534 for (j in 1:length(x_values)) {
535   ##### Make a decision based on current state
536   i_values[j] = sample(c(0,1),1,replace=T,prob = c(p_ix_hat[x_values[j] +
537     1,1],p_ix_hat[x_values[j] + 1,2]))
538   ##### If decision is to not replace, continue on and increment x randomly
539   if (i_values[j] == 0) {
540     x_values[j+1] = x_values[j] + sample(c(0,1,2),1,replace = T, prob = c
541       (theta_30,theta_31,theta_32))
542     ##### If decision is to replace, reset state to 0
543     } else {
544       x_values[j+1] = 0
545     }
546   }
547   ## Generate a decision for the last period, even though we never see the
548   fruits of that decision
549   i_values[length(i_values)] = sample(c(0,1),1,replace=T,prob = c(p_ix_hat[
550     x_values[length(x_values)] + 1,1],
551     p_ix_hat[
552       x_
553       values
554       [
555         length
556         (x_
557         values
558         )] +
559         1,2])
560     )
561   ##### Return the states and replacement decisions in a list
562   results <- list('x_values' = x_values, 'i_values' = i_values)
563   return(results)
564 }
565 ##### Given a set of parameters, this function generates period-by-period
566 demand estimates for new buses (e.g.,
567 ##### how many buses will get their engine replaced in each period)
568 estimate_demand <- function(n_sims, n_buses, n_periods) {
569   ##### Initialize a vector to hold simulated demand
570   simulated_demand_total <- rep(0, n_periods)
571   ##### Run a bunch of simulations and simulate engine replacement decisions
572   for (j in 1:n_sims) {
573     simulated_demand_total = simulated_demand_total + simulate_sequence(n_
574       periods)$i_values
575   }
576   ##### Divide by the number of sims to get averages, multiply by number of
577   buses (this works because
578   #####buses are independent). Then return what we get.
579   return((n_buses/n_sims)*simulated_demand_total)
580 }
581 ##### Get demand as a function of RC for the first bus

```

```

571
572 ## Specify the range of RCs, as well as constants.
573 RC_range = seq(0, 15, .25)
574 n_periods = 15
575 n_sims = 1000
576 n_buses = 100
577 load('hotz_and_miller_estimate.Rdata')
578
579 ## Initialize an empty dataframe to hold results
580 estimated_demand_df <- data.frame(time_period = c(),
581                                   RC = c(),
582                                   demand = c(),
583                                   engine = c())
584
585 # Loop through the RCs, then estimate the probabilities using the Rust
586   method, then do simulation.
587 for (j in RC_range) {
588   RC = j
589
590   #### Set a critical value for to measure the deviation between iterative
591     updates of EV. The distance between the two EV matrices
592   #### is the infinity norm of the difference
593   cri <- 10^(-8)
594
595   #### Set an initial value for the EV matrix (all 0s, EV), and another EV
596     object to hold the updated estimates, EV2.
597   EV <- matrix(100,33,2)
598   EV2 <- matrix(0,33,2)
599
600   ## While the infinity norm is less than the threshold, iterate
601   while(max(abs(EV-EV2))>cri){
602     #### Set the current EV to the previous updated EV
603     EV <- EV2
604     #### Compute a new updated EV by iterating on the current EV
605     EV2 <- value.Iterate(EV)
606   }
607
608   #### Do one last update to set EV equal to the last EV2
609   EV <- EV2
610
611   # get EV(x,i) for x=0,1,2,...,30
612   #### EV contains extra states, which we needed to compute the above
613     computation. Throw them away.
614   EV <- EV[1:31,]
615
616   #### Get estimated probability based on the above EV
617   p_ix_hat <- choice.prob.Estimate()
618   #### Estimate demand using that probability
619   estimated_demand <- estimate_demand(n_sims, n_buses, n_periods)
620
621   #### Add this estimate to a temp dataframe
622   estimated_demand_df_temp <- data.frame(time_period = seq(1, n_periods, 1)
623     ,

```

```

620         RC = rep(RC, n_periods),
621         demand = estimated_demand,
622         engine = rep('Engine_1', n_periods))
623     ### Collate temp dataframe to full dataframe
624     estimated_demand_df <- rbind(estimated_demand_df, estimated_demand_df_
        temp)
625
626 }
627
628 ### Reset theta_1 to the "new engine", redo the exercise above.
629 theta_1 = .02
630
631 ### Loop through RCs
632 for (j in RC_range) {
633     RC = j
634
635     ### Set a critical value for to measure the deviation between iterative
        updates of EV. The distance between the two EV matrices
636     ### is the infinity norm of the difference
637     cri <- 10^(-8)
638
639     ### Set an initial value for the EV matrix (all 0s, EV), and another EV
        object to hold the updated estimates, EV2.
640     EV <- matrix(100,33,2)
641     EV2 <- matrix(0,33,2)
642
643     ## While the infinity norm is less than the threshold, iterate
644     while(max(abs(EV-EV2))>cri){
645
646         ### Set the current EV to the previous updated EV
647         EV <- EV2
648         ### Compute a new updated EV by iterating on the current EV
649         EV2 <- value.Iterate(EV)
650     }
651
652     ### Do one last update to set EV equal to the last EV2
653     EV <- EV2
654
655     # get EV(x,i) for x=0,1,2,...,30
656     ### EV contains extra states, which we needed to compute the above
        computation. Throw them away.
657     EV <- EV[1:31,]
658
659     ### Get probability estimates based on EV
660     p_ix_hat <- choice.prob.Estimate()
661     ### Estimate demand
662     estimated_demand <- estimate_demand(n_sims, n_buses, n_periods)
663
664     ### Add to temp dataframe
665     estimated_demand_df_temp <- data.frame(time_period = seq(1, n_periods, 1)
        ,
666
667         RC = rep(RC, n_periods),
        demand = estimated_demand,

```

```

668         engine = rep('Engine_2', n_periods
669                       ))
670     ### Collate to full dataframe
671     estimated_demand_df <- rbind(estimated_demand_df, estimated_demand_df_
672     temp)
673 }
674 ### For a reduced set of RCs, see the period-by-period demand
675 per_period_demand_plot <- estimated_demand_df %>%
676   filter(RC %in% c(1, 3, 5, 7, 10)) %>%
677   ggplot(., aes(x=time_period, y=demand, color=as.factor(RC))) + geom_line()
678   +
679   facet_wrap(~engine)
680   ggsave(per_period_demand_plot, file='per_period_demand_plot.png', height=4,
681         width=4, units='in')
682 ### Aggregate over periods to get demand as a function of RC for different
683   thetas.
684   aggregate_demand_plot <- estimated_demand_df %>%
685   group_by(RC, engine) %>%
686   summarise(total_demand = sum(demand)) %>%
687   ungroup() %>%
688   ggplot(., aes(x=RC, y=total_demand, color=engine)) + geom_line()
689   ggsave(aggregate_demand_plot, file='aggregate_demand_plot.png', height=4,
690         width=4, units='in')
691 #####
692 # Question 3.5 #
693 #####

```

Listing 1: ./rust.R