14.273 Industrial Organization: Problem Set 4

Dave Holtz, Jeremy Yang

May 18, 2017

1. Model setup.

Following the notations in Rust [2], HZ's flow utility is:

$$u(x_t, i_t, \theta_1) + \epsilon_t(i_t) = \begin{cases} -RC - c(0, \theta_1) + \epsilon_t(1) & i_t = 1 \\ -c(x_t, \theta_1) + \epsilon_t(0) & i_t = 0 \end{cases}$$

where RC is the replacement cost, x_t is the observed state variable for mileage, $c(\cdot)$ is cost function and i_t is the decision to replace engine and $\epsilon_t(\cdot)$ is action specific and type I extreme value distributed structural error (or unobserved state variable).

The state transition probability is given by:

$$\theta_{3j} = \mathbb{P}(x_{t+1} = x_t + j | x_t, i_t = 0)$$

$j \in \{0, 1, 2\}$ and if $i_t = 1$ then $x_{t+1} = 0$ with probability 1.

HZ chooses i_t in every period t to maximize an infinite sum of discounted flow utilities. The maximal value is defined as the value function (suppress the dependency on θ_1, θ_3):

$$V(x_1, \epsilon_1) := \max_{i_t, t \in \{1, 2, ..\}} \mathbb{E}[\sum_{t=1}^{\infty} \beta^{t-1}(u(x_t, i_t, \theta_1) + \epsilon_t(i_t))]$$

Rewrite the value function as in the Bellman optimality form:

$$V(x_t, \epsilon_t) = \max_{i_t} (u(x_t, i_t, \theta_1) + \epsilon_t(i_t)) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t]$$

where the expectation is with respect to (conditional) state transition probability of both x and ϵ, see Rust [2] equation (4.5). The Bellman equation breaks the dynamic optimization problem into an infinite series of static choices.

2. (1) The choice specific value function can be derived by plugging a specific action into the value function:

$$\tilde{V}(x_t, \epsilon_t, i_t) = \begin{cases} -RC - c(0, \theta_1) + \epsilon_t(1) + \beta\mathbb{E}[V(x_{t+1}, \epsilon_{t+1})|x_t, i_t = 1] \\ -c(x_t, \theta_1) + \epsilon_t(0) + \beta\mathbb{E}[V(x_{t+1}, \epsilon_{t+1})|x_t, i_t = 0] \end{cases}$$

$$V(x_t, \epsilon_t) = \max\{\tilde{V}(x_t, \epsilon_t, 1), \tilde{V}(x_t, \epsilon_t, 0)\}$$

HZ's decision is about trading off the total (future) cost of maintaining an old engine and the lump sum cost of replacing to a new one. The time to replace is the stopping time in this problem, so it can be thought as an optimal stopping time problem where the optimal policy is characterized by a cutoff in x, HZ would choose to replace the engine if x is above that threshold (the threshold depends on realized value of ϵ).

(2) It's clear from 2 (1) that the optimal stopping rule is:

$$- RC - c(0, \theta_1) + \epsilon_t(1) + \beta\mathbb{E}[V(x_{t+1}, \epsilon_{t+1})|x_t, i_t = 1] >$$
$$- c(x_t, \theta_1) + \epsilon_t(0) + \beta\mathbb{E}[V(x_{t+1}, \epsilon_{t+1})|x_t, i_t = 0]$$

or,

$$\tilde{V}(x_t, \epsilon_t, 1) > \tilde{V}(x_t, \epsilon_t, 0)$$

therefore, because the errors are type I extreme value distributed:

$$\mathbb{P}(i_t = 1|x_t) = \frac{\exp(u(x_t, 1, \theta_1) + \beta\mathbb{E}[V_{t+1}|x_t, i_t = 1])}{\sum_{k=\{0,1\}} \exp(u(x_t, k, \theta_1) + \beta\mathbb{E}[V_{t+1}|x_t, i_t = k]} \qquad (2.1)$$

where $u(x_t, i_t, \theta_1)$ is defined in 1 and for convenience:

$$V_{t+1} := V(x_{t+1}, \epsilon_{t+1})$$

(3) For discrete x, under the assumption that the errors are type I extreme value distributed, we have (Rust [2] equation (4.14)):

$$EV(x, i) = \sum_y \log\{\sum_j \exp[u(y, j) + \beta EV(y, j)]\} \cdot p(y|x, i) \qquad (2.2)$$

where

$$EV(x, i) := \mathbb{E}[V_{t+1}|x_t, i_t]$$

and x, i are the state and choice of current period and y, j are the state and choice of the next period. Also note that here the transition probability does not depend on x_t but only on j (or Δx). To compute expected value function, we first need to estimate transition probability from the data, this can be done simply by counting:

$$\hat{\theta}_{30} = \frac{\sum_b \sum_t 1_{\{x_{bt+1} - x_{bt} = 0, i_{bt} = 0\}}}{\sum_b \sum_t 1_{\{i_{bt} = 0\}}}$$

$$\hat{\theta}_{31} = \frac{\sum_b \sum_t \mathbf{1}_{\{x_{bt+1}-x_{bt}=1, i_{bt}=0\}}}{\sum_b \sum_t \mathbf{1}_{\{i_{bt}=0\}}}$$

$$\hat{\theta}_{32} = \frac{\sum_b \sum_t \mathbf{1}_{\{x_{bt+1}-x_{bt}=2, i_{bt}=0\}}}{\sum_b \sum_t \mathbf{1}_{\{i_{bt}=0\}}}$$

we compute the expected value function in the inner loop of the nested fixed point algorithm (holding the value of θ fixed), we first guess the initial values of $EV(x,i)$ for all possible values of x, i and use the equation (2.2) to iterate expected value function until it converges. The criterion is:

$$\max_{x,i} |EV^{T+1}(x,i) - EV^T(x,i)| < \eta$$

The plot for $x = 1-30$ at the true value of parameters are shown in Figure 1's left panel. The provided EV values for the Rust dataset are provided in Figure 1's right panel. The two match well, suggesting that our implementation is working as expected. There is some strange behavior in the tail (for mileage $\in$ (29, 30)). We suspect this is truncation error from not having data for higher states.
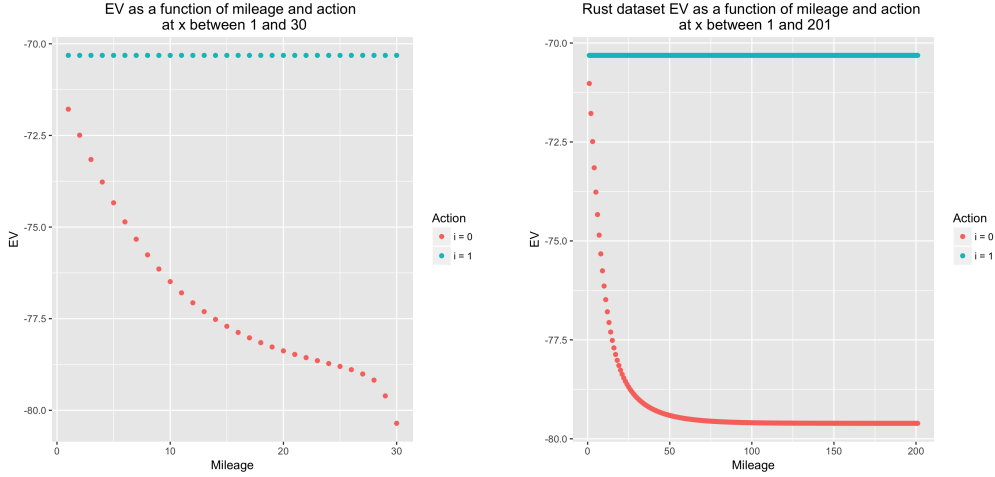

Figure 1: Expected Value Function for $i = 0$ and $i = 1$. Left panel shows results using iterative method, right panel shows provided Rust results.

(4) The provided dataset contains mileage and engine replacement information for 100 buses over 1,000 periods. The table below shows the mean mileage, maximum mileage, minimum mileage, standard deviation of the mileage, the average mileage at engine replacement across all buses and periods, and the average number of engine replacements for a particular bus over the 1,000 periods.

avg miles	max miles	min miles	s.d. miles	avg replace miles	avg replacements
8.245	33.000	0.000	5.709	15.953	52.980

We might also be interested in understanding how each of these summary statistics vary across buses. For instance, maybe some buses have their engines replaced much more often. In order to study this, Figure 2 shows the distributions of average mileage, maximum mileage, s.d. mileage, avg miles at replacement, and number of replacements across the 100 buses in the sample. In general, these distributions are quite concentrated, suggesting that there are not systematic differences across buses.

The final, bottom right plot in 2 also shows the empirically observed conditional choice probability as a function of state (mileage) that Harold Zurcher actually acts on. At a high level, Zurcher's has to make the investment decision of when to replace a given bus's engine. The mean replacement mileage plot suggests that on average he replaces a bus's engine after about 80,000 miles. The conditional choice probability plot suggests that the likelihood he increases the engine is practically zero until the bus hits 50,000 miles, after which the probability that the bus has its engine replaced climbs quickly. By the time a bus has 150,000 miles on it, it has a 50% probability of having its engine changed in a given time period.

3. (1) In the outer loop we search over a grid of values for (θ_1, β, RC), and compute the log likelihood function:

$$\log L = \sum_b \{\sum_t \log \mathbb{P}(i_{bt}|x_{bt}) + \sum_t \log \mathbb{P}(x_{bt}|x_{bt-1}, i_{t-1})\}$$

where b indexes for bus and t indexes for time period. We compute a log likelihood for each combination of values for (θ_1, β, RC) and choose the set of parameters that maximizes the log-likelihood of the data. The maximum likelihood parameters obtained with the Rust [2] method are:

$$\theta_1 = 0.05$$
$$\beta = 0.99$$
$$RC = 10$$

The maximum-likelihood parameters that we obtain using the Rust method are exactly identical to the "true" parameters.

(2) In Hotz-Miller's [1] approach, we will estimate the choice specific value function (as opposed to the expected value function as in Rust). We start by noting that conditional choice probability is observed directly from the data:

$$\hat{\mathbb{P}}(i = 1|x) = \frac{\sum_b \sum_t 1_{\{i_{bt}=1, x_{bt}=x\}}}{\sum_b \sum_t 1_{\{x_{bt}=x\}}}$$

The choice-specific value function (minus the structural error, and suppressing the dependency on θ_1, θ_3) can be presented recursively in the following form:

$$\tilde{V}(x_t, i_t) = u(x_t, i_t) + \beta \mathbb{E}_{x_{t+1}}[\mathbb{E}_{i_{t+1}}[\mathbb{E}_{\epsilon_{t+1}}[u(x_{t+1}, i_{t+1}) + \epsilon_{t+1} + \beta(\cdots)|i_{t+1}, x_{t+1}]|x_{t+1}]|x_t, i_t]$$
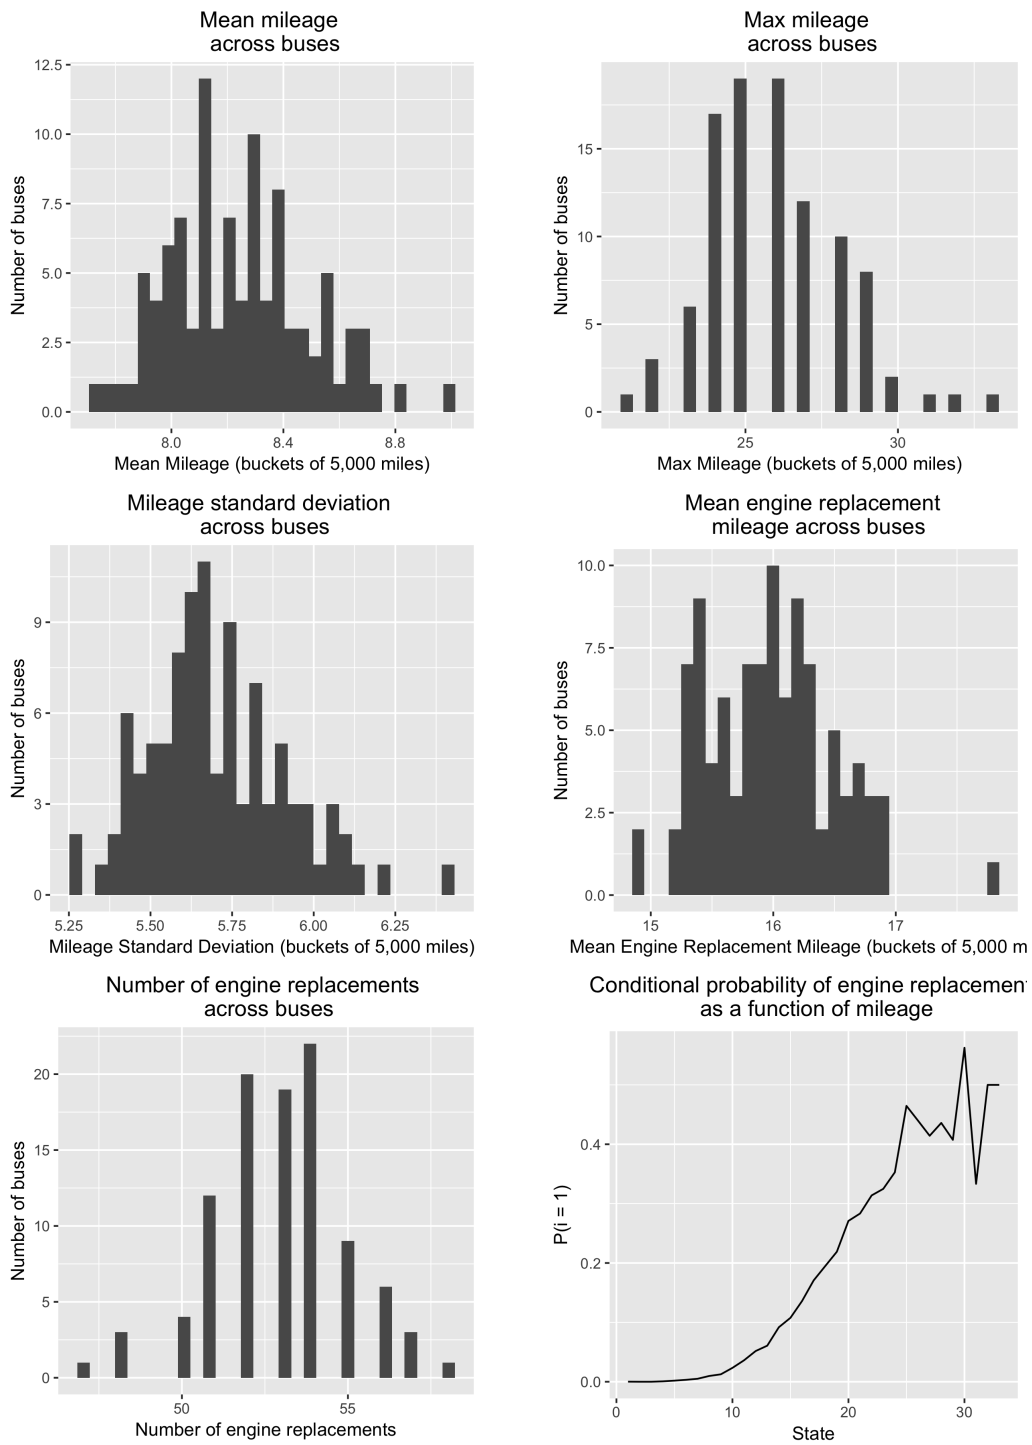
Figure 2: Distribution of various summary statistics across buses, as well as the empirical conditional choice probability for Zurcher.

where $(\cdots)$ represents higher (two and above) period forward expectations. In principle it's an infinite loop but in practice we need to stop at some T, for example, when $T = 2$, $(\cdots)$ simplifies to:

$$(\cdots) = \mathbb{E}_{x_{t+2}}[\mathbb{E}_{i_{t+2}}[\mathbb{E}_{\epsilon_{t+2}}[u(x_{t+2}, i_{t+2}) + \epsilon_{t+2}|i_{t+2}, x_{t+2}]|x_{t+2}]|x_{t+1}, i_{t+1}]$$

For simplicity, in the code we use one-period forward simulation where:

$$\tilde{V}(x_t, i_t) = u(x_t, i_t) + \beta \mathbb{E}_{x_{t+1}}[\mathbb{E}_{i_{t+1}}[\mathbb{E}_{\epsilon_{t+1}}[u(x_{t+1}, i_{t+1}) + \epsilon_{t+1}|i_{t+1}, x_{t+1}]|x_{t+1}]|x_t, i_t]$$

it is estimated as:

$$\hat{\tilde{V}}(x_t, i_t) = \frac{1}{S} \sum_s [u(x_t, i_t) + \beta[u(x_{t+1}^s, i_{t+1}^s) + \gamma - \log(\hat{\mathbb{P}}(i_{t+1}^s | x_{t+1}^s))]]$$

where x_{t+1}^s is drawn from the transition probability $\hat{\theta}_{30}, \hat{\theta}_{31}, \hat{\theta}_{32}$, and i_{t+1}^s is drawn from $\hat{\mathbb{P}}(i|x)$, γ is the Euler's constant.

We only go one period forward because we only observe data for states up to $x_t = 33$. It is possible for larger T that we would encounter a state that is not in our dataset. When this occurs, its unclear what value should be used as the conditional choice probability. While we avoid this issue by setting $T = 2$, this does reduce the precision of our estimates.

The estimates we obtain using the Hotz and Miller [1] method are

$$\begin{aligned} \theta_1 &= 0.09 \\ \beta &= 0.92 \\ RC &= 6. \end{aligned}$$

We expect that these estimates are likelihood less precise than those obtained via the Rust [2] method (in this particular case) because we so severely truncate the sum in the Hotz and Miller [1] method. The Rust approach relies on parametric assumptions that in this particular case appear to be satisfied, which is why it outperforms Hotz and Miller. However, in general, Hotz and Miller may be a "safer" approach (particularly when as a practitioner you are able to include more higher order terms in your sums).

(3) In order to determine which engine HZ prefers, we simply need to look at HZ's value function for both engines at $t = 0$ (which corresponds to $x_t = 0$ for all buses). There are a number of different mileage evolution paths that any given bus could take. However, the ex ante value function at time $= 0$ provides a weighted average of all of these scenarios. So at the outset, he will prefer whichever engine provides the most value in expectation. Given our estimation, HZ prefers the new engine over the old one, as

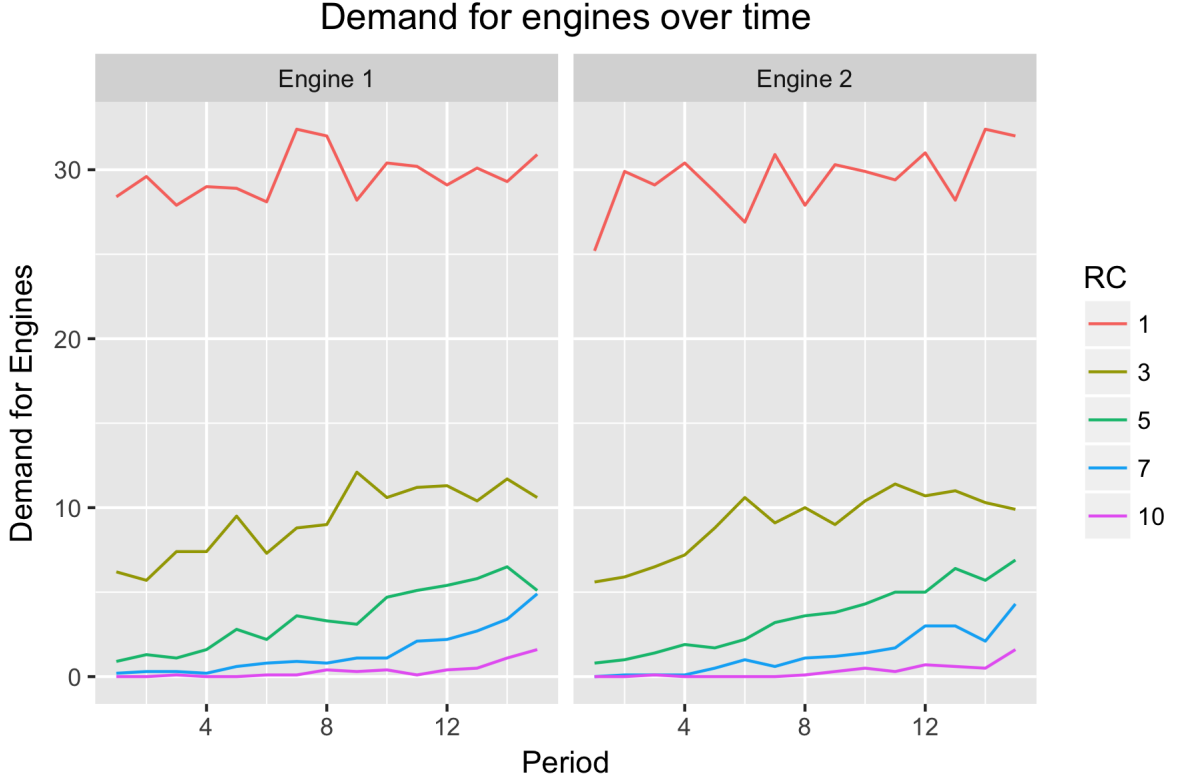$$EV_{old}(x = 0) = -71.026 < -65.214 = EV_{new}(x = 0).$$

Figure 3: The demand for engines across a fleet of 100 buses as a function of period (over the first 15 periods) for different values of RC. Unsurprisingly, when RC is lower, HZ is much more willing to change bus engines.

(4) We want to compute HZ's demand function for the two buses, which we will denote as engine 1 ($\theta_1 = 0.05, RC = 10$) and engine 2 ($\theta_1 = 0.02, RC = 20$) as a function of RC. In order to do so, we obtain conditional choice probability estimates, $\hat{\mathbb{P}}(i = 1|x)$ by using the Rust [2] method to iterate EV values. We use the Rust methodology because the Hotz and Miller [1] methodology depends on the observed conditional choice probabilities, which we know do *not* correspond to the counterfactual engine 2.

With those conditional choice probability estimates for the two engines in hand, we run 1,000 simulations of a bus's state transitions (and HZ's corresponding engine replacement decisions) over the first 15 periods. This allows us to get an expected, per-bus demand for engines over the first 15 periods. In order to get the expected demand that HZ has for engines across all buses, we simply multiply this figure by 100. So the expected demand (as a function of period t) is:

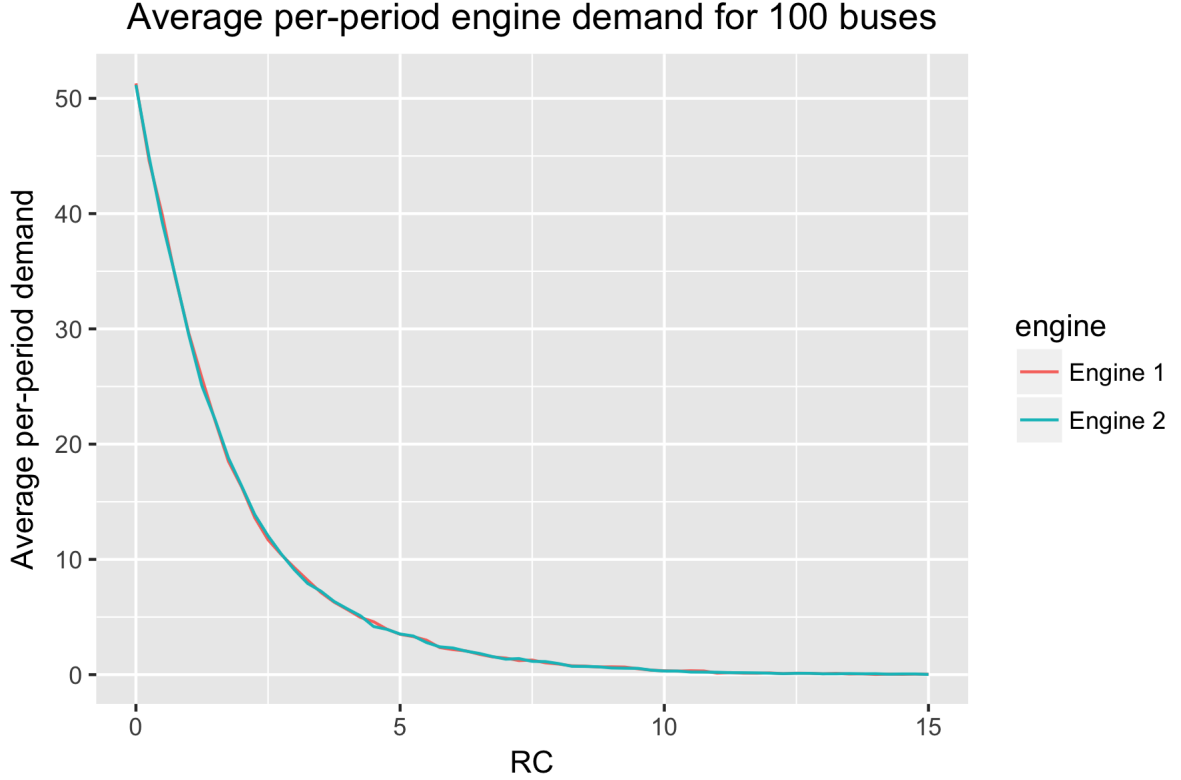$$D(t) = 100 \times \sum_{x_t} \hat{\mathbb{P}}(i = 1|x = x_t)\hat{\mathbb{P}}(x = x_t|t) \tag{1}$$

Figure 4: Aggregate per-period demand for new engines across a fleet of 100 buses as a function of RC. For engine 1, $\theta_1 = 0.05$. For engine 2, $\theta_1 = 0.02$.

HZ's demand for engines as a function of the period, t for a few values of RC can be found in Figure 3. The average per-period demand for the two engines (averaged across 15 periods) for different values of RC can be found in Figure 4. it's worth noting that the demand curves for the two engines appear almost identical - there are only small differences (on the order of a tenths of an engine) between the two. This could be a true difference, or it could be simulation error. Although we're not sure why these demand curves are so similar, we have two hypotheses:

1. Our Rust EV estimates are highly sensitive to initialization. It's possible our estimates for engine 2 are not correct.

2. The increase in RC from engine 1 to engine 2 is almost perfectly offset by the decrease in θ_1, creating two extremely similar demand curves.

(5) To determine the total value of the engines, assuming marginal cost RC, we can simply compute the total area to the right of a given RC in a demand curve that looks like Figure 4. This area will give the total surplus that HZ gets from the engine in a given period. In order to get a total value, we simply multiply this by the number of periods we want to consider. Mathematically,

it is socially optimal to produce the more efficient engine if the total value is greater than the total cost:

$$V_{engine}(RC) - C(RC) = n \cdot \int_{RC}^{\infty} D_{engine}(p)dp - n \cdot RC \cdot D_{engine}(RC) - c > 0 \quad (2)$$

where $n = 15$ is the number of periods and $D(p)$ is the amount of demand that HZ would have a given RC and c is the fixed R&D cost. Let's assume that c is 0 for engine 1 (the status quo engine). Figure 5 shows the value produced by both engines as a function of RC, assuming that $c = 0$. We can see here that, similarly to demand, the value produced by the two engines is very similar.
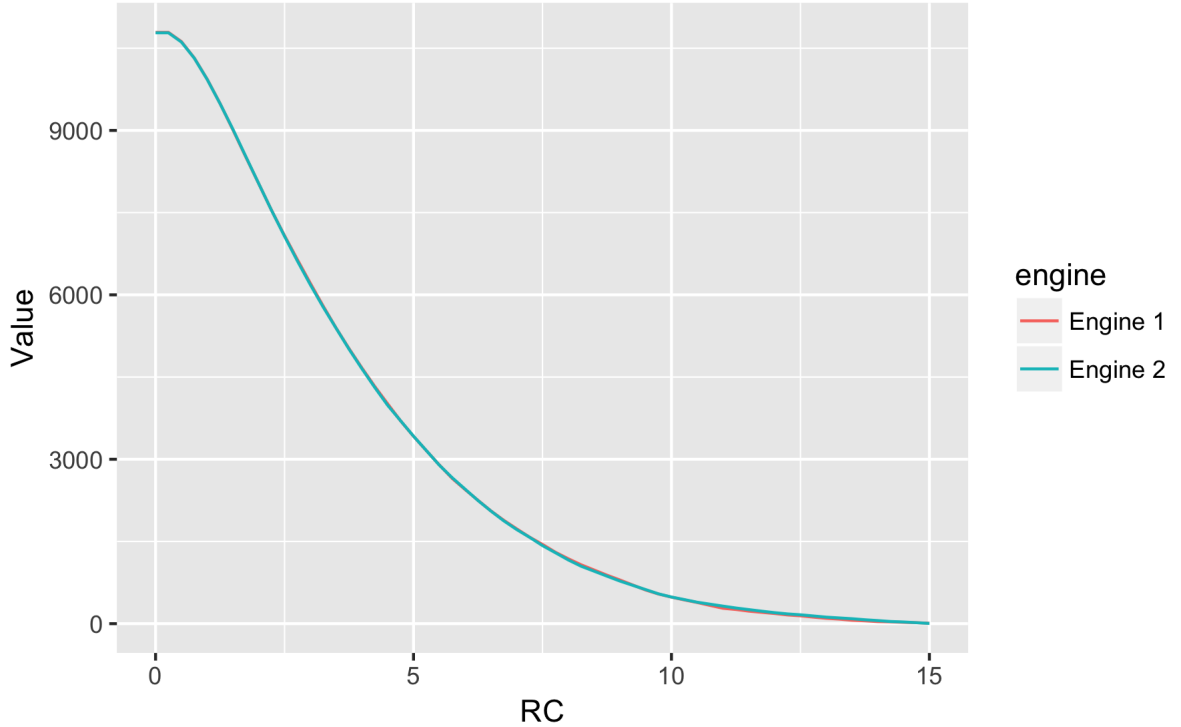

Figure 5: Value produced by the two different engines as a function of RC. We assume that $c = 0$ (i.e., R&D costs are zero) for both engines. For engine 1, $\theta_1 = 0.05$. For engine 2, $\theta_1 = 0.02$.

Theoretically, the maximum justifiable development would be $value_2 - value_1$, if that quantity were positive (otherwise it is zero). For engine 1 at $RC = 10$, the total value is 2,447.1. For engine 2 at $RC = 20$, the total value is 0 (because our estimated demand is zero at $RC = 20$). In light of this, there is no R&D cost at which development of engine 2 would be justified.

9

```r
1   #### This code uses the methodologies of both Rust (1987) and Hotz and
        Miller (1993) to estimate the parameters of a single
2   #### agent dynamic problem where an agent (Harold Zurcher) must choose when
        to have the engines replaced in a fleet of buses.
3
4   ## Import libraries
5   library(R.matlab)
6   library(ggplot2)
7   library(dplyr)
8
9   ## Read in data
10  setwd('~/Dropbox_(MIT)/MIT/Spring_2017/14.273/HW4/273-pset4/')
11  data <- readMat('../rust.mat')
12  gamma = .577
13
14  ## Extract bus replacement events
15  i <- data$it
16  ## Extract bus mileage counts (in increments of 5,000 miles)
17  x <- data$xt
18
19
20  ## Buses transition from different mileage states, and can jump forward
        zero, one, or two 5,000 mile buckets. This block
21  ## of code estimates the transition probabilities empirically from the data
        .
22
23  ### Initialize empty vectors to hold a count of how many times each jump
        happens
24  zero <- vector()
25  one <- vector()
26  two <- vector()
27
28  ### Loop through the 100 buses in the dataset
29  for (k in 1:100) {
30
31          ### Given a bus k, grab the mileage counts
32          xk <- x[,k]
33          ### Also grab the engine replacement events
34          ik <- i[,k]
35          ### Get a modified array which gives the change in mileage buckets
                from period j to period j+1
36          jk <- xk[-1]-xk[-1000]
37
38          ### We only care about periods where i=0 for transition
                probabilities, since i=1 will always send
39          ### x back to 0. This selects out only time periods for this bus
                where i = 0
40          j <- jk[ik==0]
41
42          ### This counts up how many times the mileage bucket counter, x,
                moves up by 0, 1, or 2 when i=0
```

```r
43              zero[k] <- length(j[j==0])
44              one[k] <- length(j[j==1])
45              two[k] <- length(j[j==2])
46  }
47
48  ## Estimate the x_t-independent transition probabilities by dividing the
        number of times for each transition by the
49  ## total number of transitions
50  theta_30 = sum(zero)/(sum(zero)+sum(one)+sum(two))
51  theta_31 = sum(one)/(sum(zero)+sum(one)+sum(two))
52  theta_32 = sum(two)/(sum(zero)+sum(one)+sum(two))
53
54  ##################
55  # Question 2.3 #
56  ##################
57
58  #### We'll now take the true values of the parameter values as given, and
        use the method described in Rust (1987) to iteratively
59  #### estimate the value function (or in this case, the EV function).
60
61  ## Initialize parameters to their true values
62  theta_1 = .05
63  theta_30 = .3
64  theta_31 = .5
65  theta_32 = .2
66  beta =.99
67  RC = 10
68
69  ### Define the linear cost function. If an engine is not replaced, the bus
        incurs cost theta_1*x, so cost
70  ### increases linearly as a bus gets older.
71  cost <- function(x){
72              return (theta_1*x)}
73
74  ### Define the utility function at mileage x from action i. If the agent
        chooses to replace the engine in a bus,
75  ### it costs RC. If they choose _not_ to replace the engine, they incur the
         cost of running the bus at mileage x.
76  u <- function(x,i){
77              -RC*i - cost(x*(1-i))
78  }
79
80
81  ### The value function can be estimated through an iteration procedure. We
        start with some initial guess for EV,
82  ### calculate EV with an expression that includes our initial guess of EV,
        and continue iterating until the difference
83  ### between subsequent EV estimates becomes small.
84
85  #### value.Iterate is a function to iteratively update the value function
        according to the methodology in Rust. The function
86  #### takes as an argument a current estimate of EV, and returns an updated
        estimate of EV. EV is an x by d matrix - we want the
```

```r
####  EV values for each decision d at every possible current mileage value
     x.
value.Iterate <- function(EV){

  ### First iterate through each of the 30 x states
  for (x in 1:31){
    ## Update the EV value corresponding to not replacing the engine. There
          are three contributions here - one from the
    ## j = 0 case, one from the j = 1 case, and one from the j = 2 case.
         Note the indexing here. When x =1, the state is
    ## equal to 0 (this is the x that needs to be passed into u()), but we
         want to grab the EV corresponding to the 1st entry.
    EV2[x,1] <- log(exp(u(x-1,0)+beta*EV[x,1])+exp(u(x-1,1)+beta*EV[x,2]))*
         theta_30  +
      log(exp(u(x,0)+beta*EV[x+1,1])+exp(u(x,1)+beta*EV[x+1,2]))*theta_31
         +
      log(exp(u(x+1,0)+beta*EV[x+2,1])+exp(u(x+1,1)+beta*EV[x+2,2]))*theta_
         32

    ## Update the EV value corresponding to replacing the engine. When the
         engine is replaced, x at the next period will
    ## deterministically reset to x = 0.
    EV2[x,2] <- log(exp(u(0,0)+beta*EV[1,1])+exp(u(0,1)+beta*EV[1,2]))
  }

  ## Return the updated EV values.
  return(EV2)
}

### Set a critical value for to measure the deviation between iterative
     updates of EV. The distance between the two EV matrices
### is the infinity norm of the difference
cri <- 10^(-8)

### Set an initial value for the EV matrix (all 0s, EV), and another EV
     object to hold the updated estimates, EV2.
EV <- matrix(0,33,2)
EV2 <- matrix(-80,33,2)

## While the infinity norm is less than the threshold, iterate
while(max(abs(EV-EV2))>cri){

  ### Set the current EV to the previous updated EV
  EV <- EV2
  ### Compute a new updated EV by iterating on the current EV
  EV2 <- value.Iterate(EV)
}

### Do one last update to set EV equal to the last EV2
EV <- EV2

# get EV(x,i) for x=0,1,2,..,30
### EV contains extra states, which we needed to compute the above
     computation. Throw them away.
```

```r
130 EV <- EV[1:31,]
131
132 ### Plot the EV of both replacing the engine (i = 1) and not replacing the
        engine (i = 0) at every x
133 ### between 1 and 30
134 df <- data.frame('x'=c(1:30, 1:30),'EV'=c(EV[2:31,1], EV[2:31,2]),'Action'
        = c(rep('i_=_0', 30), rep('i_=_1', 30)))
135
136 ### Generate a plot that compares the EV of replacing the engine and not
        replacing the engine
137 ev_plot <- ggplot(df, aes(x=x, y=EV, color=Action)) + geom_point() + xlab('
        Mileage') + ylab('EV') +
138   ggtitle('EV_as_a_function_of_mileage_and_action_\n_at_x_between_1_and_30'
        ) +
139   theme(plot.title = element_text(hjust = 0.5))
140 ggsave(ev_plot, file='ev_plot.png', height=6, width=6, units='in')
141
142 ### This is a plot to see the EV data in the attached rust matlab file. The
         state space is different than ours (200 states),
143 ### so its hard to compare. Our's is linear (seems wrong), whereas the
        provided data is not. However, the first 30 states
144 ### _do_ look approximately linear, so maybe we're not so far off.
145
146 df_rust <- data.frame('x'=c(seq(1,201), seq(1, 201)), 'EV'=c(data$EV[,1],
        data$EV[,2]), 'Action' = c(rep('i_=_0', 201),
147
148 ev_plot_rust <- ggplot(df_rust, aes(x=x, y=EV, color=Action)) + geom_point
        () + xlab('Mileage') + ylab('EV') +
149   ggtitle('Rust_dataset_EV_as_a_function_of_mileage_and_action_\n_at_x_
            between_1_and_201') +
150   theme(plot.title = element_text(hjust = 0.5))
151 ggsave(ev_plot_rust, file='ev_plot_rust.png', height=6, width=6, units='in'
        )
152
153 ###################
154 # Question 2.4 #
155 ###################
156
157 ### Calculate the mean mileage, mean time to engine replacement, max
        mileage, min mileage, and sd mileage over the whole sample
```

```r
158   mean_x <- mean(x)
159   mean_engine_replacement_age <- mean(x[i == 1])
160   max_x <- max(x)
161   min_x <- min(x)
162   sd_x <- sd(x)
163   avg_replacements <- mean(apply(i, 2, function(x) {sum(x)}))
164
165   aggregate_stats <- c(mean_x, max_x, min_x, sd_x, mean_engine_)
166   aggregate_stats %>%
167     round(., 3) %>%
168     kable(., format='latex')
169
170   ### Calculate the per bus mean mileage, mean time to engine replacement,
          max mileage, min mileage, and sd mileage
171   mean_x_per_bus <- apply(x, 2, function(x) {mean(x)})
172   max_x_per_bus <- apply(x, 2, function(x) {max(x)})
173   min_x_per_bus <- apply(x, 2, function(x) {min(x)})
174   sd_x_per_bus <- apply(x, 2, function(x) {sd(x)})
175   mean_engine_replacement_per_bus <- apply(x*i, 2, function(x) {sum(x)/sum(x
          != 0)})
176   replacements <- apply(i, 2, function(x) {sum(x)})
177
178
179   ### Collate per bus information into a dataframe
180   per_bus_statistics <- data.frame(bus = seq(1, 100, 1),
181                                     mean_x_per_bus = mean_x_per_bus,
182                                     max_x_per_bus = max_x_per_bus,
183                                     min_x_per_bus = min_x_per_bus,
184                                     sd_x_per_bus = sd_x_per_bus,
185                                     mean_engine_replacement_per_bus = mean_
                                          engine_replacement_per_bus,
186                                   replacements = replacements)
187
188   ### Create some plots
189   mean_mileage_plot <- ggplot(per_bus_statistics, aes(x=mean_x_per_bus)) +
          geom_histogram() +
190     xlab('Mean Mileage (buckets of 5,000 miles)') + ylab('Number of buses') +
            ggtitle('Mean mileage \n across buses') +
191     theme(plot.title = element_text(hjust = 0.5))
192   ggsave(mean_mileage_plot, file='mean_mileage_plot.png', height=4, width=4,
          units='in')
193
194   max_mileage_plot <- ggplot(per_bus_statistics, aes(x=max_x_per_bus)) + geom
          _histogram() +
195     xlab('Max Mileage (buckets of 5,000 miles)') + ylab('Number of buses') +
            ggtitle('Max mileage \n across buses') +
196     theme(plot.title = element_text(hjust = 0.5))
197   ggsave(max_mileage_plot, file='max_mileage_plot.png', height=4, width=4,
          units='in')
198
199   sd_mileage_plot <- ggplot(per_bus_statistics, aes(x=sd_x_per_bus)) + geom_
          histogram() +
200     xlab('Mileage Standard Deviation (buckets of 5,000 miles)') + ylab('
          Number of buses') +
```

```
201      ggtitle('Mileage_standard_deviation_\n_across_buses') +
202      theme(plot.title = element_text(hjust = 0.5))
203  ggsave(sd_mileage_plot, file='sd_mileage_plot.png', height=4, width=4,
         units='in')
204
205  time_to_engine_replacement_plot <- ggplot(per_bus_statistics, aes(x=mean_
         engine_replacement_per_bus)) + geom_histogram() +
206      xlab('Mean_Engine_Replacement_Mileage_(buckets_of_5,000_miles)') + ylab('
            Number_of_buses') +
207      ggtitle('Mean_engine_replacement_\n_mileage_across_buses') +
208      theme(plot.title = element_text(hjust = 0.5))
209  ggsave(time_to_engine_replacement_plot, file='time_to_engine_replacement_
         plot.png', height=4, width=4, units='in')
210
211  replacements_plot <- ggplot(per_bus_statistics, aes(x=replacements)) + geom
         _histogram() +
212      xlab('Number_of_engine_replacements') + ylab('Number_of_buses') +
213      ggtitle('Number_of_engine_replacements_\n_across_buses') +
214      theme(plot.title = element_text(hjust = 0.5))
215  ggsave(replacements_plot, file='replacements_plot.png', height=4, width=4,
         units='in')
216
217
218  ##################
219  # Question 3.1 #
220  ##################
221
222  ### This code will estimate the parameters beta, theta_1, and RC using the
         nested fixed-point algorithm
223  ### described in Rust.
224
225  ### A function to compute the probability of Zurcher's choices using the
         EVs we calculate using the EV calculation
226  ### framework above. The probability of choosing different actions
         basically acts like a multichoice logit function.
227  choice.prob.Estimate <- function(){
228
229      ### Initialize an empty matrix to hold choice probability estimates.
230      p_i <- matrix(0,31,2)
231
232      ### Iterate through all of the possible mileage states, x.
233      for (x in 1:31){
234        ## For each mileage state x, calculate the probability that Zurcher
              will choose i = 0.
235        p_i[x,1] <- exp(u(x-1,0)+beta*EV[x,1])/(exp(u(x-1,0)+beta*EV[x,1])+exp(
              u(x-1,1)+beta*EV[x,2]))
236        ## Calculate the probability of choosing i = 1 at each state, which is
              just 1 - P(i = 0).
237        p_i[x,2] <- 1- p_i[x,1]
238      }
239
240      ### Return the updated p_i object
241      return(p_i)
242  }
```

```r
243
244 ### A function to calculate the total log likelihood of the observed data
        given a set of parameters. This method assumes that
245 ### the probabilities across periods and buses are independent, so we can
        just add up all of the log probabilities.
246 log.likelihood.Compute <- function() {
247
248   ### Initialize 0-valued variables to hold the log choice probability, the
            log transition probability,
249   ### and the sum of the two.
250   log_choice_prob <- 0
251   log_transition_prob <- 0
252   total <- 0
253
254   ### Iterate over buses
255   for (bus in 1:100) {
256     ### Iterate over time periods
257     for (t in 1:999) {
258       ### We special case mileage states greater than 30, since they are a
                bit strange in our data. Otherwise, we calculate
259       ### the choice probability using the current value of p_i according
                to the EV values we calculated to get the
260       ### choice probability. Take the log and add it to the current
                running value.
261       if (x[t,bus] <= 30){
262         log_choice_prob <- log(p_i[x[t,bus]+1,i[t,bus]+1]) + log_choice_
                prob
263       ### Do the same thing for our special cased, x > 30 case.
264       } else {
265         log_choice_prob <- log(p_i[31,i[t,bus]+1]) + log_choice_prob
266       }
267
268       ### Calculate over the transitions for each bus the sum of the log
                transition probabilities. We have our estimates of
269       ### theta_3 given the empirical transition probabilities. So we can
                just grab that for each observed transition and add it
270       ### to the total log transition probability.
271
272       ### First we do the j = 0 case.
273       if (x[t+1,bus]-x[t,bus]==0) {
274         log_transition_prob <- log(theta_30)+log_transition_prob
275       ### Then the j = 1 case.
276       } else if (x[t+1,bus]-x[t,bus]==1) {
277         log_transition_prob <- log(theta_31)+log_transition_prob
278       ### And finally the j = 2 case.
279       } else if (x[t+1,bus]-x[t,bus]==2) {
280         log_transition_prob <- log(theta_32)+log_transition_prob
281       }
282     }
283
284     ### Now, get the total log likelihood by adding up all of the
            transition components and the choice components.
285     total <- (log_choice_prob+log_transition_prob) + total
286   }
```

```
287    return (total)
288  }
289
290  ### Now we're actually going to use the nested fixed point algorithm to get
             the maximum likelihood estimates of the parameters
291  ### that we care about. This process has three steps.
292
293  ### Step 1: We would calculate theta_30, theta_31, and theta_31 directly
         from the data. This step is not in the loop, and we've
294  ### actually already done this and it doesn't change, so we don't need to
         do it again.
295
296  ### Step 2: Next, we are going to set up a grid over values of theta_1,
         beta, and RC that we will calculate the
297  ### log likelihood to determine the maximum likelihood parameter values. We
         'll also initialize a dataframe
298  ### to hold the parameter values and the log likelihoods.
299
300  theta_1_range <- seq(.01,.10,.01)
301  beta_range <- seq(.90,.99,.01)
302  RC_range <- seq(6,15,1)
303  likelihood <- data.frame('theta_1'=rep(0),'beta'=rep(0),'RC'=rep(0),'log.
         likelihood'=rep(0))
304
305  ### Step 3: Now we actually do the nested fixed point computation.
306
307  ### Loop through theta_1
308  for (theta_1 in theta_1_range) {
309    ### Loop through beta
310    for (beta in beta_range) {
311      ### Loop through RC
312      for (RC in RC_range) {
313        print(paste(c(theta_1, beta, RC), collapse='_'))
314
315        ### Initialize the EV functions to the initial values we used above.
316        EV <- matrix(0,33,2)
317        EV2 <- matrix(-80,33,2)
318
319        ### Iteratively compute the EV values.
320        while(max(abs(EV-EV2))>cri){
321          EV <- EV2
322          EV2 <- value.Iterate(EV)
323        }
324
325        EV <- EV2
326        EV <- EV[1:31,]
327
328        ### Given these values of EV, calculated the choice probabilities
329        p_i <- choice.prob.Estimate()
330
331        ### Given the EV values, the choice probabilities and the parameters,
                 calculate
332        ### the log-likelihood of the data.
```

```r
333          likelihood <- rbind(likelihood ,c(theta_1,beta ,RC,log.likelihood.
                 Compute()))
334       }
335     }
336 }
337
338 ### Retrieve the row in the likelihood dataframe corresponding to the
        maximum likelihood estimate
339 likelihood <- likelihood[-1,]
340 parameter_estimates <- likelihood[which.max(likelihood[,4]),]
341
342 ### Use these parameters and get the relevant estimate of EV and p_i
343 theta_1 = parameter_estimates$theta_1
344 beta = parameter_estimates$beta
345 RC = parameter_estimates$RC
346
347 EV <- matrix(0,33,2)
348 EV2 <- matrix(-80,33,2)
349 ### Iteratively compute the EV values.
350 while(max(abs(EV-EV2))>cri){
351   EV <- EV2
352   EV2 <- value.Iterate(EV)
353 }
354 EV <- EV2
355 EV <- EV[1:31,]
356 ### Given these values of EV, calculated the choice probabilities
357 p_i <- choice.prob.Estimate()
358 ### Given the EV values, the choice probabilities and the parameters,
        calculate
359 ### the log-likelihood of the data.
360 likelihood <- rbind(likelihood ,c(theta_1,beta ,RC,log.likelihood.Compute()))
361
362 save(EV, p_i, likelihood, parameter_estimates, file='rust_estimate.Rdata')
363
364 ##################
365 # Question 3.2 #
366 ##################
367
368 ### Now we wil get estimates of the parameters using the Hotz and Miller
        conditional choice probability approach. This will
369 ### allow us to compare these parameter estimates to those obtained using
        the Rust approach.
370
371 ### First, we need to calculate the probability of the agent choosing
        either i = 0 or i = 1 based on the state that they find
372 ### a given bus in, x, at some time period t. This will be the baseline
        that we use to try and find the best parameter values
373 ### (i.e., which parameter values minimize the infinity norm between these
        true probabilities and the estimated probabilities)
374
375 ### The probability matrix
376 p_ix <- matrix(0,33,2)
377 ### The vector of how often the agent chooses i=1 given state x
378 ones <- vector()
```

```r
379  ### The vector of how often the agent finds a bus in state x
380  total <- vector()
381
382  ### Loop through the states
383  for (state in 0:32){
384
385    ### For a given state, a will track how many times i = 1 and b will track
              how many times that state occurs.
386    ### Initialize them to 0 for the given state.
387    a <- 0
388    b <- 0
389
390    ### Loop over the buses
391    for (bus in 1:100){
392
393      ### Increment how many times the agent chooses i = 1 in state x
394      a <- sum(i[which(x[,bus]==state),bus]) + a
395      ### Increment how many times the state x occurs
396      b <- length(i[which(x[,bus]==state),bus]) + b
397    }
398
399    ### Add the most recent estimates to the vector.
400    ones[state+1] <- a
401    total[state+1] <- b
402  }
403
404  ### Based on the ones and total vectors, updated the choice probability
        matrix.
405  p_ix[,1] <- 1-ones/total
406  p_ix[,2] <- ones/total
407
408  ### Plot conditional choice probabilities
409  p_ix_df <- as.data.frame(p_ix)
410  p_ix_df$state <- as.numeric(rownames(p_ix_df))
411  names(p_ix_df) <- c('P(i_=_0)', 'P(i_=_1)', 'State')
412  ccp_plot <- p_ix_df %>%
413    ggplot(., aes(x=State, y='P(i = 1)')) + geom_line() +
414    ggtitle('Conditional_probability_of_engine_replacement_\n_as_a_function_
            of_mileage') +
415    theme(plot.title = element_text(hjust = 0.5))
416  ggsave(ccp_plot, file='ccp_plot.png', height=4, width=4, units='in')
417
418  ### The function below uses the Hotz and Miller method to estimate V and p_
        ix_hat for every state and period
419  ### given a set of model parameters (beta, theta_1, and RC).
420  approximate.V_pixhat <- function() {
421    ### Initialize an empty valuation matrix
422    V <- matrix(0,33,2)
423    ### Initialize an empty conditional choice probability matrix
424    p_ix_hat <- matrix(0,33,2)
425
426    ### Iterate through the states
427    for (state in 0:30){
```

```r
428        ### Initialize a and b, which will basically track a running total of V
                for different choices over simulations, to 0.
429        a = 0
430        b = 0
431        ### Iterate through the simulations. Note that ideal we would probably
                want to go more than one time step into the
432        ### future. However, because of the limitations in our dataset, we only
                go one time step forward. This is mainly because
433        ### it's unclear how we would draw i (the choice) for states that do
                not appear in our data (i.e., x = 34).
434        for (s in 1:S){
435          ## Conditional on choosing i = 0, simulate the next state that a
                given bus will end up in by drawing from the
436          ## transition probabilities.
437          x_prime_0 = state + sample(c(0,1,2),1,replace = T, prob = c(theta_30,
                theta_31,theta_32))
438          ## Conditional on choosing i = 0 and ending up in some state in the
                next time period, randomly simulate a draw from
439          ## i based on the conditional choice probabilities
440          i_prime_0 = sample(c(0,1),1,replace=T,prob = c(p_ix[x_prime_0+1,1],p_
                ix[x_prime_0+1,2]))
441          # Figure out the expected utility from this truncated sequence of
                choices.
442          a = (u(state,0) + beta*(u(x_prime_0,i_prime_0)+gamma-log(p_ix[x_prime
                _0+1,i_prime_0+1]))) + a
443
444          ## Conditional on choosing i = 1, we don't need to simulate the next
                state that a bus will end up in. It will always
445          ## be x = 0. So we jump right to simulating the draw from i for x =
                0.
446          i_prime_1 = sample(c(0,1),1,replace=T,prob = c(p_ix[1,1],p_ix[1,2]))
447          ## Figure out the expected utility from this truncated sequence of
                choices.
448          b = (u(state,1) + beta*(u(0,i_prime_1)+gamma-log(p_ix[1,i_prime_1+1])
                )) + b
449        }
450
451        ## Set the value of V to be the average over all S of our simulations
                for both the i = 0 and i = 1 choices.
452        V[state+1,1] = a/S
453        V[state+1,2] = b/S
454        ## Use the multinomial logit-esque probability expression to figure out
                the probability of choosing i = 0 or i = 1
455        ## given that the bus is in state x.
456        p_ix_hat[state+1,1] <- exp(V[state+1,1])/(exp(V[state+1,1])+exp(V[state
                +1,2]))
457        p_ix_hat[state+1,2] <- 1- p_ix_hat[state+1,1]
458      }
459
460    # Put final output into a list and return it
461    results <- list('V' = V, 'p_ix_hat' = p_ix_hat)
462    return(results)
463 }
464
```

```r
### Specify a number of constants that will be used in the Hotz and Miller
    algorithm:
### S: The number of "simulations" to do per state / decision
### gamma: This should be Euler's constant
### theta_1_range: The range of theta_1 values to test
### beta_range: The range of beta_values to test
### RC_range: The range of RC values to test
S = 1000
theta_1_range <- seq(.01,.10,.01)
beta_range <- seq(.90,.99,.01)
RC_range <- seq(6,15,1)

### Initialize a dataframe to hold different parameter combinations and the
    infinity-norm between the actual conditional
### choice probabilities and the estimated ones
difference <- data.frame('theta_1'=rep(0),'beta'=rep(0),'RC'=rep(0),'
    difference'=rep(0))

### Loop through theta_1
for (theta_1 in theta_1_range) {
  ### Loop through theta_2
  for (beta in beta_range) {
    ### Loop through RC
    for (RC in RC_range) {
      # Check progress
      print(paste(c(theta_1, beta, RC), collapse='_'))

      ### Get estimates of V and P_ix_hat using the Hotz and Miller method
      v_and_p_ix_hat <- approximate.V_pixhat()
      V = v_and_p_ix_hat$V
      p_ix_hat <- v_and_p_ix_hat$p_ix_hat

      ### Now that we have a full conditional choice probability matrix,
          calculate the infinity norm (i.e., largest
      ### absolute difference between the empirical conditional choice
          probabilities and those estimated with the
      ### given parameters)
      difference <- rbind(difference,c(theta_1,beta,RC,max(abs(p_ix[1:31,]-
          p_ix_hat[1:31,]))))
    }
  }
}

### Find the set of parameters that minimizes this difference
difference <- difference[-1,]
parameter_estimates <- difference[which.min(difference[,4]),]

### Use these parameters and get the relevant estimate of V and p_ix_hat
theta_1 = parameter_estimates$theta_1
beta = parameter_estimates$beta
RC = parameter_estimates$RC
best_guesses <- approximate.V_pixhat()
V <- best_guesses$V
p_ix_hat <- best_guesses$p_ix_hat
```

```r
513
514  save(V, p_ix_hat, difference, parameter_estimates, file='hotz_and_miller_
         estimate.Rdata')
515
516  ##################
517  # Question 3.3 #
518  ##################
519
520  ### old engine
521  theta_1 = .05
522  RC = 10
523
524  ### apply Rust's approach
525  cri <- 10^(-8)
526  EV <- matrix(0,33,2)
527  EV2 <- matrix(-80,33,2)
528
529  while(max(abs(EV-EV2))>cri){
530    ### Set the current EV to the previous updated EV
531    EV <- EV2
532    ### Compute a new updated EV by iterating on the current EV
533    EV2 <- value.Iterate(EV)
534  }
535
536  ### Do one last update to set EV equal to the last EV2
537  EV_old <- EV2
538
539  ### new engine
540  theta_1 = .02
541  RC = 20
542
543  ### repeat the exercise above
544  cri <- 10^(-8)
545  EV <- matrix(0,33,2)
546  EV2 <- matrix(-80,33,2)
547
548  while(max(abs(EV-EV2))>cri){
549    ### Set the current EV to the previous updated EV
550    EV <- EV2
551    ### Compute a new updated EV by iterating on the current EV
552    EV2 <- value.Iterate(EV)
553  }
554
555  ### Do one last update to set EV equal to the last EV2
556  EV_new <- EV2
557
558  ### Returns which engine is preferred (old or new)
559  c('old','new')[which.max(c(EV_old[1,1],EV_new[1,1]))]
560
561  ##################
562  # Question 3.4 #
563  ##################
564
```

```r
565  ### This function simulates, for one agent, a sequence of state transitions
         and also engine replacement decisions
566  simulate_sequence <- function(n_periods) {
567    ### Initialize empty vectors to hold states and engine replacement
           transitions
568    x_values <- rep(0, n_periods)
569    i_values <- rep(0, n_periods)
570    ### Every bus starts at state 0
571    x_values[1] <- 0
572    ### Go through the progression
573    for (j in 1:length(x_values)) {
574      ### Make a decision based on current state
575      i_values[j] = sample(c(0,1),1,replace=T,prob = c(p_ix_hat[x_values[j] +
               1,1],p_ix_hat[x_values[j] + 1,2]))
576      ### If decision is to not replace, continue on and increment x randomly
577      if (i_values[j] == 0) {
578        x_values[j+1] = x_values[j] + sample(c(0,1,2),1,replace = T, prob = c
                (theta_30,theta_31,theta_32))
579      ### If decision is to replace, reset state to 0
580      } else {
581        x_values[j+1] = 0
582      }
583    }
584    ## Generate a decision for the last period, even though we never see the
           fruits of that decision
585    i_values[length(i_values)] = sample(c(0,1),1,replace=T,prob = c(p_ix_hat[
           x_values[length(x_values)] + 1,1],
586                                                                    p_ix_hat[
                                                                        x_
                                                                        values
                                                                        [
                                                                        length
                                                                        (x_
                                                                        values
                                                                        )] +
                                                                        1,2])
                                                                        )
587    ### Return the states and replacement decisions in a list
588    results <- list('x_values' = x_values, 'i_values' = i_values)
589    return(results)
590  }
591
592  ### Given a set of parameters, this function generates period-by-period
       demand estimates for new buses (e.g.,
593  ### how many buses will get their engine replaced in each period)
594  estimate_demand <- function(n_sims, n_buses, n_periods) {
595
596    ### Initialize a vector to hold simulated demand
597    simulated_demand_total <- rep(0, n_periods)
598
599    ### Run a bunch of simulations and simulate engine replacement decisions
600    for (j in 1:n_sims) {
601      simulated_demand_total = simulated_demand_total + simulate_sequence(n_
             periods)$i_values
```

```r
602      }

604      ### Divide by the number of sims to get averages, multiply by number of
             buses (this works because
605      ###buses are independent). Then return what we get.
606      return((n_buses/n_sims)*simulated_demand_total)
607    }

609    ### Get demand as a function of RC for the first bus

611    ## Specify the range of RCs, as well as constants.
612    RC_range = seq(0, 15, .25)
613    n_periods = 15
614    n_sims = 1000
615    n_buses = 100
616    load('rust_estimate.Rdata')

618    ## Initialize an empty dataframe to hold results
619    estimated_demand_df <- data.frame(time_period = c(),
620                                        RC = c(),
621                                        demand = c(),
622                                        engine = c())

624    # Loop through the RCs, then estimate the probabilities using the Rust
           method, then do simulation.
625    for (j in RC_range) {
626      RC = j

628      ### Set a critical value for to measure the deviation between iterative
             updates of EV. The distance between the two EV matrices
629      ### is the infinity norm of the difference
630      cri <- 10^(-8)

632      ### Set an initial value for the EV matrix (all 0s, EV), and another EV
             object to hold the updated estimates, EV2.
633      EV <- matrix(0,33,2)
634      EV2 <- matrix(-80,33,2)

636      ## While the infinity norm is less than the threshold, iterate
637      while(max(abs(EV-EV2))>cri){

639        ### Set the current EV to the previous updated EV
640        EV <- EV2
641        ### Compute a new updated EV by iterating on the current EV
642        EV2 <- value.Iterate(EV)
643      }

645      ### Do one last update to set EV equal to the last EV2
646      EV <- EV2

648      # get EV(x,i) for x=0,1,2,..,30
649      ### EV contains extra states, which we needed to compute the above
             computation. Throw them away.
650      EV <- EV[1:31,]
```

```r
651
652    ### Get estimated probability based on the above EV
653    p_ix_hat <- choice.prob.Estimate()
654    ### Estimate demand using that probability
655    estimated_demand <- estimate_demand(n_sims, n_buses, n_periods)
656
657    ### Add this estimate to a temp dataframe
658    estimated_demand_df_temp <- data.frame(time_period = seq(1, n_periods, 1)
            ,
659                                           RC = rep(RC, n_periods),
660                                           demand = estimated_demand,
661                                           engine = rep('Engine_1', n_periods))
662    ### Collate temp dataframe to full dataframe
663    estimated_demand_df <- rbind(estimated_demand_df, estimated_demand_df_
          temp)
664
665  }
666
667  ### Reset theta_1 to the "new engine", redo the exercise above.
668  theta_1 = .02
669
670  ### Loop through RCs
671  for (j in RC_range) {
672    RC = j
673
674    ### Set a critical value for to measure the deviation between iterative
            updates of EV. The distance between the two EV matrices
675    ### is the infinity norm of the difference
676    cri <- 10^(-8)
677
678    ### Set an initial value for the EV matrix (all 0s, EV), and another EV
            object to hold the updated estimates, EV2.
679    EV <- matrix(0,33,2)
680    EV2 <- matrix(-80,33,2)
681
682    ## While the infinity norm is less than the threshold, iterate
683    while(max(abs(EV-EV2))>cri){
684
685      ### Set the current EV to the previous updated EV
686      EV <- EV2
687      ### Compute a new updated EV by iterating on the current EV
688      EV2 <- value.Iterate(EV)
689    }
690
691    ### Do one last update to set EV equal to the last EV2
692    EV <- EV2
693
694    # get EV(x,i) for x=0,1,2,..,30
695    ### EV contains extra states, which we needed to compute the above
            computation. Throw them away.
696    EV <- EV[1:31,]
697
698    ### Get probability estimates based on EV
699    p_ix_hat <- choice.prob.Estimate()
```

```
700    ### Estimate demand
701    estimated_demand <- estimate_demand(n_sims, n_buses, n_periods)
702
703    ### Add to temp dataframe
704    estimated_demand_df_temp <- data.frame(time_period = seq(1, n_periods, 1)
          ,
705                                             RC = rep(RC, n_periods),
706                                             demand = estimated_demand,
707                                             engine = rep('Engine_2', n_periods
                                                 ))
708    ### Collate to full dataframe
709    estimated_demand_df <- rbind(estimated_demand_df, estimated_demand_df_
          temp)
710
711  }
712
713  ### For a reduced set of RCs, see the period-by-period demand
714  per_period_demand_plot <- estimated_demand_df %>%
715    filter(RC %in% c(1, 3, 5, 7, 10)) %>%
716    mutate(RC = as.factor(RC))  %>%
717  ggplot(., aes(x=time_period, y=demand, color=RC)) + geom_line() +
718    facet_wrap(~engine) + xlab('Period') + ylab('Demand_for_Engines') +
          ggtitle('Demand_for_engines_over_time') +
719    theme(plot.title = element_text(hjust = 0.5))
720  ggsave(per_period_demand_plot, file='per_period_demand_plot.png', height=4,
        width=6, units='in')
721
722  ### Aggregate over periods to get demand as a function of RC for different
        thetas.
723  aggregate_demand_plot <- estimated_demand_df %>%
724    group_by(RC, engine) %>%
725    summarise(total_demand = sum(demand)/n_periods) %>%
726    ungroup() %>%
727    ggplot(., aes(x=RC, y=total_demand, color=engine)) + geom_line() + xlab('
          RC') +
728    ylab('Average_per-period_demand') + ggtitle('Average_per-period_engine_
          demand_for_100_buses') +
729    theme(plot.title = element_text(hjust = 0.5))
730  ggsave(aggregate_demand_plot, file='aggregate_demand_plot.png', height=4,
        width=6, units='in')
731
732
733  ##################
734  # Question 3.5 #
735  ##################
736
737  n_periods = 15
738
739  ### for a given RC and engine, average demand across time periods
740  engine_value <- estimated_demand_df %>%
741  group_by(RC, engine) %>%
742   summarize(avg_demand=mean(demand))
743
```

```R
# Initialize an empty dataframe to hold value as a function of RC for both
      engines
engine_value_df <- data.frame(RC = c(), engine_value = c(), engine = c())

## Loop over RC values
for (j in RC_range) {
  # Get the engine 1 value at this RC
  engine_1_value <- engine_value %>%
    filter(engine=='Engine_1',RC>=j) %>%
    summarize(value = RC*avg_demand) %>%
    summarize(total_value = sum(value)) %>%
    as.numeric()* n_periods

  # Get the engine 2 value at this RC
  engine_2_value <- engine_value %>%
    filter(engine=='Engine_2',RC>=j) %>%
    summarize(value = RC*avg_demand) %>%
    summarize(total_value = sum(value)) %>%
    as.numeric()* n_periods

  # Add these two engine values to the dataframe
  engine_value_df <- rbind(engine_value_df, data.frame(RC = j, engine_value
      = engine_1_value, engine = 'Engine_1'))
  engine_value_df <- rbind(engine_value_df, data.frame(RC = j, engine_value
      = engine_2_value, engine = 'Engine_2'))
}

# Create a plot of value as a function of RC for both engines
value_plot <- ggplot(engine_value_df, aes(x=RC, y=engine_value, color=
    engine)) + geom_line() +
  xlab('RC') + ylab('Value') + ggtitle('Value_produced_by_each_engine_as_a_
      function_of_RC_(c=0)') +
  theme(plot.title = element_text(hjust = 0.5))
ggsave(value_plot, file='value_plot.png', height=4, width=6, units='in')

### for engine 1 (original one), calculate value at its true RC
engine_value %>%
filter(engine=='Engine_1',RC>=6) %>%
summarize(value = RC*avg_demand) %>%
summarize(total_value = sum(value)) %>%
as.numeric()* n_periods

### for engine 2 (new one), calculate value at its true RC
engine_value %>%
filter(engine=='Engine_2',RC>=20) %>%
summarize(value = RC*avg_demand) %>%
summarize(total_value = sum(value)) %>%
as.numeric()* n_periods
```

Listing 1: ./rust.R

27

References

[1] V Joseph Hotz and Robert A Miller. Conditional choice probabilities and the estimation of dynamic models. *The Review of Economic Studies*, 60(3):497–529, 1993.

[2] John Rust. Optimal replacement of gmc bus engines: An empirical model of harold zurcher. *Econometrica: Journal of the Econometric Society*, pages 999–1033, 1987.