
14.273 Industrial Organization: Pset4

Dave Holtz, Jeremy Yang

May 18, 2017

1. Model setup.

Following the notations in Rust (1987), HZ's flow utility is:

$$u(x_t, i_t, \theta_1) + \epsilon_t(i_t) = \begin{cases} -RC - c(0, \theta_1) + \epsilon_t(1) & i_t = 1 \\ -c(x_t, \theta_1) + \epsilon_t(0) & i_t = 0 \end{cases}$$

where RC is the replacement cost, x_t is the observed state variable for mileage, $c(\cdot)$ is cost function and i_t is the decision to replace engine and $\epsilon_t(\cdot)$ is action specific and type I extreme value distributed structural error (or unobserved state variable).

The state transition probability is given by:

$$\theta_{3j} = \mathbb{P}(x_{t+1} = x_t + j | x_t, i_t = 0)$$

$j \in \{0, 1, 2\}$ and if $i_t = 1$ then $x_{t+1} = 0$ with probability 1.

HZ chooses i_t in every period t to maximize an infinite sum of discounted flow utilities. The maximal value is defined as the value function (suppress the dependency on θ_1, θ_3):

$$V(x_1, \epsilon_1) := \max_{i_t, t \in \{1, 2, \dots\}} \mathbb{E} \left[\sum_{t=1}^{\infty} \beta^{t-1} (u(x_t, i_t, \theta_1) + \epsilon_t(i_t)) \right]$$

Rewrite the value function as in the Bellman optimality form:

$$V(x_t, \epsilon_t) = \max_{i_t} (u(x_t, i_t, \theta_1) + \epsilon_t(i_t)) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t]$$

where the expectation is with respect to (conditional) state transition probability of both x and ϵ , see Rust (1987) equation (4.5). The Bellman equation breaks the dynamic optimization problem into an infinite series of static choices.

2. (1) The choice specific value function can be derived by plugging a specific action into the value function:

$$\tilde{V}(x_t, \epsilon_t, i_t) = \begin{cases} -RC - c(0, \theta_1) + \epsilon_t(1) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 1] \\ -c(x_t, \theta_1) + \epsilon_t(0) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 0] \end{cases}$$

$$V(x_t, \epsilon_t) = \max\{\tilde{V}(x_t, \epsilon_t, 1), \tilde{V}(x_t, \epsilon_t, 0)\}$$

HZ's decision is about trading off the total (future) cost of maintaining an old engine and the lump sum cost of replacing to a new one. The time to replace is the stopping time in this problem, so it can be thought as an optimal stopping time problem where the optimal policy is characterized by a cutoff in x , HZ would choose to replace the engine if x is above that threshold (the threshold depends on realized value of ϵ).

- (2) It's clear from 2 (1) that the optimal stopping rule is:

$$\begin{aligned} & -RC - c(0, \theta_1) + \epsilon_t(1) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 1] > \\ & -c(x_t, \theta_1) + \epsilon_t(0) + \beta \mathbb{E}[V(x_{t+1}, \epsilon_{t+1}) | x_t, i_t = 0] \end{aligned}$$

or,

$$\tilde{V}(x_t, \epsilon_t, 1) > \tilde{V}(x_t, \epsilon_t, 0)$$

therefore, because the errors are type I extreme value distributed:

$$\mathbb{P}(i_t = 1 | x_t) = \frac{\exp(u(x_t, 1, \theta_1) + \beta \mathbb{E}[V_{t+1} | x_t, i_t = 1])}{\sum_{k=\{0,1\}} \exp(u(x_t, k, \theta_1) + \beta \mathbb{E}[V_{t+1} | x_t, i_t = k])} \quad (2.1)$$

where $u(x_t, i_t, \theta_1)$ is defined in 1 and for convenience:

$$V_{t+1} := V(x_{t+1}, \epsilon_{t+1})$$

- (3) For discrete x , under the assumption that the errors are type I extreme value distributed, we have (Rust (1987) equation (4.14)):

$$EV(x, i) = \sum_y \log\left\{ \sum_j \exp[u(y, j) + \beta EV(y, j)] \right\} \cdot p(y | x, i) \quad (2.2)$$

where

$$EV(x, i) := \mathbb{E}[V_{t+1} | x_t, i_t]$$

and x, i are the state and choice of current period and y, j are the state and choice of the next period. Also note that here the transition probability does not depend on x_t but only on j (or Δx). To compute expected value function, we first need to estimate transition probability from the data, this can be done simply by counting:

$$\hat{\theta}_{30} = \frac{\sum_b \sum_t 1_{\{x_{bt+1} - x_{bt} = 0, i_{bt} = 0\}}}{\sum_b \sum_t 1_{\{i_{bt} = 0\}}}$$

$$\hat{\theta}_{31} = \frac{\sum_b \sum_t 1_{\{x_{bt+1}-x_{bt}=1, i_{bt}=0\}}}{\sum_b \sum_t 1_{\{i_{bt}=0\}}}$$

$$\hat{\theta}_{32} = \frac{\sum_b \sum_t 1_{\{x_{bt+1}-x_{bt}=2, i_{bt}=0\}}}{\sum_b \sum_t 1_{\{i_{bt}=0\}}}$$

we compute the expected value function in the inner loop of the nested fixed point algorithm (holding the value of θ fixed), we first guess the initial values of $EV(x, i)$ for all possible values of x, i and use the equation (2.2) to iterate expected value function until it converges. The criterion is:

$$\max_{x,i} |EV^{T+1}(x, i) - EV^T(x, i)| < \eta$$

The plot for $x = 1 - 30$ at the true value of parameters are shown in Figure 1.

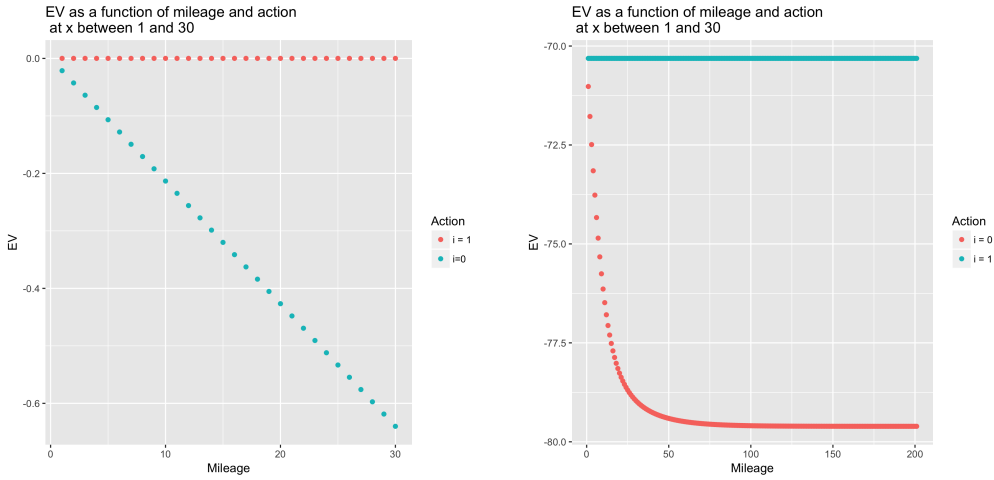


Figure 1: Expected Value Function for $i = 0$ and $i = 1$. Left panel shows results using iterative method, right panel shows provided Rust results.

Interestingly, our EV results are linear in mileage, which is probably not expected. Despite a good amount of debugging, we have been unable to identify a problem. However, it's also unclear how our calculated results for *just* 30 states should compare to the provided EV results, which provide information on 200 states. The first 30 states of the provided Rust EV estimates are decreasing in approximately linear fashion, suggesting our estimates might not be *so* bad. However, the order of magnitude of our EV values (e.g., 10^{-1}) is much smaller than the order of magnitude of EV values in the provided dataset (e.g., ~ 70), suggesting something is probably wrong. However, we don't have any more time to debug this, so we simply moved on.

- (4) The provided dataset contains mileage and engine replacement information for 100 buses over 1,000 periods. The table below shows the mean mileage, maximum mileage, minimum mileage, standard deviation of the mileage, the average mileage at engine replacement across all buses and periods, and the

average number of engine replacements for a particular bus over the 1,000 periods.

avg miles	max miles	min miles	s.d. miles	avg replace miles	avg replacements
8.245	33.000	0.000	5.709	15.953	52.980

We might also be interested in understanding how each of these summary statistics vary across buses. For instance, maybe some buses have their engines replaced much more often. In order to study this, Figure 2 shows the distributions of average mileage, maximum mileage, s.d. mileage, avg miles at replacement, and number of replacements across the 100 buses in the sample. In general, these distributions are quite concentrated, suggesting that there are not systematic differences across buses.

3. (1) In the outer loop we search over a grid of values for (θ_1, β, RC) , and compute the log likelihood function:

$$\log L = \sum_b \left\{ \sum_t \log \mathbb{P}(i_{bt} | x_{bt}) + \sum_t \log \mathbb{P}(x_{bt} | x_{bt-1}, i_{t-1}) \right\}$$

where b indexes for bus and t indexes for time period. We compute a log likelihood for each combination of values for (θ_1, β, RC) and choose the one that has the maximal value as our maximum likelihood estimation.

- (2) In Hotz-Miller's approach, we will estimate the choice specific value function (as opposed to the expected value function as in Rust). We start by noting that conditional choice probability is observed directly from the data:

$$\hat{\mathbb{P}}(i = 1 | x) = \frac{\sum_b \sum_t 1_{\{i_{bt}=1, x_{bt}=x\}}}{\sum_b \sum_t 1_{\{x_{bt}=x\}}}$$

The choice-specific value function (minus the structural error, and suppressing the dependency on θ_1, θ_3) can be presented recursively in the following form:

$$\tilde{V}(x_t, i_t) = u(x_t, i_t) + \beta \mathbb{E}_{x_{t+1}} [\mathbb{E}_{i_{t+1}} [\mathbb{E}_{\epsilon_{t+1}} [u(x_{t+1}, i_{t+1}) + \epsilon_{t+1} + \beta(\dots) | i_{t+1}, x_{t+1}] | x_{t+1}] | x_t, i_t]$$

where (\dots) represents higher (two and above) period forward expectations. In principle it's an infinite loop but in practice we need to stop at some T , for example, when $T = 2$, (\dots) simplifies to:

$$(\dots) = \mathbb{E}_{x_{t+2}} [\mathbb{E}_{i_{t+2}} [\mathbb{E}_{\epsilon_{t+2}} [u(x_{t+2}, i_{t+2}) + \epsilon_{t+2} | i_{t+2}, x_{t+2}] | x_{t+2}] | x_{t+1}, i_{t+1}]$$

For simplicity, in the code we use one-period forward simulation where:

$$\tilde{V}(x_t, i_t) = u(x_t, i_t) + \beta \mathbb{E}_{x_{t+1}} [\mathbb{E}_{i_{t+1}} [\mathbb{E}_{\epsilon_{t+1}} [u(x_{t+1}, i_{t+1}) + \epsilon_{t+1} | i_{t+1}, x_{t+1}] | x_{t+1}] | x_t, i_t]$$

it is estimated as:

$$\hat{\tilde{V}}(x_t, i_t) = \frac{1}{S} \sum_s [u(x_t, i_t) + \beta [u(x_{t+1}^s, i_{t+1}^s) + \gamma - \log(\hat{\mathbb{P}}(i_{t+1}^s | x_{t+1}^s))]]$$

where x_{t+1}^s is drawn from the transition probability $\hat{\theta}_{30}, \hat{\theta}_{31}, \hat{\theta}_{32}$, and i_{t+1}^s is drawn from $\hat{\mathbb{P}}(i | x)$, γ is the Euler's constant.

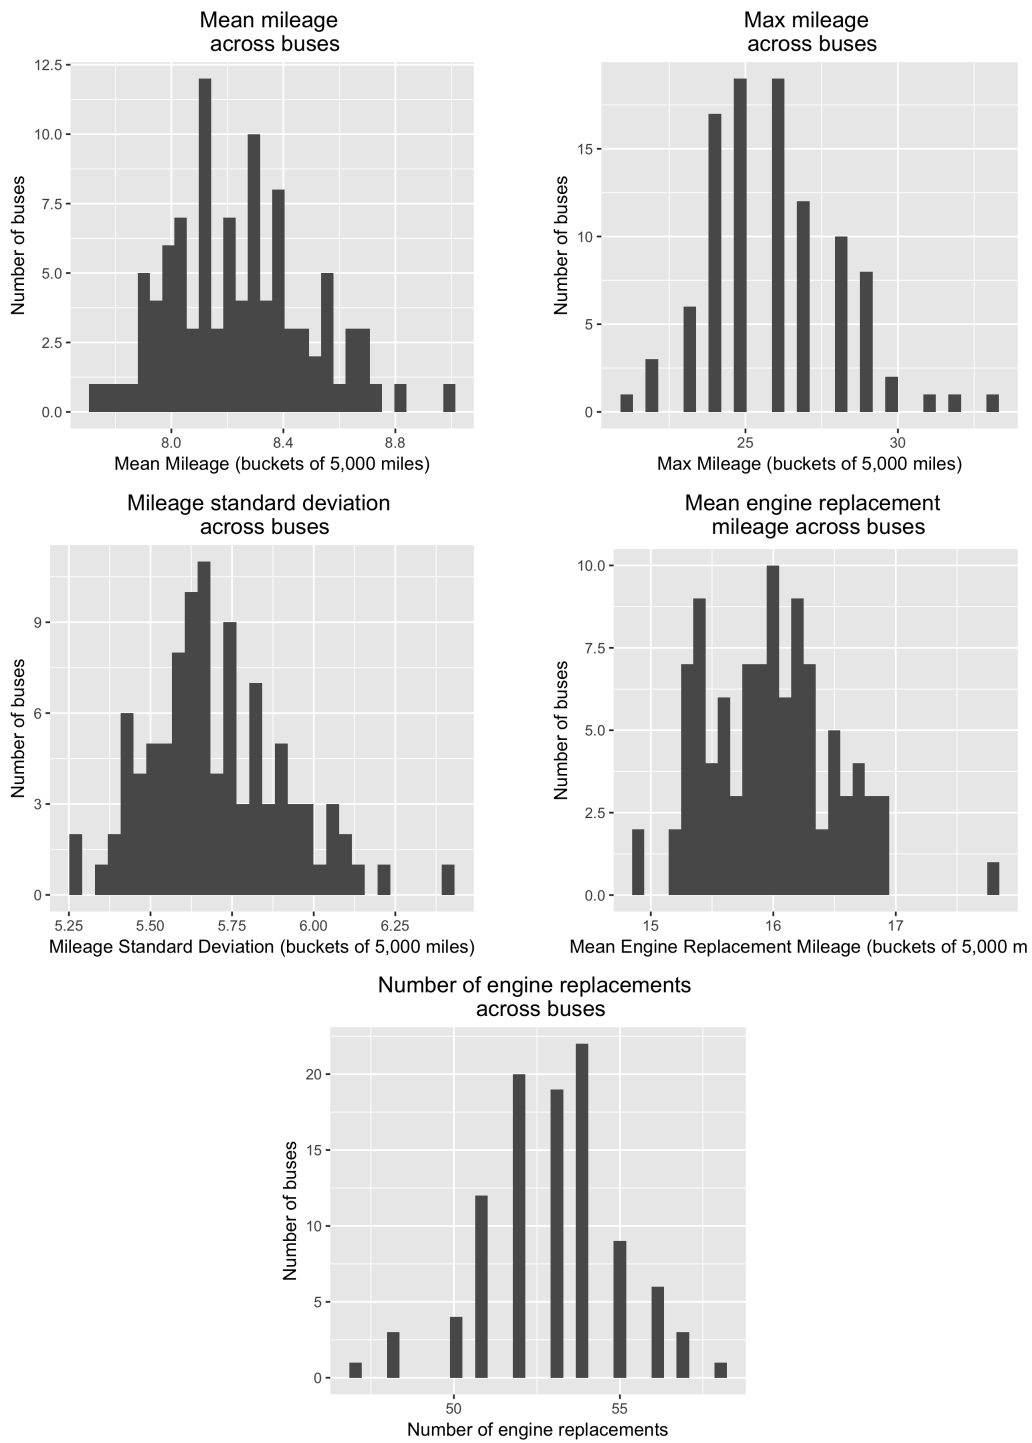


Figure 2: Distribution of various summary statistics across buses.

APPENDIX: CODE

```
1 ##### This code uses the methodologies of both Rust (1987) and Hotz and
   Miller (1993) to estimate the parameters of a single
2 ##### agent dynamic problem where an agent (Harold Zurcher) must choose when
   to have the engines replaced in a fleet of buses.
3
4 ## Import libraries
5 library(R.matlab)
6 library(ggplot2)
7 library(dplyr)
8
9 ## Read in data
10 setwd('~/.Dropbox/(MIT)/MIT/Spring_2017/14.273/HW4/273-pset4/')
11 data <- readMat('rust.mat')
12
13 ## Extract bus replacement events
14 i <- data$it
15 ## Extract bus mileage counts (in increments of 5,000 miles)
16 x <- data$xt
17
18
19 ## Buses transition from different mileage states, and can jump forward
   zero, one, or two 5,000 mile buckets. This block
20 ## of code estimates the transition probabilities empirically from the data
   .
21
22 ##### Initialize empty vectors to hold a count of how many times each jump
   happens
23 zero <- vector()
24 one <- vector()
25 two <- vector()
26
27 ##### Loop through the 100 buses in the dataset
28 for (k in 1:100) {
29
30     ##### Given a bus k, grab the mileage counts
31     xk <- x[,k]
32     ##### Also grab the engine replacement events
33     ik <- i[,k]
34     ##### Get a modified array which gives the change in mileage buckets
       from period j to period j+1
35     jk <- xk[-1]-xk[-1000]
36
37     ##### We only care about periods where i=0 for transition
       probabilities, since i=1 will always send
38     ##### x back to 0. This selects out only time periods for this bus
       where i = 0
39     j <- jk[ik==0]
40
41     ##### This counts up how many times the mileage bucket counter, x,
       moves up by 0, 1, or 2 when i=0
42     zero[k] <- length(j[j==0])
```

```

43     one[k] <- length(j[j==1])
44     two[k] <- length(j[j==2])
45 }
46
47 ## Estimate the x_t-independent transition probabilities by dividing the
48   number of times for each transition by the
49   total number of transitions
49 theta_30 = sum(zero)/(sum(zero)+sum(one)+sum(two))
50 theta_31 = sum(one)/(sum(zero)+sum(one)+sum(two))
51 theta_32 = sum(two)/(sum(zero)+sum(one)+sum(two))
52
53 #####
54 # Question 2.3 #
55 #####
56
57 ##### We'll now take the true values of the parameter values as given, and
58   use the method described in Rust (1987) to iteratively
59   estimate the value function (or in this case, the EV function).
60
61 ## Initialize parameters to their true values
61 theta_1 = .05
62 theta_30 = .3
63 theta_31 = .5
64 theta_32 = .2
65 beta = .99
66 RC = 10
67
68 ##### Define the linear cost function. If an engine is not replaced, the bus
69   incurs cost theta_1*x, so cost
70   increases linearly as a bus gets older.
70 cost <- function(x){
71     return (theta_1*x)}
72
73 ##### Define the utility function at mileage x from action i. If the agent
74   chooses to replace the engine in a bus,
75   it costs RC. If they choose _not_ to replace the engine, they incur the
76   cost of running the bus at mileage x.
75 u <- function(x,i){
76     -RC*i - cost(x*(1-i))
77 }
78
79
80 ##### The value function can be estimated through an iteration procedure. We
81   start with some initial guess for EV,
82   calculate EV with an expression that includes our initial guess of EV,
83   and continue iterating until the difference
84   between subsequent EV estimates becomes small.
85
86 ##### value.Iterate is a function to iteratively update the value function
87   according to the methodology in Rust. The function
88   takes as an argument a current estimate of EV, and returns an updated
89   estimate of EV. EV is an x by d matrix - we want the
90   EV values for each decision d at every possible current mileage value
91   x.

```

```

87 value.Iterate <- function(EV){
88
89   ### First iterate through each of the 30 x states
90   for (x in 1:31){
91     ## Update the EV value corresponding to not replacing the engine. There
92     ## are three contributions here – one from the
93     ## j = 0 case, one from the j = 1 case, and one from the j = 2 case.
94     ## Note the indexing here. When x =1, the state is
95     ## equal to 0 (this is the x that needs to be passed into u()), but we
96     ## want to grab the EV corresponding to the 1st entry.
97     EV2[x,1] <- log(exp(u(x-1,0)+beta*EV[x,1])+exp(u(x-1,1)+beta*EV[x,2]))*
98     theta_30
99     + log(exp(u(x,0)+beta*EV[x+1,1])+exp(u(x,1)+beta*EV[x+1,2]))*theta_31
100    + log(exp(u(x+1,0)+beta*EV[x+2,1])+exp(u(x+1,1)+beta*EV[x+2,2]))*theta_
101    32
102
103    ## Update the EV value corresponding to replacing the engine. When the
104    ## engine is replaced, x at the next period will
105    ## deterministically reset to x = 0.
106    EV2[x,2] <- log(exp(u(0,0)+beta*EV[1,1])+exp(u(0,1)+beta*EV[1,2]))
107  }
108
109  ## Return the updated EV values.
110  return(EV2)
111 }
112
113 ### Set a critical value for to measure the deviation between iterative
114 ### updates of EV. The distance between the two EV matrices
115 ### is the infinity norm of the difference
116 cri <- 10^(-8)
117
118 ### Set an initial value for the EV matrix (all 0s, EV), and another EV
119 ### object to hold the updated estimates, EV2.
120 EV <- matrix(100,33,2)
121 EV2 <- matrix(0,33,2)
122
123 ## While the infinity norm is less than the threshold, iterate
124 while(max(abs(EV-EV2))>cri){
125
126   ### Set the current EV to the previous updated EV
127   EV <- EV2
128   ### Compute a new updated EV by iterating on the current EV
129   EV2 <- value.Iterate(EV)
130 }
131
132 ### Do one last update to set EV equal to the last EV2
133 EV <- EV2
134
135 # get EV(x,i) for x=0,1,2,...,30
136 ### EV contains extra states, which we needed to compute the above
137 ### computation. Throw them away.
138 EV <- EV[1:31,]
139

```



```

131 ##### Plot the EV of both replacing the engine (i = 1) and not replacing the
      engine (i = 0) at every x
132 ##### between 1 and 30
133 df <- data.frame('x'=c(1:30, 1:30), 'EV'=c(EV[2:31,1], EV[2:31,2]), 'Action'
      = c(rep('i_0', 30), rep('i_1', 30)))
134
135 ##### Generate a plot that compares the EV of replacing the engine and not
      replacing the engine
136 ev_plot <- ggplot(df, aes(x=x, y=EV, color=Action)) + geom_point() + xlab('
      Mileage') + ylab('EV') +
137   ggtitle('EV_as_a_function_of_mileage_and_action\nat_x_between_1_and_30'
      ) +
138   theme(plot.title = element_text(hjust = 0.5))
139
140 ##### This is a plot to see the EV data in the attached rust matlab file. The
      state space is different than ours (200 states),
141 ##### so its hard to compare. Our's is linear (seems wrong), whereas the
      provided data is not. However, the first 30 states
142 ##### _do_ look approximately linear, so maybe we're not so far off.
143
144 df_rust <- data.frame('x'=c(seq(1,201), seq(1, 201)), 'EV'=c(data$EV[,1],
      data$EV[,2]), 'Action' = c(rep('i_0', 201),
145
146
147
148
149
150 #####
151 # Question 2.4 #
152 #####
153
154 ##### Calculate the mean mileage, mean time to engine replacement, max
      mileage, min mileage, and sd mileage over the whole sample
155 mean_x <- mean(x)
156 mean_engine_replacement_age <- mean(x[i == 1])
157 max_x <- max(x)
158 min_x <- min(x)
159 sd_x <- sd(x)

```

```

160 avg_replacements <- mean(apply(i, 2, function(x) {sum(x)}))
161
162 aggregate_stats <- c(mean_x, max_x, min_x, sd_x, mean_engine_)
163 aggregate_stats %>%
164   round(., 3) %>%
165   kable(., format='latex')
166
167 ### Calculate the per bus mean mileage, mean time to engine replacement,
168   max mileage, min mileage, and sd mileage
169 mean_x_per_bus <- apply(x, 2, function(x) {mean(x)})
170 max_x_per_bus <- apply(x, 2, function(x) {max(x)})
171 min_x_per_bus <- apply(x, 2, function(x) {min(x)})
172 sd_x_per_bus <- apply(x, 2, function(x) {sd(x)})
173 mean_engine_replacement_per_bus <- apply(x*i, 2, function(x) {sum(x)/sum(x
174   != 0)})
175 replacements <- apply(i, 2, function(x) {sum(x)})
176
177 ### Collate per bus information into a dataframe
178 per_bus_statistics <- data.frame(bus = seq(1, 100, 1),
179   mean_x_per_bus = mean_x_per_bus,
180   max_x_per_bus = max_x_per_bus,
181   min_x_per_bus = min_x_per_bus,
182   sd_x_per_bus = sd_x_per_bus,
183   mean_engine_replacement_per_bus = mean_
184     engine_replacement_per_bus,
185   replacements = replacements)
186
187 ### Create some plots
188 mean_mileage_plot <- ggplot(per_bus_statistics, aes(x=mean_x_per_bus)) +
189   geom_histogram() +
190   xlab('Mean_Mileage_(buckets_of_5,000_miles)') + ylab('Number_of_buses') +
191   ggtitle('Mean_mileage_across_buses') +
192   theme(plot.title = element_text(hjust = 0.5))
193 ggsave(mean_mileage_plot, file='mean_mileage_plot.png', height=4, width=4,
194   units='in')
195
196 max_mileage_plot <- ggplot(per_bus_statistics, aes(x=max_x_per_bus)) + geom
197   _histogram() +
198   xlab('Max_Mileage_(buckets_of_5,000_miles)') + ylab('Number_of_buses') +
199   ggtitle('Max_mileage_across_buses') +
200   theme(plot.title = element_text(hjust = 0.5))
201 ggsave(max_mileage_plot, file='max_mileage_plot.png', height=4, width=4,
202   units='in')
203
204 sd_mileage_plot <- ggplot(per_bus_statistics, aes(x=sd_x_per_bus)) + geom_
205   histogram() +
206   xlab('Mileage_Standard_Deviation_(buckets_of_5,000_miles)') + ylab('
207     Number_of_buses') +
208   ggtitle('Mileage_standard_deviation_across_buses') +
209   theme(plot.title = element_text(hjust = 0.5))
210 ggsave(sd_mileage_plot, file='sd_mileage_plot.png', height=4, width=4,
211   units='in')

```

```

202 | time_to_engine_replacement_plot <- ggplot(per_bus_statistics , aes(x=mean_
    engine_replacement_per_bus)) + geom_histogram() +
203 | xlab('Mean_Engine_Replacement_Mileage_(buckets_of_5,000_miles)') + ylab('
    Number_of_buses') +
204 | ggtitle('Mean_engine_replacement_mileage_across_buses') +
205 | theme(plot.title = element_text(hjust = 0.5))
206 | ggsave(time_to_engine_replacement_plot , file='time_to_engine_replacement_
    plot.png' , height=4, width=4, units='in')
207 |
208 | replacements_plot <- ggplot(per_bus_statistics , aes(x=replacements)) + geom
    _histogram() +
209 | xlab('Number_of_engine_replacements') + ylab('Number_of_buses') +
210 | ggtitle('Number_of_engine_replacements_across_buses') +
211 | theme(plot.title = element_text(hjust = 0.5))
212 | ggsave(replacements_plot , file='replacements_plot.png' , height=4, width=4,
    units='in')
213 |
214 |
215 | #####
216 | # Question 3.1 #
217 | #####
218 |
219 | ### This code will estimate the parameters beta , theta_1, and RC using the
    nested fixed-point algorithm
220 | ### described in Rust.
221 |
222 | ### A function to compute the probability of Zurcher's choices using the
    EVs we calculate using the EV calculation
223 | ### framework above. The probability of choosing different actions
    basically acts like a multichoice logit function.
224 | choice_prob.Estimate <- function(){
225 |
226 |     ### Initialize an empty matrix to hold choice probability estimates.
227 |     p_i <- matrix(0,31,2)
228 |
229 |     ### Iterate through all of the possible mileage states , x.
230 |     for (x in 1:31){
231 |         ## For each mileage state x, calculate the probability that Zurcher
            will choose i = 0.
232 |         p_i[x,1] <- exp(u(x-1,0)+beta*EV[x,1]) / (exp(u(x-1,0)+beta*EV[x,1])+exp(
            u(x-1,1)+beta*EV[x,2]))
233 |         ## Calculate the probability of choosing i = 1 at each state , which is
            just 1 - P(i = 0).
234 |         p_i[x,2] <- 1- p_i[x,1]
235 |     }
236 |
237 |     ### Return the updated p_i object
238 |     return(p_i)
239 | }
240 |
241 | ### A function to calculate the total log likelihood of the observed data
    given a set of parameters. This method assumes that
242 | ### the probabilities across periods and buses are independent , so we can
    just add up all of the log probabilities.

```

```

243 log.likelihood.Compute <- function() {
244
245     ### Initialize 0-valued variables to hold the log choice probability, the
        log transition probability,
246     ### and the sum of the two.
247     log_choice_prob <- 0
248     log_transition_prob <- 0
249     total <- 0
250
251     ### Iterate over buses
252     for (bus in 1:100) {
253         ### Iterate over time periods
254         for (t in 1:999) {
255             ### We special case mileage states greater than 30, since they are a
                bit strange in our data. Otherwise, we calculate
256             ### the choice probability using the current value of p_i according
                to the EV values we calculated to get the
257             ### choice probability. Take the log and add it to the current
                running value.
258             if (x[t,bus] <= 30){
259                 log_choice_prob <- log(p_i[x[t,bus]+1,i[t,bus]+1]) + log_choice_
                    prob
260             ### Do the same thing for our special cased, x > 30 case.
261             } else {
262                 log_choice_prob <- log(p_i[31,i[t,bus]+1]) + log_choice_prob
263             }
264
265             ### Calculate over the transitions for each bus the sum of the log
                transition probabilities. We have our estimates of
266             ### theta_3 given the empirical transition probabilities. So we can
                just grab that for each observed transition and add it
267             ### to the total log transition probability.
268
269             ### First we do the j = 0 case.
270             if (x[t+1,bus]-x[t,bus]==0) {
271                 log_transition_prob <- log(theta_30)+log_transition_prob
272             ### Then the j = 1 case.
273             } else if (x[t+1,bus]-x[t,bus]==1) {
274                 log_transition_prob <- log(theta_31)+log_transition_prob
275             ### And finally the j = 2 case.
276             } else if (x[t+1,bus]-x[t,bus]==2) {
277                 log_transition_prob <- log(theta_32)+log_transition_prob
278             }
279         }
280
281         ### Now, get the total log likelihood by adding up all of the
                transition components and the choice components.
282         total <- (log_choice_prob+log_transition_prob) + total
283     }
284     return (total)
285 }
286
287 ### Now we're actually going to use the nested fixed point algorithm to get
        the maximum likelihood estimates of the parameters

```

```

288 ##### that we care about. This process has three steps.
289
290 ##### Step 1: We would calculate theta_30, theta_31, and theta_31 directly
      from the data. This step is not in the loop, and we've
291 ##### actually already done this and it doesn't change, so we don't need to
      do it again.
292
293 ##### Step 2: Next, we are going to set up a grid over values of theta_1,
      beta, and RC that we will calculate the
294 ##### log likelihood to determine the maximum likelihood parameter values. We
      'll also initialize a dataframe
295 ##### to hold the parameter values and the log likelihoods.
296
297 theta_1_range <- seq(.01,.10,.01)
298 beta_range <- seq(.90,.99,.01)
299 RC_range <- seq(6,15,1)
300 likelihood <- data.frame( 'theta_1'=rep(0), 'beta'=rep(0), 'RC'=rep(0), 'log.
      likelihood'=rep(0))
301
302 ##### Step 3: Now we actually do the nested fixed point computation.
303
304 ##### Loop through theta_1
305 for (theta_1 in theta_1_range) {
306   ##### Loop through beta
307   for (beta in beta_range) {
308     ##### Loop through RC
309     for (RC in RC_range) {
310       print(paste(c(theta_1, beta, RC), collapse='_'))
311
312       ##### Initialize the EV functions to the initial values we used above.
313       EV <- matrix(100,33,2)
314       EV2 <- matrix(0,33,2)
315
316       ##### Iteratively compute the EV values.
317       while(max(abs(EV-EV2))>cri){
318         EV <- EV2
319         EV2 <- value.Iterate(EV)
320       }
321
322       EV <- EV2
323       EV <- EV[1:31,]
324
325       ##### Given these values of EV, calculate the choice probabilities
326       p_i <- choice.prob.Estimate()
327
328       ##### Given the EV values, the choice probabilities and the parameters,
      calculate
329       ##### the log-likelihood of the data.
330       likelihood <- rbind(likelihood, c(theta_1, beta, RC, log.likelihood.
      Compute()))
331     }
332   }
333 }
334

```

```

335 ##### Retrieve the row in the likelihood dataframe corresponding to the
      maximum likelihood estimate
336 likelihood <- likelihood[-1,]
337 parameter_estimates <- likelihood[which.max(likelihood[,4]),]
338
339 ##### Use these parameters and get the relevant estimate of EV and p_i
340 theta_1 = parameter_estimates$theta_1
341 beta = parameter_estimates$beta
342 RC = parameter_estimates$RC
343
344 EV <- matrix(100,33,2)
345 EV2 <- matrix(0,33,2)
346 ##### Iteratively compute the EV values.
347 while(max(abs(EV-EV2))>cri){
348     EV <- EV2
349     EV2 <- value.Iterate(EV)
350 }
351 EV <- EV2
352 EV <- EV[1:31,]
353 ##### Given these values of EV, calculated the choice probabilities
354 p_i <- choice.prob.Estimate()
355 ##### Given the EV values, the choice probabilities and the parameters,
      calculate
356 ##### the log-likelihood of the data.
357 likelihood <- rbind(likelihood,c(theta_1,beta,RC,log.likelihood.Compute()))
358
359 save(EV, p_i, likelihood, parameter_estimates, file='rust_estimate.Rdata')
360
361 #####
362 # Question 3.2 #
363 #####
364
365 ##### Now we will get estimates of the parameters using the Hotz and Miller
      conditional choice probability approach. This will
366 ##### allow us to compare these parameter estimates to those obtained using
      the Rust approach.
367
368 ##### First, we need to calculate the probability of the agent choosing
      either i = 0 or i = 1 based on the state that they find
369 ##### a given bus in, x, at some time period t. This will be the baseline
      that we use to try and find the best parameter values
370 ##### (i.e., which parameter values minimize the infinity norm between these
      true probabilities and the estimated probabilities)
371
372 ##### The probability matrix
373 p_ix <- matrix(0,33,2)
374 ##### The vector of how often the agent chooses i=1 given state x
375 ones <- vector()
376 ##### The vector of how often the agent finds a bus in state x
377 total <- vector()
378
379 ##### Loop through the states
380 for (state in 0:32){
381

```

```

382  ### For a given state, a will track how many times i = 1 and b will track
      how many times that state occurs.
383  ### Initialize them to 0 for the given state.
384  a <- 0
385  b <- 0
386
387  ### Loop over the buses
388  for (bus in 1:100){
389
390      ### Increment how many times the agent chooses i = 1 in state x
391      a <- sum(i[which(x[,bus]==state),bus]) + a
392      ### Increment how many times the state x occurs
393      b <- length(i[which(x[,bus]==state),bus]) + b
394  }
395
396  ### Add the most recent estimates to the vector.
397  ones[state+1] <- a
398  total[state+1] <- b
399  }
400
401  ### Based on the ones and total vectors, updated the choice probability
      matrix.
402  p_ix[,1] <- 1-ones/total
403  p_ix[,2] <- ones/total
404
405  ### The function below uses the Hotz and Miller method to estimate V and p_
      ix_hat for every state and period
406  ### given a set of model parameters (beta, theta_1, and RC).
407  approximate.V_pihat <- function() {
408      ### Initialize an empty valuation matrix
409      V <- matrix(0,33,2)
410      ### Initialize an empty conditional choice probability matrix
411      p_ix_hat <- matrix(0,33,2)
412
413      ### Iterate through the states
414      for (state in 0:30){
415          ### Initialize a and b, which will basically track a running total of V
              for different choices over simulations, to 0.
416          a = 0
417          b = 0
418          ### Iterate through the simulations. Note that ideal we would probably
              want to go more than one time step into the
419          ### future. However, because of the limitations in our dataset, we only
              go one time step forward. This is mainly because
420          ### it's unclear how we would draw i (the choice) for states that do
              not appear in our data (i.e., x = 34).
421          for (s in 1:S){
422              ## Conditional on choosing i = 0, simulate the next state that a
                  given bus will end up in by drawing from the
423              ## transition probabilities.
424              x_prime_0 = state + sample(c(0,1,2),1,replace = T, prob = c(theta_30,
                  theta_31,theta_32))
425              ## Conditional on choosing i = 0 and ending up in some state in the
                  next time period, randomly simulate a draw from

```

```

426     ## i based on the conditional choice probabilities
427     i_prime_0 = sample(c(0,1),1,replace=T,prob = c(p_ix[x_prime_0+1,1],p_
428         ix[x_prime_0+1,2]))
429     # Figure out the expected utility from this truncated sequence of
430     choices.
431     a = (u(state,0) + beta*(u(x_prime_0,i_prime_0)+gamma-log(p_ix[x_prime
432         _0+1,i_prime_0+1]))) + a
433     ## Conditional on choosing i = 1, we don't need to simulate the next
434     state that a bus will end up in. It will always
435     be x = 0. So we jump right to simulating the draw from i for x =
436     0.
437     i_prime_1 = sample(c(0,1),1,replace=T,prob = c(p_ix[1,1],p_ix[1,2]))
438     ## Figure out the expected utility from this truncated sequence of
439     choices.
440     b = (u(state,1) + beta*(u(0,i_prime_1)+gamma-log(p_ix[1,i_prime_1+1])
441         )) + b
442 }
443
444     ## Set the value of V to be the average over all S of our simulations
445     for both the i = 0 and i = 1 choices.
446     V[state+1,1] = a/S
447     V[state+1,2] = b/S
448     ## Use the multinomial logit-esque probability expression to figure out
449     the probability of choosing i = 0 or i = 1
450     ## given that the bus is in state x.
451     p_ix_hat[state+1,1] <- exp(V[state+1,1])/(exp(V[state+1,1])+exp(V[state
452         +1,2]))
453     p_ix_hat[state+1,2] <- 1- p_ix_hat[state+1,1]
454 }
455
456 # Put final output into a list and return it
457 results <- list('V' = V, 'p_ix_hat' = p_ix_hat)
458 return(results)
459 }
460
461 #### Specify a number of constants that will be used in the Hotz and Miller
462 algorithm:
463 #### S: The number of "simulations" to do per state / decision
464 #### gamma: This should be Euler's constant
465 #### theta_1_range: The range of theta_1 values to test
466 #### beta_range: The range of beta values to test
467 #### RC_range: The range of RC values to test
468 S = 1000
469 gamma = .577
470 theta_1_range <- seq(.01,.10,.01)
471 beta_range <- seq(.90,.99,.01)
472 RC_range <- seq(6,15,1)
473
474 #### Initialize a dataframe to hold different parameter combinations and the
475 infinity-norm between the actual conditional
476 choice probabilities and the estimated ones
477 difference <- data.frame('theta_1'=rep(0),'beta'=rep(0),'RC'=rep(0),'
478     difference'=rep(0))

```



```

467
468 #### Loop through theta_1
469 for (theta_1 in theta_1_range) {
470   #### Loop through theta_2
471   for (beta in beta_range) {
472     #### Loop through RC
473     for (RC in RC_range) {
474       # Check progress
475       print(paste(c(theta_1, beta, RC), collapse='_'))
476
477       #### Get estimates of V and P_ix_hat using the Hotz and Miller method
478       v_and_p_ix_hat <- approximate.V_pixhat()
479       V = v_and_p_ix_hat$V
480       p_ix_hat <- v_and_p_ix_hat$p_ix_hat
481
482       #### Now that we have a full conditional choice probability matrix,
483       calculate the infinity norm (i.e., largest
484       #### absolute difference between the empirical conditional choice
485       probabilities and those estimated with the
486       #### given parameters)
487       difference <- rbind(difference, c(theta_1, beta, RC, max(abs(p_ix[1:31,] -
488         p_ix_hat[1:31,])))
489     }
490   }
491 }
492
493 #### Find the set of parameters that minimizes this difference
494 difference <- difference[-1,]
495 parameter_estimates <- difference[which.min(difference[,4]),]
496
497 #### Use these parameters and get the relevant estimate of V and p_ix_hat
498 theta_1 = parameter_estimates$theta_1
499 beta = parameter_estimates$beta
500 RC = parameter_estimates$RC
501 best_guesses <- approximate.V_pixhat()
502 V <- best_guesses$V
503 p_ix_hat <- best_guesses$p_ix_hat
504
505 save(V, p_ix_hat, difference, parameter_estimates, file='hotz_and_miller_
506   estimate.Rdata')
507
508 #####
509 # Question 3.3 #
510 #####
511
512 #####
513 # Question 3.4 #
514 #####
515
516 #### This function simulates, for one agent, a sequence of state transitions
517 and also engine replacement decisions
518 simulate_sequence <- function(n_periods) {

```

```

516   ### Initialize empty vectors to hold states and engine replacement
      transitions
517   x_values <- rep(0, n_periods)
518   i_values <- rep(0, n_periods)
519   ### Every bus starts at state 0
520   x_values[1] <- 0
521   ### Go through the progression
522   for (j in 1:length(x_values)) {
523     ### Make a decision based on current state
524     i_values[j] = sample(c(0,1),1,replace=T,prob = c(p_ix_hat[x_values[j] +
      1,1],p_ix_hat[x_values[j] + 1,2]))
525     ### If decision is to not replace, continue on and increment x randomly
526     if (i_values[j] == 0) {
527       x_values[j+1] = x_values[j] + sample(c(0,1,2),1,replace = T, prob = c
      (theta_30,theta_31,theta_32))
528     } else {
529       x_values[j+1] = 0
530     }
531   }
532 }
533 ## Generate a decision for the last period, even though we never see the
      fruits of that decision
534 i_values[length(i_values)] = sample(c(0,1),1,replace=T,prob = c(p_ix_hat[
      x_values[length(x_values)] + 1,1],
535                                p_ix_hat[
      x_
      values
      [
      length
      (x_
      values
      )] +
      1,2])
      )

536   ### Return the states and replacement decisions in a list
537   results <- list('x_values' = x_values, 'i_values' = i_values)
538   return(results)
539 }
540
541 ### Given a set of parameters, this function generates period-by-period
      demand estimates for new buses (e.g.,
542 ### how many buses will get their engine replaced in each period)
543 estimate_demand <- function(n_sims, n_buses, n_periods) {
544
545   ### Initialize a vector to hold simulated demand
546   simulated_demand_total <- rep(0, n_periods)
547
548   ### Run a bunch of simulations and simulate engine replacement decisions
549   for (j in 1:n_sims) {
550     simulated_demand_total = simulated_demand_total + simulate_sequence(n_
      periods)$i_values
551   }
552

```

```

553     ### Divide by the number of sims to get averages, multiply by number of
554     buses (this works because
555     ###buses are independent). Then return what we get.
556     return((n_buses/n_sims)*simulated_demand_total)
557 }
558 ### Get demand as a function of RC for the first bus
559
560 ## Specify the range of RCs, as well as constants.
561 RC_range = seq(0, 15, .25)
562 n_periods = 15
563 n_sims = 1000
564 n_buses = 100
565 load('hotz_and_miller_estimate.Rdata')
566
567 ## Initialize an empty dataframe to hold results
568 estimated_demand_df <- data.frame(time_period = c(),
569                                   RC = c(),
570                                   demand = c(),
571                                   engine = c())
572
573 # Loop through the RCs, then estimate the probabilities using the Rust
574   method, then do simulation.
575 for (j in RC_range) {
576   RC = j
577
578   ### Set a critical value for to measure the deviation between iterative
579   updates of EV. The distance between the two EV matrices
580   ### is the infinity norm of the difference
581   cri <- 10^(-8)
582
583   ### Set an initial value for the EV matrix (all 0s, EV), and another EV
584   object to hold the updated estimates, EV2.
585   EV <- matrix(100,33,2)
586   EV2 <- matrix(0,33,2)
587
588   ## While the infinity norm is less than the threshold, iterate
589   while(max(abs(EV-EV2))>cri){
590     ### Set the current EV to the previous updated EV
591     EV <- EV2
592     ### Compute a new updated EV by iterating on the current EV
593     EV2 <- value.Iterate(EV)
594   }
595
596   ### Do one last update to set EV equal to the last EV2
597   EV <- EV2
598
599   # get EV(x,i) for x=0,1,2,...,30
600   ### EV contains extra states, which we needed to compute the above
601   computation. Throw them away.
602   EV <- EV[1:31,]
603
604   ### Get estimated probability based on the above EV

```

```

602 p_ix_hat <- choice.prob.Estimate()
603 #### Estimate demand using that probability
604 estimated_demand <- estimate_demand(n_sims, n_buses, n_periods)
605
606 #### Add this estimate to a temp dataframe
607 estimated_demand_df_temp <- data.frame(time_period = seq(1, n_periods, 1)
608                                     ,
609                                     RC = rep(RC, n_periods),
610                                     demand = estimated_demand,
611                                     engine = rep('Engine_1', n_periods))
612 #### Collate temp dataframe to full dataframe
613 estimated_demand_df <- rbind(estimated_demand_df, estimated_demand_df_
614                               temp)
615
616 }
617
618 #### Reset theta_1 to the "new engine", redo the exercise above.
619 theta_1 = .02
620
621 #### Loop through RCs
622 for (j in RC_range) {
623   RC = j
624
625   #### Set a critical value for to measure the deviation between iterative
626   updates of EV. The distance between the two EV matrices
627   #### is the infinity norm of the difference
628   cri <- 10^(-8)
629
630   #### Set an initial value for the EV matrix (all 0s, EV), and another EV
631   object to hold the updated estimates, EV2.
632   EV <- matrix(100,33,2)
633   EV2 <- matrix(0,33,2)
634
635   ## While the infinity norm is less than the threshold, iterate
636   while(max(abs(EV-EV2))>cri){
637
638     #### Set the current EV to the previous updated EV
639     EV <- EV2
640     #### Compute a new updated EV by iterating on the current EV
641     EV2 <- value.Iterate(EV)
642   }
643
644   #### Do one last update to set EV equal to the last EV2
645   EV <- EV2
646
647   # get EV(x,i) for x=0,1,2,...,30
648   #### EV contains extra states, which we needed to compute the above
649   computation. Throw them away.
650   EV <- EV[1:31,]
651
652   #### Get probability estimates based on EV
653   p_ix_hat <- choice.prob.Estimate()
654   #### Estimate demand
655   estimated_demand <- estimate_demand(n_sims, n_buses, n_periods)

```

```

651
652 ### Add to temp dataframe
653 estimated_demand_df_temp <- data.frame(time_period = seq(1, n_periods, 1)
654                                     ,
655                                     RC = rep(RC, n_periods),
656                                     demand = estimated_demand,
657                                     engine = rep('Engine_2', n_periods)
658                                     )
659
660 ### Collate to full dataframe
661 estimated_demand_df <- rbind(estimated_demand_df, estimated_demand_df_
662                               temp)
663 }
664
665 ### For a reduced set of RCs, see the period-by-period demand
666 estimated_demand_df %>%
667   filter(RC %in% c(1, 3, 5, 7, 10)) %>%
668   ggplot(., aes(x=time_period, y=demand, color=as.factor(RC))) + geom_line()
669   +
670   facet_wrap(~engine)
671
672 ### Aggregate over periods to get demand as a function of RC for different
673   thetas.
674 estimated_demand_df %>%
675   group_by(RC, engine) %>%
676   summarise(total_demand = sum(demand)) %>%
677   ungroup() %>%
678   ggplot(., aes(x=RC, y=total_demand, color=engine)) + geom_line()
679
680 #####
681 # Question 3.5 #
682 #####

```

Listing 1: ./rust.R